

# Network Security - part I

Srdjan Matic



16th September 2024

based on the material prepared by:  
Emiliano de Cristofaro and Juan Caballero

# Outline

## Background

## Issues of the Different Network Layers

The Data Link Layer

The Network Layer

The Transport Layer

## Scanning and OS Fingerprinting

Port Scanning

OS Fingerprinting

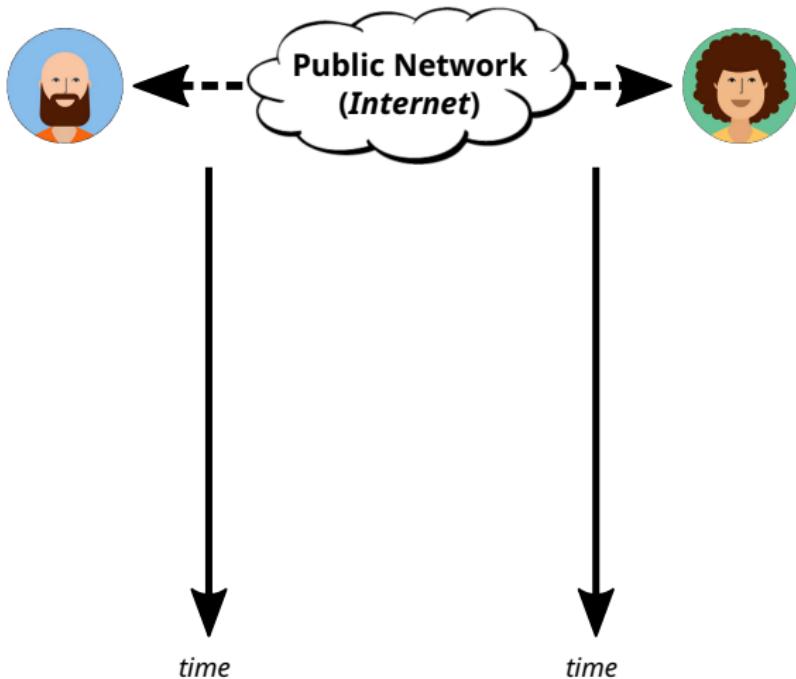
## Denial of Service

Denial of Service (DoS)

Spoofing

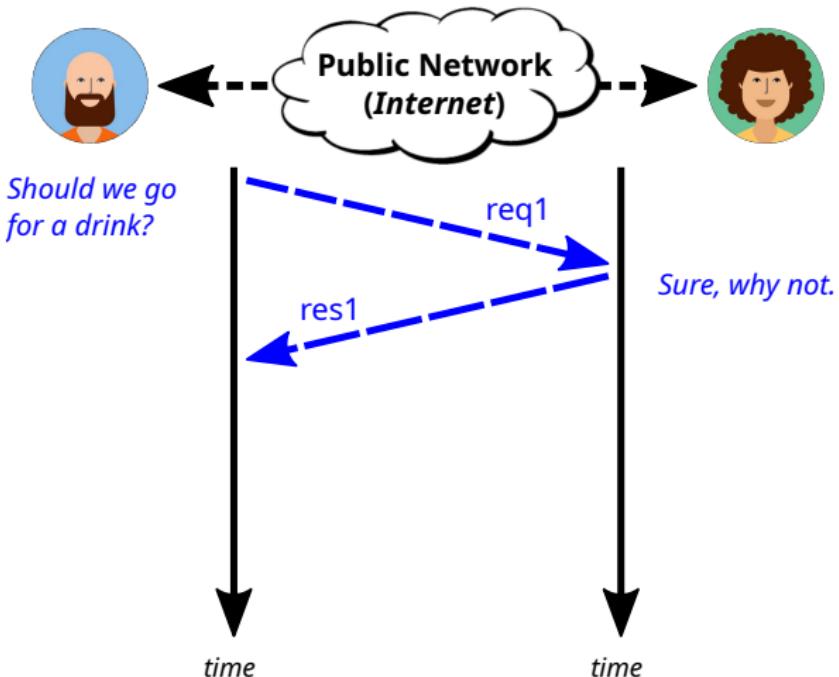
Packet Flooding

# An Example of a (Network) Communication



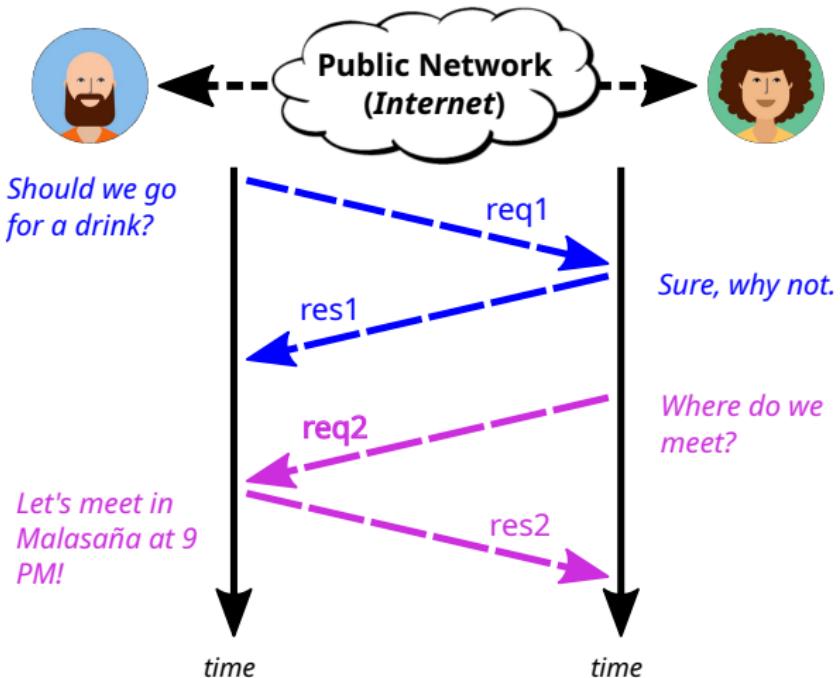
- Only two participants.

# An Example of a (Network) Communication



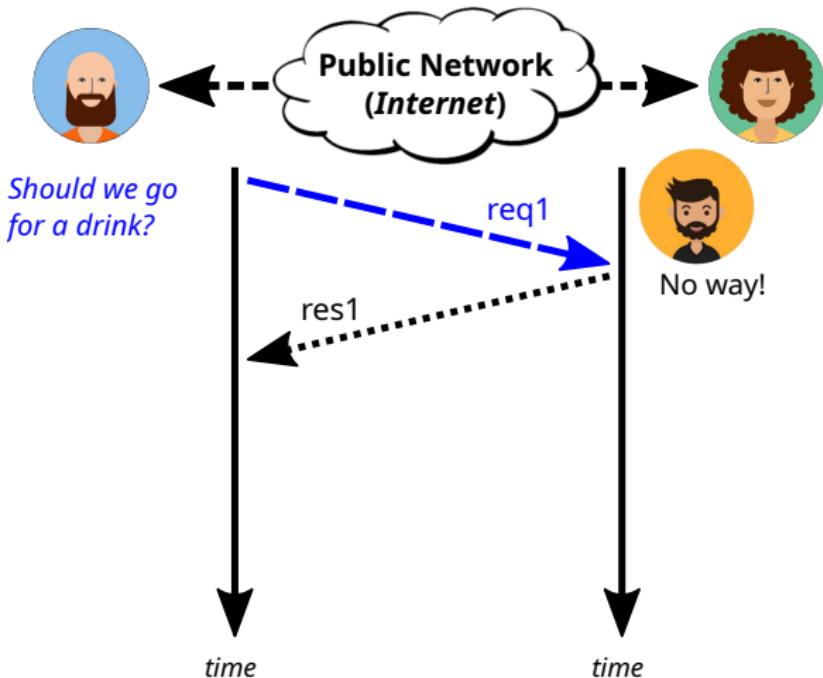
- Only two participants.
- Bi-directional communications.

# An Example of a (Network) Communication



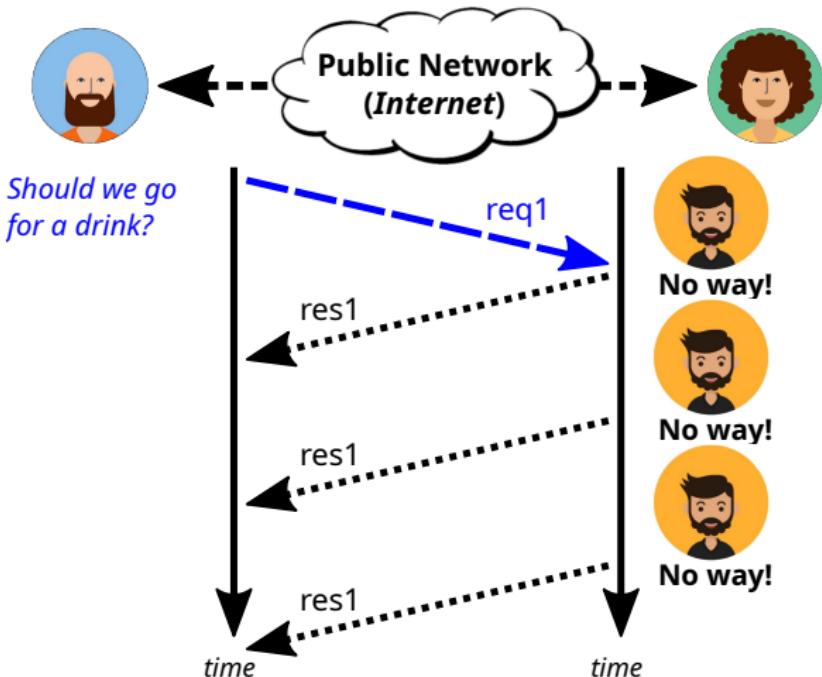
- Only two participants.
- Bi-directional communications.
- Request  $\implies$  response.

# An Example of a (Network) Communication



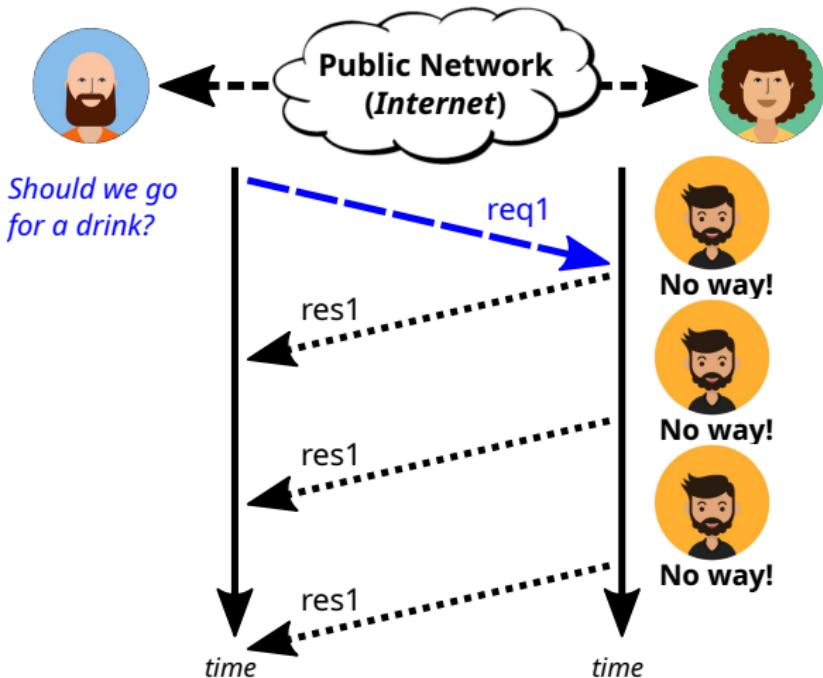
- Only two participants.
- Bi-directional communications.
- Request  $\implies$  response.

# An Example of a (Network) Communication



- Only two participants.
- Bi-directional communications.
- Request  $\implies$  response.

# An Example of a (Network) Communication



- Only two participants.
- Bi-directional communications.
- Request  $\implies$  response.
- *Confidentiality, integrity, authenticity.*

# Communications over Internet

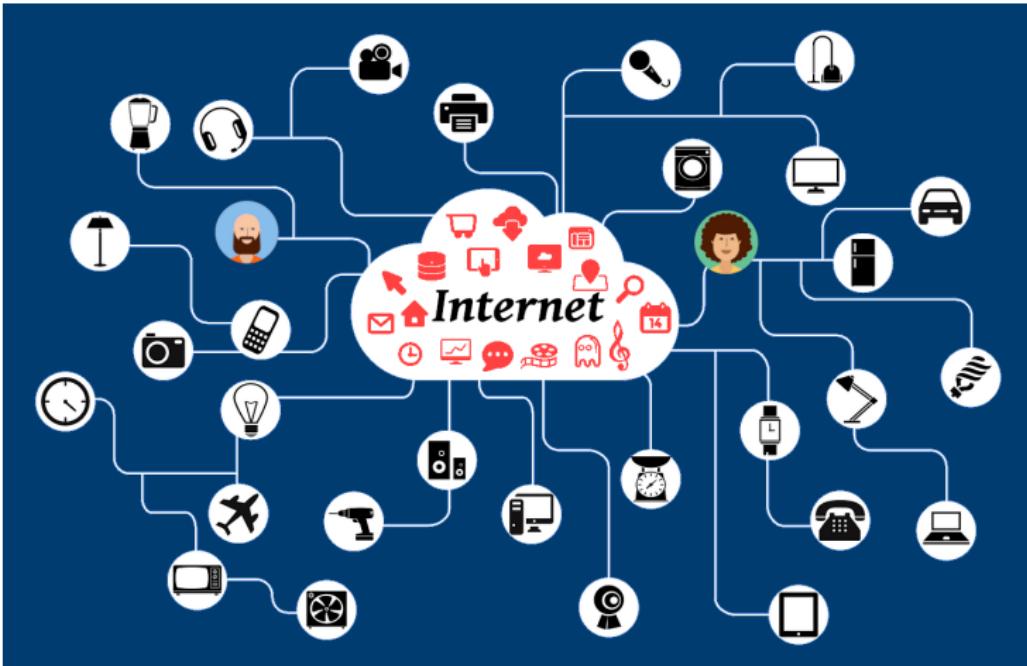


# Communications over Internet



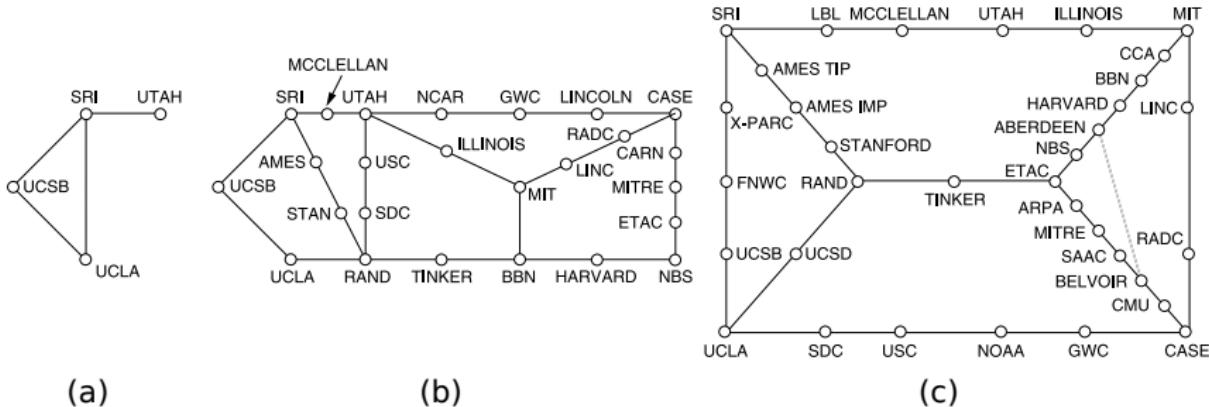
1. Internet: thousands of protocols and services.

# Communications over Internet



1. Internet: thousands of protocols and services.
2. Users: billions of devices.

# History of Internet



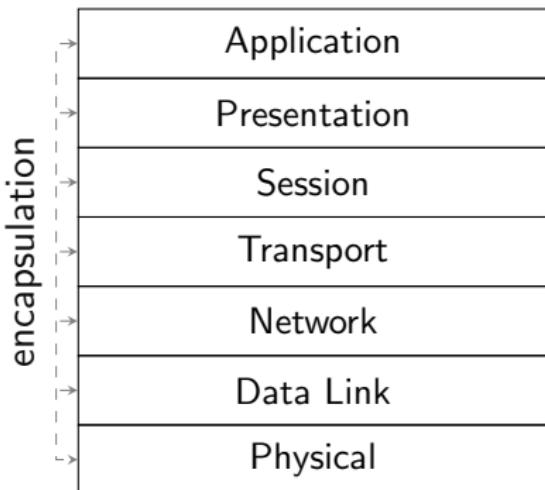
- ARPANET growth:
  - (a) December 1969. (b) July 1970. (c) September 1972<sup>1</sup>.
- Introduction of new protocols:
  - TCP/IP (1974)
  - DNS (1980s)

<sup>1</sup>Source: A. Tanenbaum, "Computer Networks", 5ed (2011)

# ISO/OSI Model

The *ISO OSI (Open Systems Interconnection) Reference Model* was developed in 1983 to allow connecting open systems.

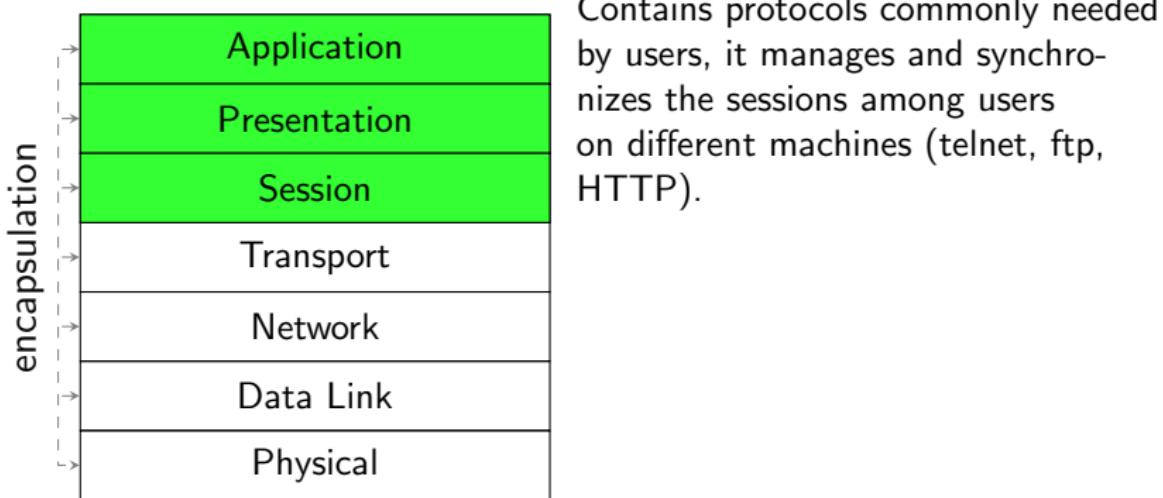
It contains seven layers, and each layer is defined as a separate international standard.



# ISO/OSI Model

The *ISO OSI (Open Systems Interconnection) Reference Model* was developed in 1983 to allow connecting open systems.

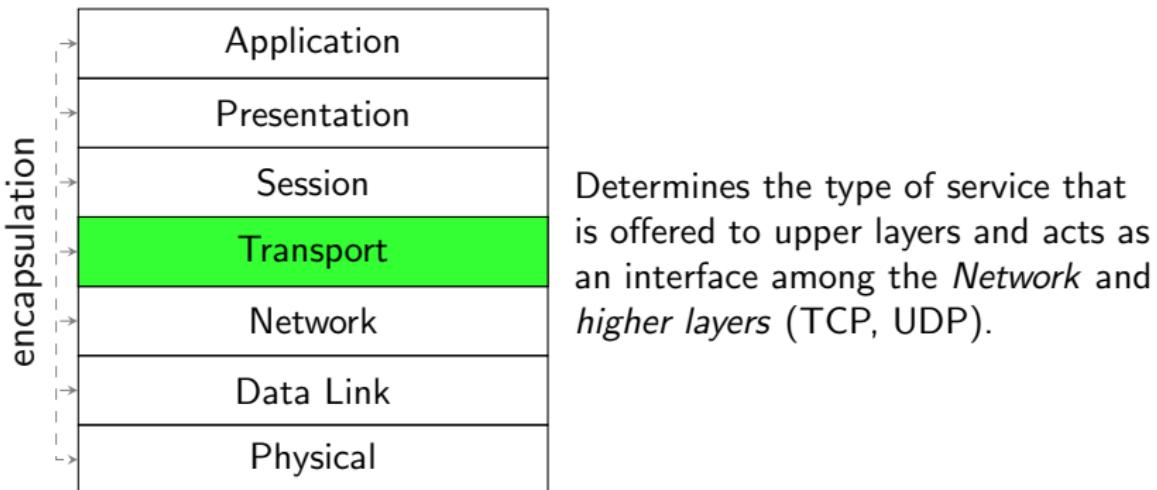
It contains seven layers, and each layer is defined as a separate international standard.



# ISO/OSI Model

The *ISO OSI (Open Systems Interconnection) Reference Model* was developed in 1983 to allow connecting open systems.

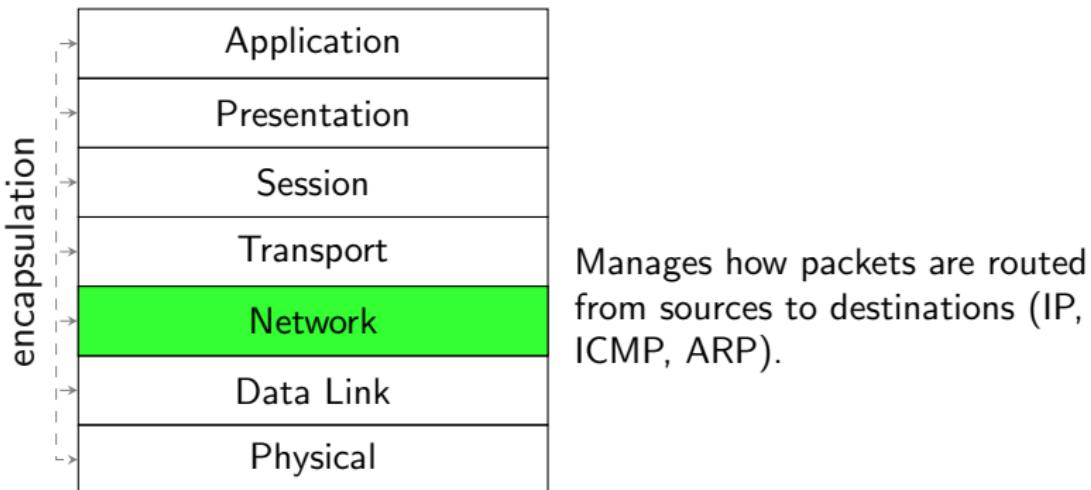
It contains seven layers, and each layer is defined as a separate international standard.



# ISO/OSI Model

The *ISO OSI (Open Systems Interconnection) Reference Model* was developed in 1983 to allow connecting open systems.

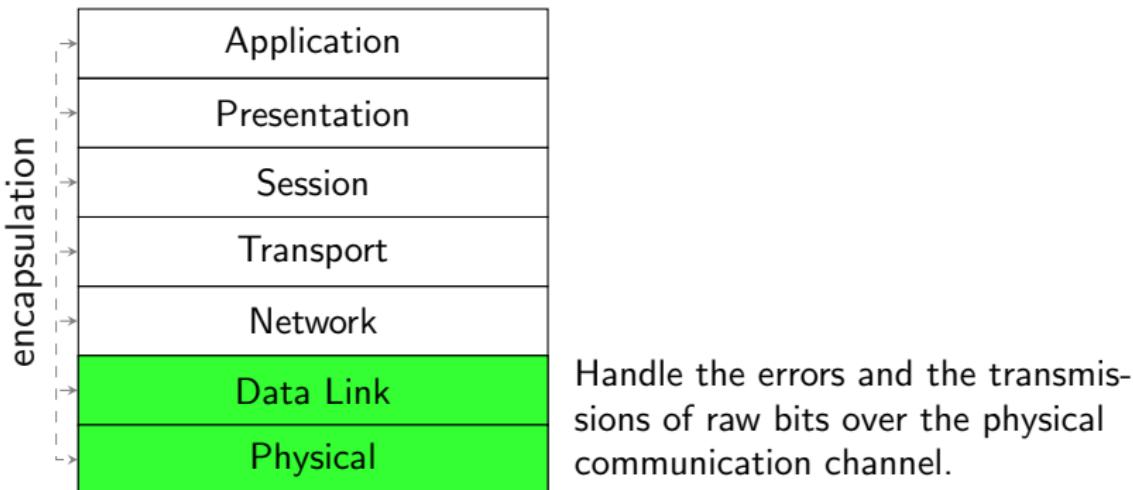
It contains seven layers, and each layer is defined as a separate international standard.



# ISO/OSI Model

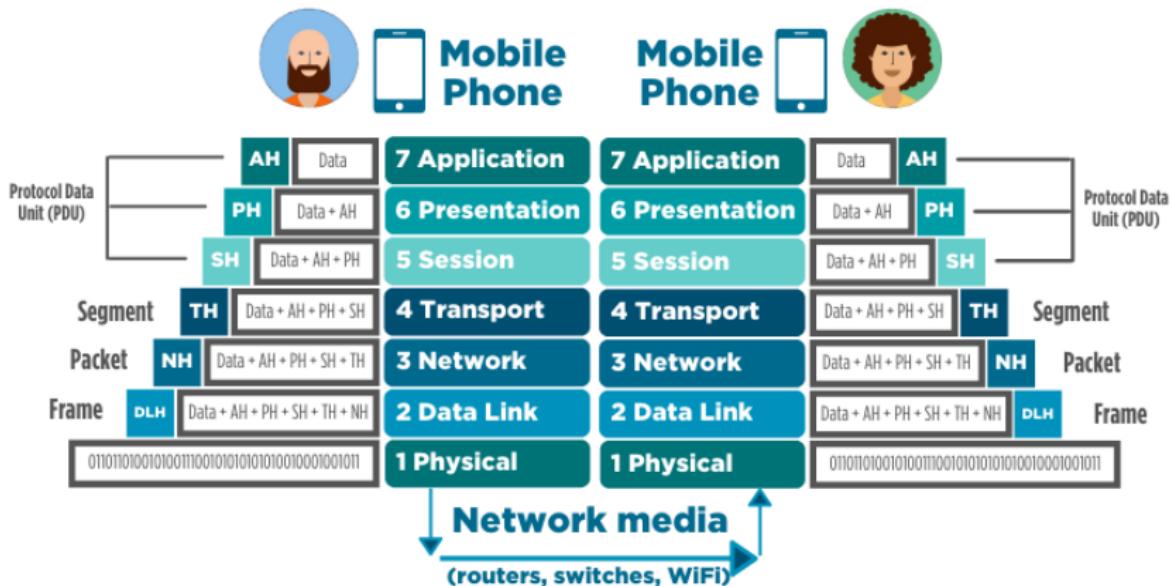
The *ISO OSI (Open Systems Interconnection) Reference Model* was developed in 1983 to allow connecting open systems.

It contains seven layers, and each layer is defined as a separate international standard.



# Information Flow

How messages are exchanged in in the ISO/OSI model:



- ↓ data **encapsulation**

- ↑ data **de-capsulation**

# Possible Source of Issues

1. No authentication  $\Rightarrow$  **spoofing, hijacking**

## Possible Source of Issues

1. No authentication  $\Rightarrow$  spoofing, hijacking
2. No timestamps  $\Rightarrow$  replay attacks

## Possible Source of Issues

1. No authentication  $\Rightarrow$  spoofing, hijacking
2. No timestamps  $\Rightarrow$  replay attacks
3. Insufficient checks  $\Rightarrow$  poisoning

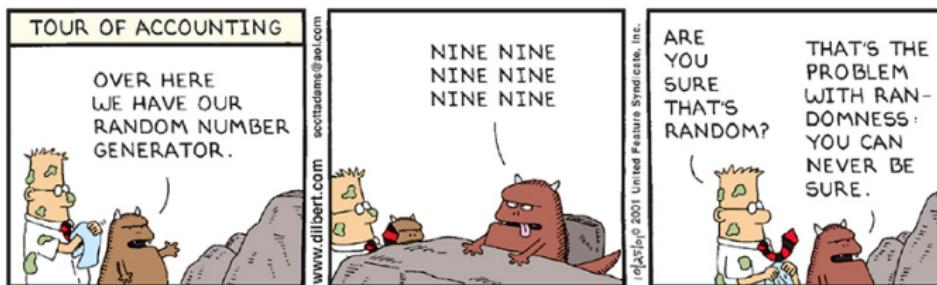
# Possible Source of Issues

1. No authentication  $\Rightarrow$  spoofing, hijacking
2. No timestamps  $\Rightarrow$  replay attacks
3. Insufficient checks  $\Rightarrow$  poisoning
4. Poor implementations  $\Rightarrow$  memory corruption, weak randomness



# Possible Source of Issues

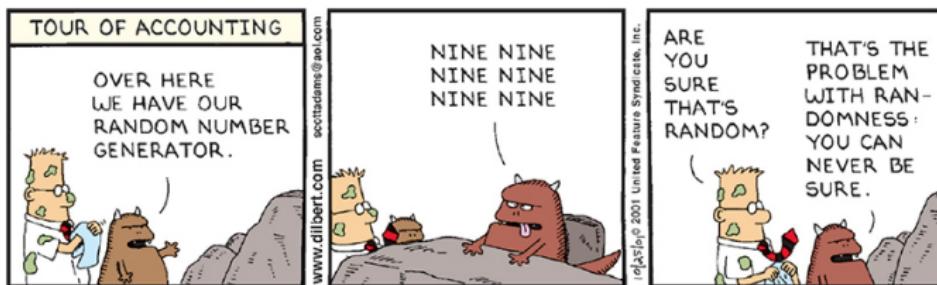
1. No authentication  $\Rightarrow$  spoofing, hijacking
2. No timestamps  $\Rightarrow$  replay attacks
3. Insufficient checks  $\Rightarrow$  poisoning
4. Poor implementations  $\Rightarrow$  memory corruption, weak randomness



Why do these issues still persist?

# Possible Source of Issues

1. No authentication  $\Rightarrow$  spoofing, hijacking
2. No timestamps  $\Rightarrow$  replay attacks
3. Insufficient checks  $\Rightarrow$  poisoning
4. Poor implementations  $\Rightarrow$  memory corruption, weak randomness

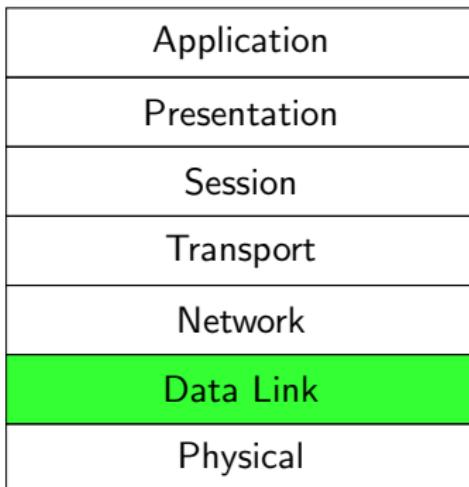


## Why do these issues still persist?

Legacy requirements:

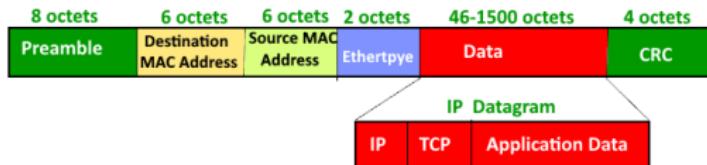
fixing issues might require to re-deploy **millions** of servers.

# The Data Link Layer

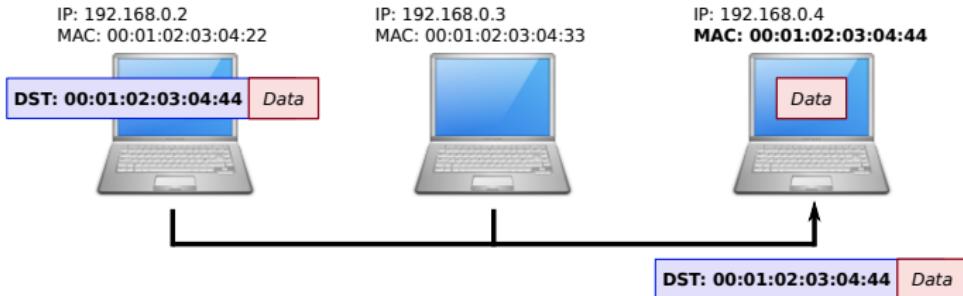


# Local Network - Delivery

To reach machines in the *same network*, packets (e.g., IP datagrams) are encapsulated in *Data Link frames*. The Data Link layer implements basic functionalities such as *collision avoidance* and *error handling* (CRC, checksum). The most common type is the *Ethernet frame*.

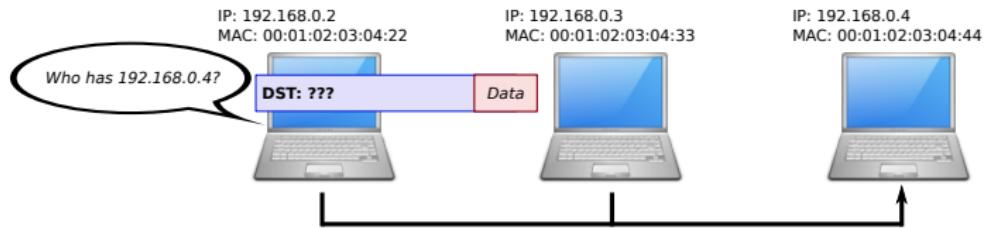


The frame includes the MAC address of the destination. Frames are **broadcasted** and the network cards pass to the operating system only those frames that *are indented for a particular host* (i.e., with the same MAC address of the host).



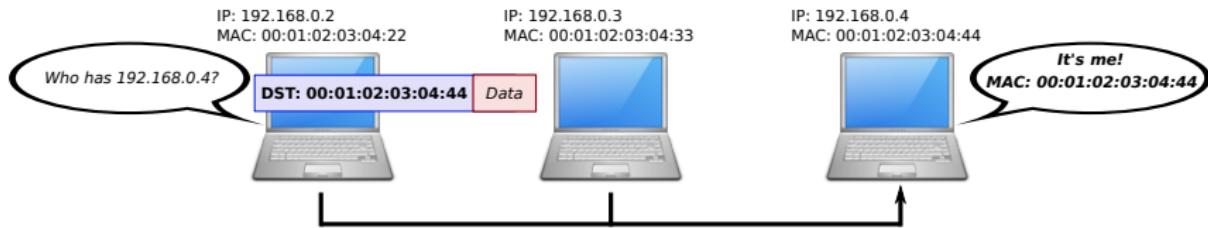
# The Address Resolution Protocol (ARP)

Hosts leverage ARP to learn which MAC address is associated to a given IP.  
The ARP response is cached when the host needs to contact the same machine.



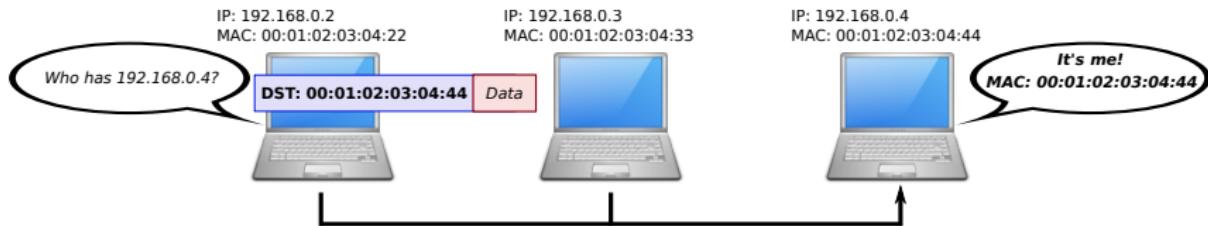
# The Address Resolution Protocol (ARP)

Hosts leverage ARP to learn which MAC address is associated to a given IP.  
The ARP response is cached when the host needs to contact the same machine.



# The Address Resolution Protocol (ARP)

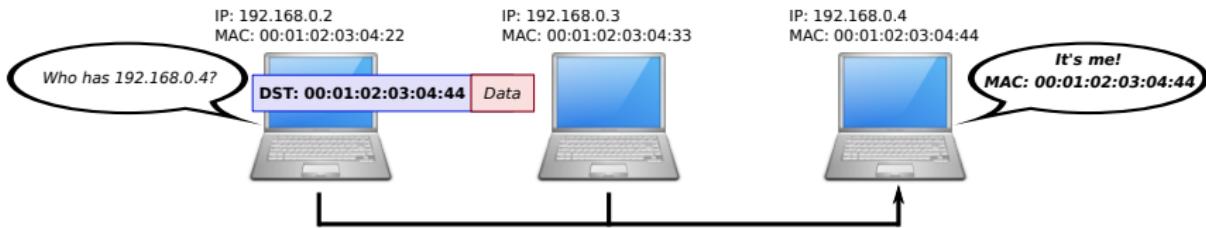
Hosts leverage ARP to learn which MAC address is associated to a given IP.  
The ARP response is cached when the host needs to contact the same machine.



Can you spot any issue?

# The Address Resolution Protocol (ARP)

Hosts leverage ARP to learn which MAC address is associated to a given IP.  
The ARP response is cached when the host needs to contact the same machine.

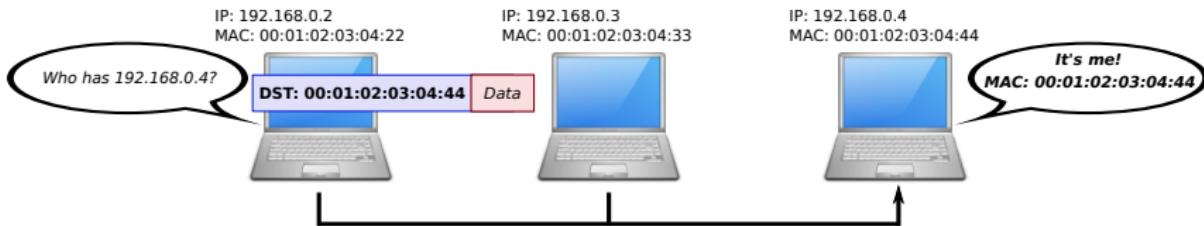


ARP does not provide authentication:

- attackers compete with legitimate hosts to provide a *fake MAC/IP mapping*  
⇒ the traffic gets redirected to the attacker

# The Address Resolution Protocol (ARP)

Hosts leverage ARP to learn which MAC address is associated to a given IP.  
The ARP response is cached when the host needs to contact the same machine.

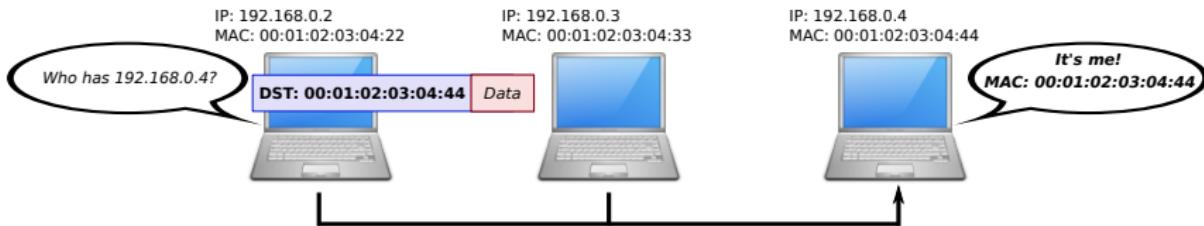


ARP does not provide authentication:

- attackers compete with legitimate hosts to provide a *fake MAC/IP mapping*  
⇒ the traffic gets redirected to the attacker
- legitimate ARP replies will restore the IP/MAC mapping

# The Address Resolution Protocol (ARP)

Hosts leverage ARP to learn which MAC address is associated to a given IP.  
The ARP response is cached when the host needs to contact the same machine.

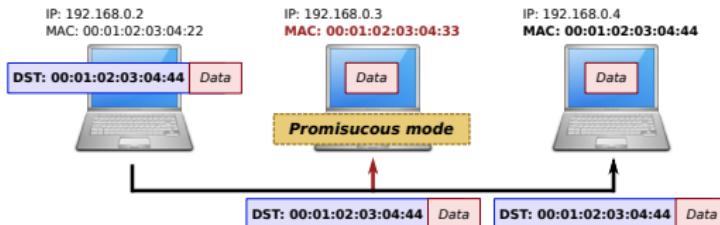


ARP does not provide authentication:

- attackers compete with legitimate hosts to provide a *fake MAC/IP mapping*  
⇒ the traffic gets redirected to the attacker
- legitimate ARP replies will restore the IP/MAC mapping
  - solution: the attacker *continuously sends* spoofed ARP replies

# Local Network - Sniffing

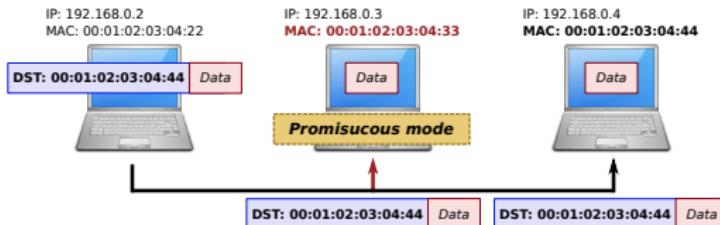
All the hosts on the local network “see” the traffic, but they discard it. If a host sets his network card in *promiscuous mode*, it will gain access to all the traffic.  $\Rightarrow$  the host becomes a **network sniffer**.



Additional information: <https://www.just.edu.jo/~tawalbeh/nyit/incs745/presentations/Sniffers.pdf>

# Local Network - Sniffing

All the hosts on the local network “see” the traffic, but they discard it. If a host sets his network card in *promiscuous mode*, it will gain access to all the traffic.  $\Rightarrow$  the host becomes a **network sniffer**.



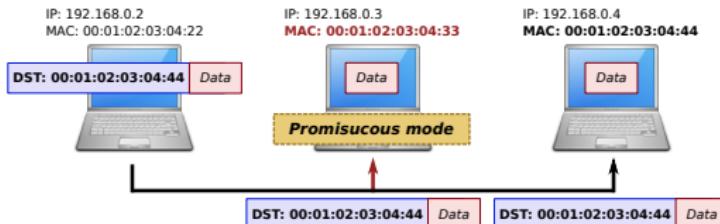
What is the issue?

Unencrypted traffic (e.g, passwords, usernames, files content)

Additional information: <https://www.just.edu.jo/~tawalbeh/nyit/incs745/presentations/Sniffers.pdf>

# Local Network - Sniffing

All the hosts on the local network “see” the traffic, but they discard it. If a host sets his network card in *promiscuous mode*, it will gain access to all the traffic.  $\Rightarrow$  the host becomes a **network sniffer**.



What is the issue?

Unencrypted traffic (e.g, passwords, usernames, files content)

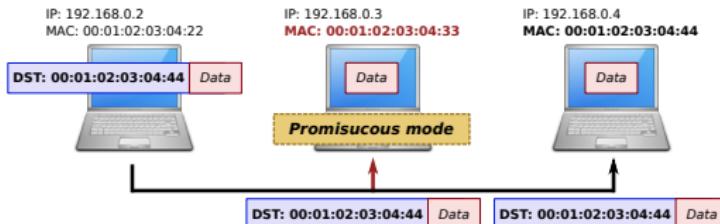
Sniffer detection:

- **ping:** send a frame with the correct IP and incorrect MAC  $\rightarrow$  only sniffers will reply;

Additional information: <https://www.just.edu.jo/~tawalbeh/nyit/incs745/presentations/Sniffers.pdf>

# Local Network - Sniffing

All the hosts on the local network “see” the traffic, but they discard it. If a host sets his network card in *promiscuous mode*, it will gain access to all the traffic.  $\Rightarrow$  the host becomes a **network sniffer**.



What is the issue?

Unencrypted traffic (e.g, passwords, usernames, files content)

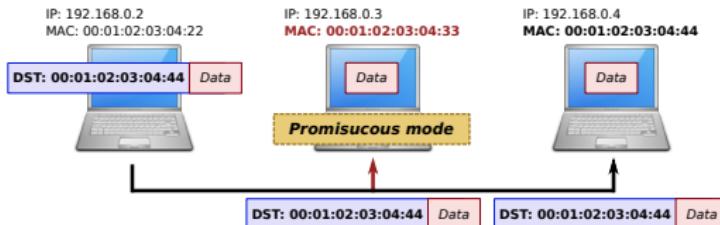
Sniffer detection:

- **ping:** send a frame with the correct IP and incorrect MAC  $\rightarrow$  only sniffers will reply;
- **latency:** send huge amount of data, and check latency on the candidate sniffer *before* and *during* the data flooding;

Additional information: <https://www.just.edu.jo/~tawalbeh/nyit/incs745/presentations/Sniffers.pdf>

# Local Network - Sniffing

All the hosts on the local network “see” the traffic, but they discard it. If a host sets his network card in *promiscuous mode*, it will gain access to all the traffic. ⇒ the host becomes a **network sniffer**.



What is the issue?

Unencrypted traffic (e.g, passwords, usernames, files content)

Sniffer detection:

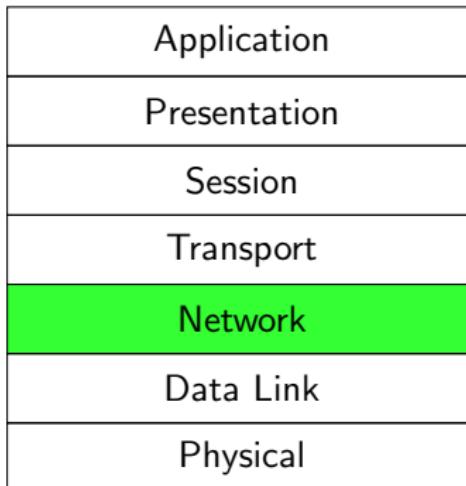
- **ping**: send a frame with the correct IP and incorrect MAC → only sniffers will reply;
- **latency**: send huge amount of data, and check latency on the candidate sniffer *before* and *during* the data flooding;

Sniffer prevention:

- use **switched Ethernet** instead of a hub: maps MAC(s) ↔ physical port(s) and forwards the frames to the right host.
  - can be bypassed with MAC flooding

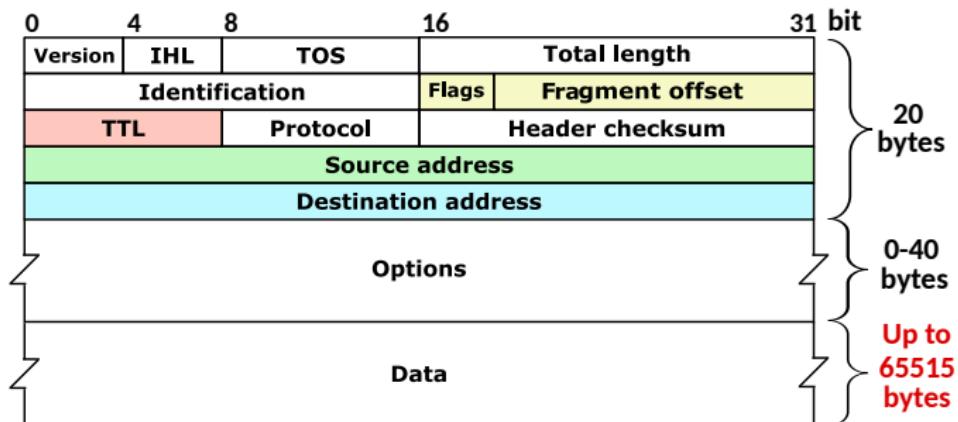
Additional information: <https://www.just.edu.jo/~tawalbeh/nyit/incs745/presentations/Sniffers.pdf>

# The Network Layer



# IPv4 Datagram

Encapsulates data from higher layers - best effort, no reliability guaranteed.

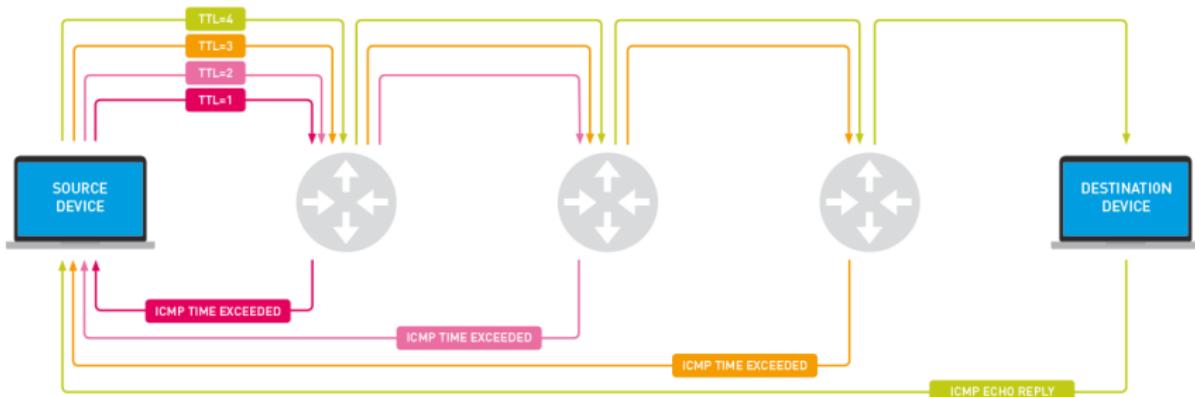


- **Source/Destination addresses**: where the packet originates/goes to; no authentication provided.
- **Time To Live (TTL)**: decreased at each router; ensures that datagrams do not roam forever.
- **Flags + Fragment Offset**: datagrams can be fragmented if the *lower transmission layers* can handle only packets of a certain size. These fields are used to reassemble the datagram on the receiving end.

# ICMP and Traceroute

ICMP (Internet Control Message Protocol) is used to *test connectivity*.

About a dozen types of ICMP messages are defined, and each ICMP message type is carried encapsulated in an IP packet.

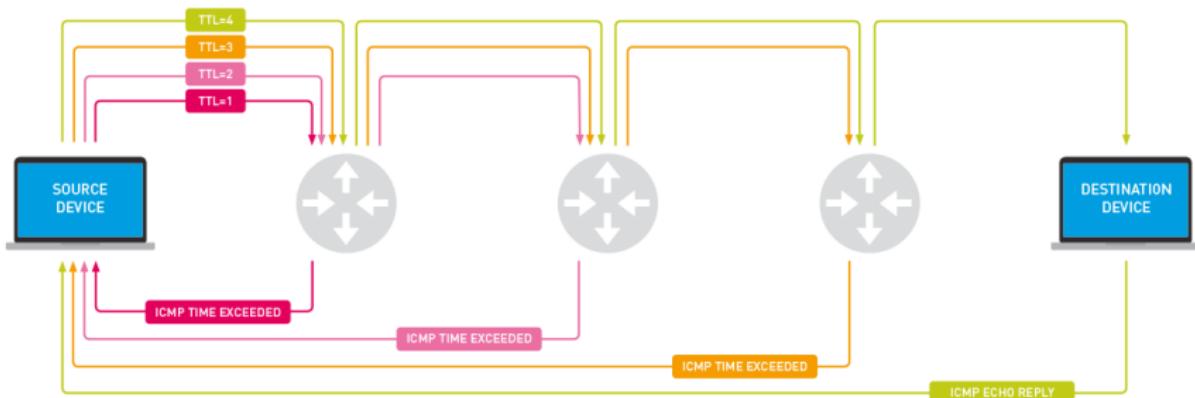


1. Routers *decrement* the TTL by one every time they route a datagram.

# ICMP and Traceroute

ICMP (Internet Control Message Protocol) is used to *test connectivity*.

About a dozen types of ICMP messages are defined, and each ICMP message type is carried encapsulated in an IP packet.

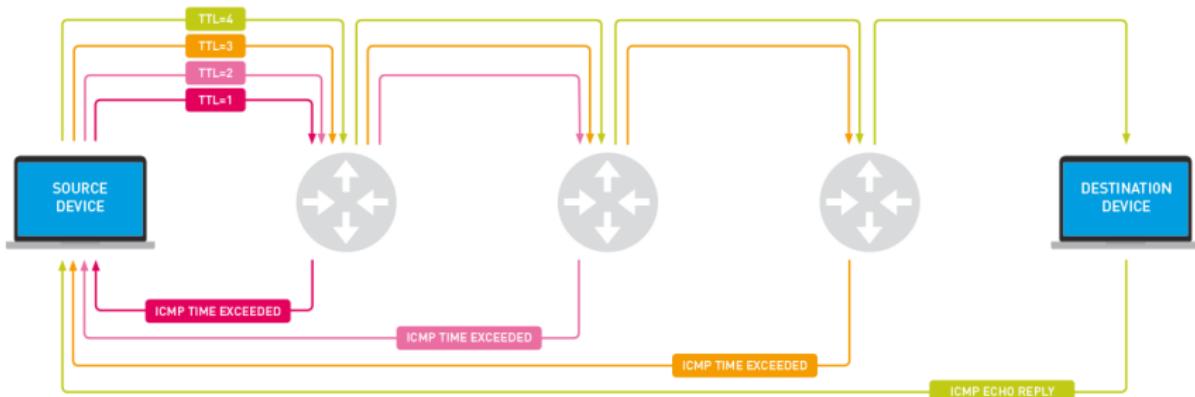


1. Routers *decrement* the TTL by one every time they route a datagram.
2. When TTL is 0, the datagram is discarded and the router sends a “*ICMP Time Exceeded*” message to the sender.

# ICMP and Traceroute

ICMP (Internet Control Message Protocol) is used to *test connectivity*.

About a dozen types of ICMP messages are defined, and each ICMP message type is carried encapsulated in an IP packet.



1. Routers *decrement* the TTL by one every time they route a datagram.
2. When TTL is 0, the datagram is discarded and the router sends a “*ICMP Time Exceeded*” message to the sender.
3. A host can enumerate the routers on a path by sending datagrams with increasing TTL ( $\rightarrow$  *Traceroute*).

# Ping of Death (1997)

1. The attacker wants to send a datagram larger than 65KB, but this would violate the TCP/IP rules;



90,000 bytes

**90KB >> 65KB**

---

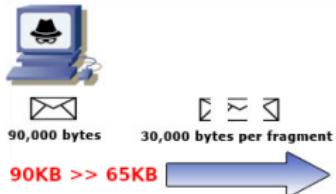
<sup>2</sup> [https://en.wikipedia.org/wiki/Ping\\_of\\_death](https://en.wikipedia.org/wiki/Ping_of_death)

<sup>3</sup> <https://www.cve.org/CVERecord?id=CVE-2013-3183>

<sup>4</sup> <https://www.cve.org/CVERecord?id=CVE-2020-16898>

# Ping of Death (1997)

1. The attacker wants to send a datagram larger than 65KB, but this would violate the TCP/IP rules;
2. The attacker splits the datagram into *multiple fragments*;

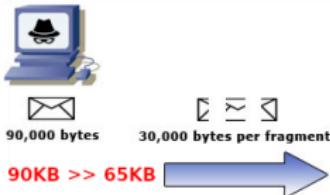


<sup>2</sup>[https://en.wikipedia.org/wiki/Ping\\_of\\_death](https://en.wikipedia.org/wiki/Ping_of_death)

<sup>3</sup><https://www.cve.org/CVERecord?id=CVE-2013-3183>

<sup>4</sup><https://www.cve.org/CVERecord?id=CVE-2020-16898>

# Ping of Death (1997)



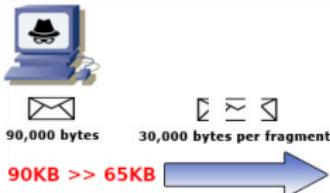
1. The attacker wants to send a datagram larger than 65KB, but this would violate the TCP/IP rules;
2. The attacker splits the datagram into *multiple fragments*;
3. The total length is not known a priori and the receiver needs to process all the fragments before it can reassemble them. Once reassembled at the destination, the datagram will exceed the maximum allowed size, triggering a buffer overflow!

<sup>2</sup>[https://en.wikipedia.org/wiki/Ping\\_of\\_death](https://en.wikipedia.org/wiki/Ping_of_death)

<sup>3</sup><https://www.cve.org/CVERecord?id=CVE-2013-3183>

<sup>4</sup><https://www.cve.org/CVERecord?id=CVE-2020-16898>

# Ping of Death (1997)



1. The attacker wants to send a datagram larger than 65KB, but this would violate the TCP/IP rules;
2. The attacker splits the datagram into *multiple fragments*;
3. The total length is not known a priori and the receiver needs to process all the fragments before it can reassemble them. Once reassembled at the destination, the datagram will exceed the maximum allowed size, triggering a buffer overflow!

A few decades ago, this vulnerability affected the majority of operating systems, but nowadays is patched (*almost<sup>2,3,4</sup>*) everywhere...

**Good example of how sometimes standards ≠ implementation**

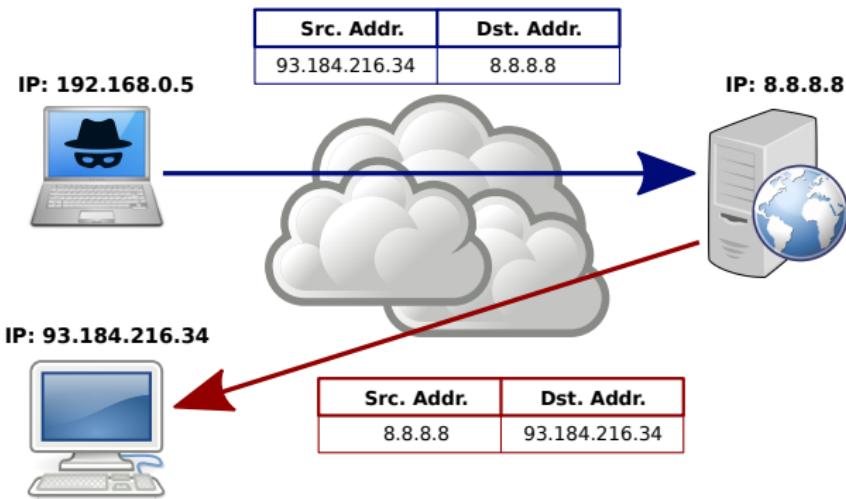
<sup>2</sup> [https://en.wikipedia.org/wiki/Ping\\_of\\_death](https://en.wikipedia.org/wiki/Ping_of_death)

<sup>3</sup> <https://www.cve.org/CVERecord?id=CVE-2013-3183>

<sup>4</sup> <https://www.cve.org/CVERecord?id=CVE-2020-16898>

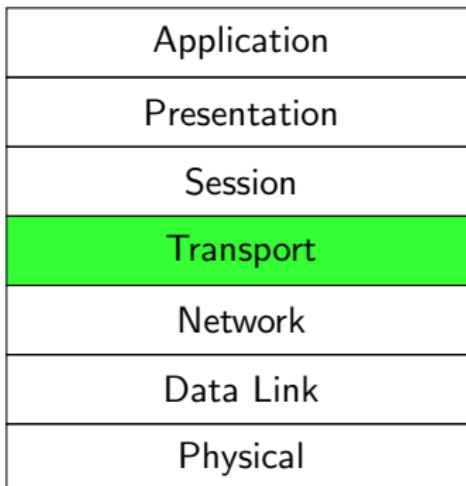
# IP Spoofing

1. The attacker sends a message from a *different IP* from the one that he controls.
2. The recipient of the message will reply to the *spoofed IP* address.



- In the local network: IP spoofing == ARP spoofing.
- IP spoofing is commonly used in **Denial of Service (DoS) attacks** → we will see more about this in the next slides...

# The Transport Layer



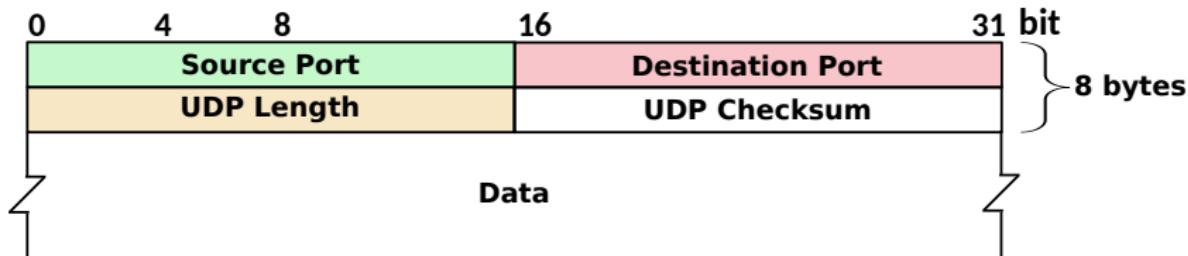
# UDP Datagram

Is the simplest transport protocol that provides port abstraction to programs.

**Best effort:** messages are not acknowledged, reception is not guaranteed.

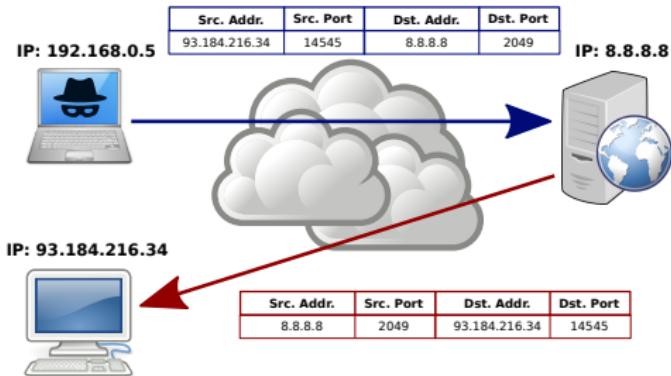
Popular protocols relying on UDP:

- Domain Name System (port 53);
- Dynamic Host Configuration Protocol (ports 67, 68);
- Network File System (port 2049).



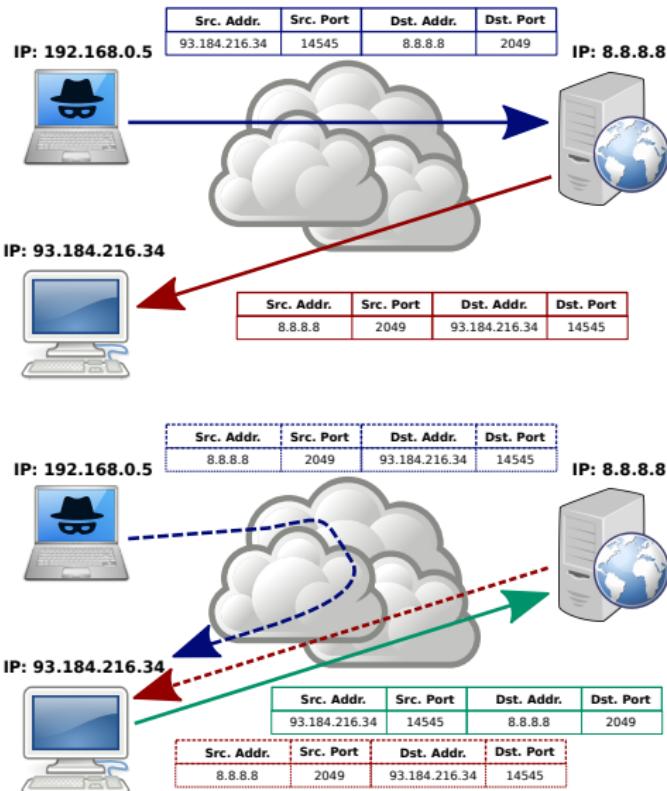
- **Source/Destination Port:** 65,535 possible options.
  - The service(s) listen on a *given port(s)*;
  - The client(s) create connections from *random ports*.
- **UDP Length:** total number of bytes (header + payload).  
The minimum length of an UDP packet is 8 bytes.

# UDP Spoofing/Hijacking



**UDP spoofing:** identical approach as for IP spoofing.

# UDP Spoofing/Hijacking



**UDP spoofing:** identical approach as for IP spoofing.

**UDP hijacking:** the attacker *competes* with the server to send to the victim a reply before the the packet from server is received.

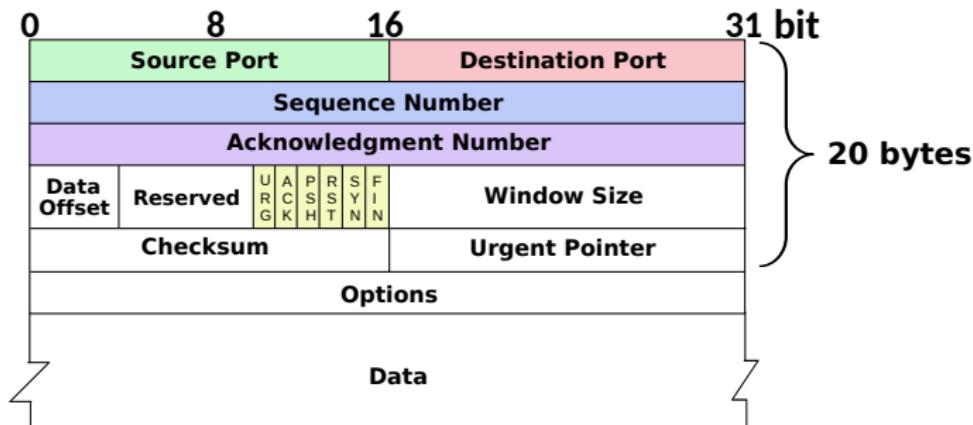
The attacker sends to the victim the same header as the server.

# TCP Segments

Identical port abstraction as UDP, but with a **connection-oriented and reliable** stream delivery service. It leverages *virtual circuits* which are composed of two streams and a full-duplex connection.

Popular protocols relying on TCP:

- Secure Shell (22);
- Hyper Text Transfer Protocol (port 80 and 443).

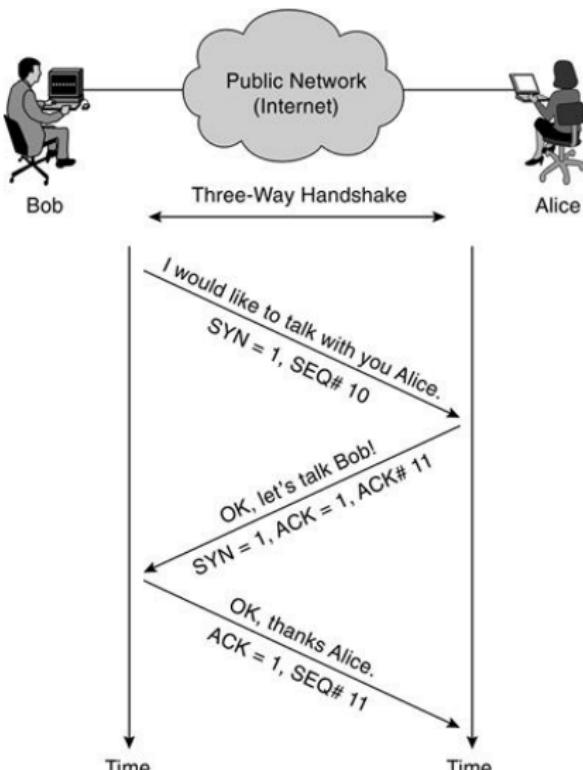


- **Source/Destination Port**: service that.
- **(Initial) Sequence Number (ISN)**: segment position in the stream.
- **Acknowledgment Number**: last segment received.
- **TCP flags**: used to setup and shutdown the virtual circuit.

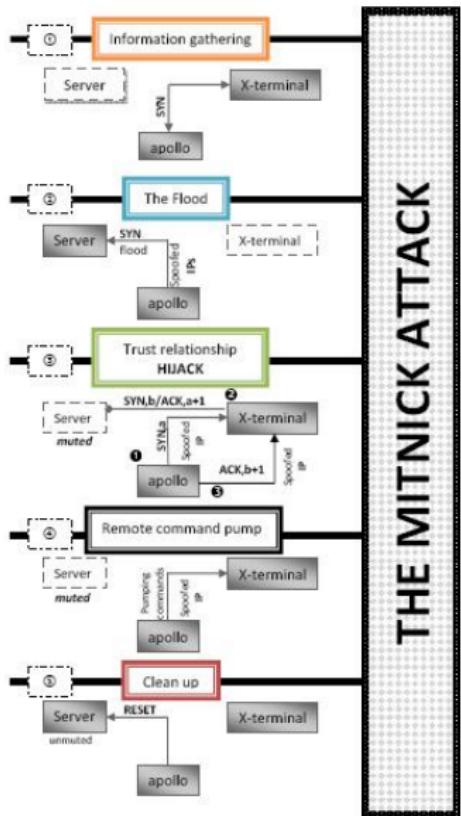
# TCP Flags and the Three-way Handshake

Flags usage:

1. **SYN**: requests for synchronization of sequence and acknowledgment (ACK) numbers.
2. **ACK**: acknowledges the receipt of a segment.
  - ACK = sequence number + 1
3. **FIN**: requests the shutdown of the circuit (*gracious*).
4. **RST**: requests to immediately shutdown the circuit (*abrupt*).
5. **PSH**: requests to pass the data stream to the user level (program) as soon as possible.
6. **URG**: specifies that the data is urgent.



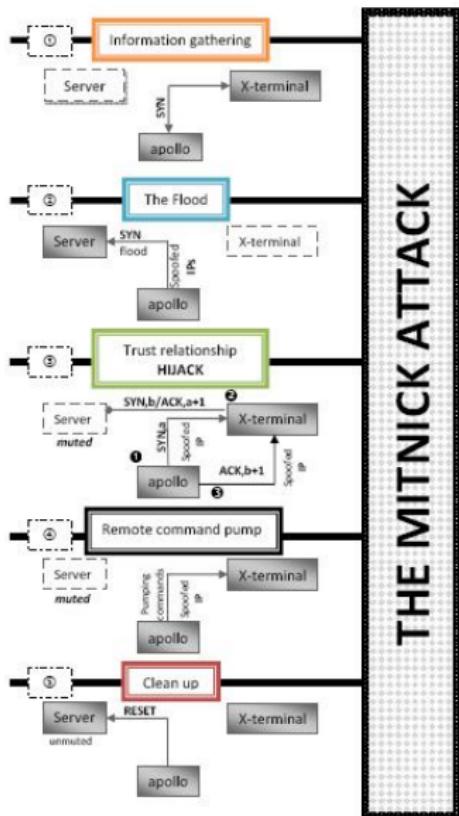
# Mitnick's TCP Session Hijack Attack (1994)



## 1. Information gathering:

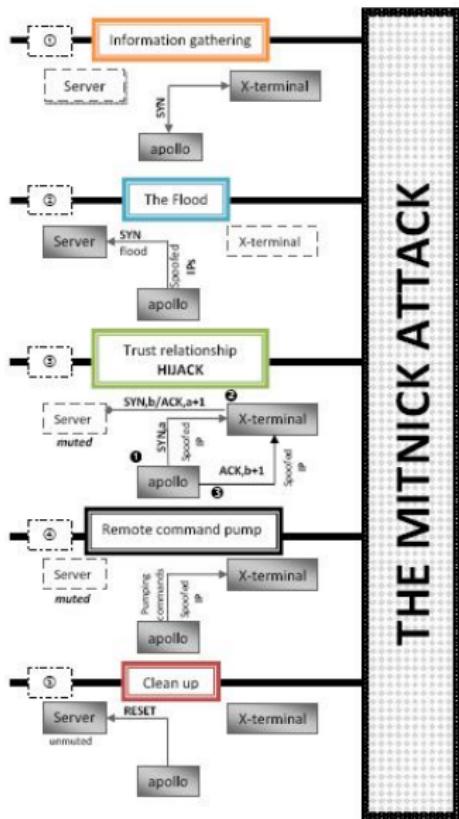
- determine trusted relationships (IP, TCP ISN);
- determine the TCP ISN

# Mitnick's TCP Session Hijack Attack (1994)



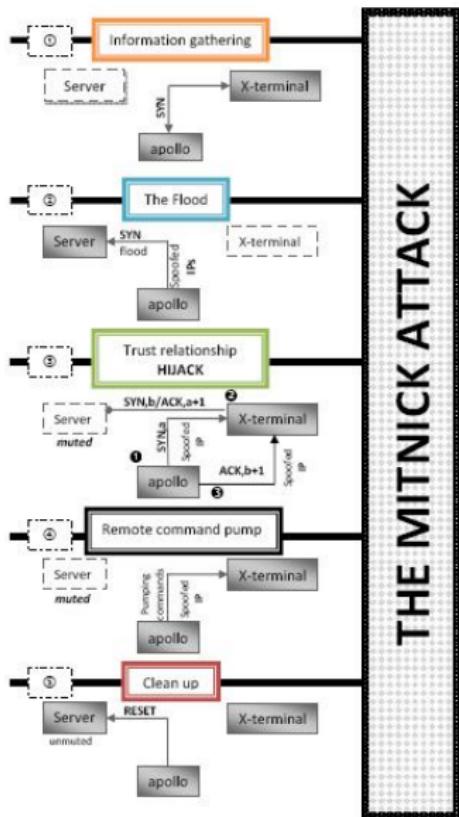
1. Information gathering:
  - determine trusted relationships (IP, TCP ISN);
  - determine the TCP ISN
2. Server flooding:
  - fill the memory of server the with half-open SYN requests from spoofed IPs.

# Mitnick's TCP Session Hijack Attack (1994)



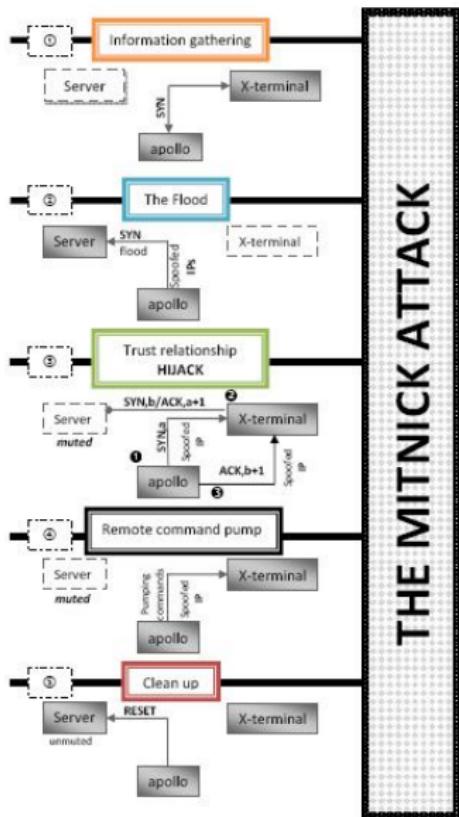
1. Information gathering:
  - determine trusted relationships (IP, TCP ISN);
  - determine the TCP ISN
2. Server flooding:
  - fill the memory of server the with half-open SYN requests from spoofed IPs.
3. Session hijacking:

# Mitnick's TCP Session Hijack Attack (1994)



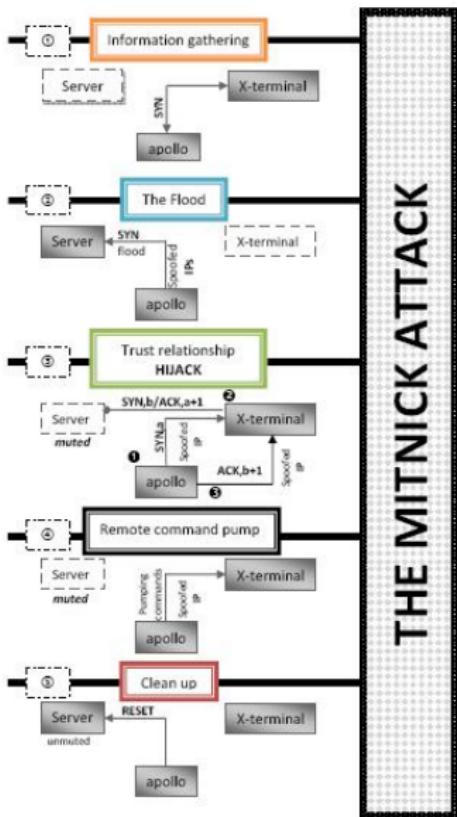
1. Information gathering:
  - determine trusted relationships (IP, TCP ISN);
  - determine the TCP ISN
2. Server flooding:
  - fill the memory of server the with half-open SYN requests from spoofed IPs.
3. Session hijacking:
  - send SYN to the client with the spoofed address of the server;

# Mitnick's TCP Session Hijack Attack (1994)



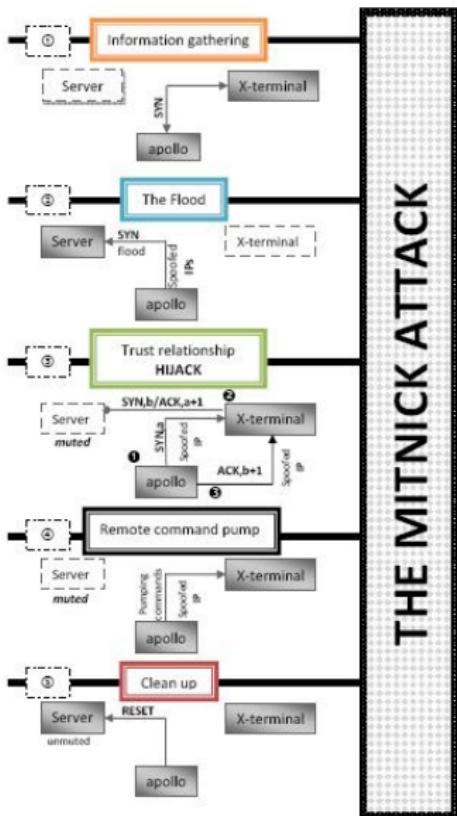
1. Information gathering:
  - determine trusted relationships (IP, TCP ISN);
  - determine the TCP ISN
2. Server flooding:
  - fill the memory of server the with half-open SYN requests from spoofed IPs.
3. Session hijacking:
  - send SYN to the client with the spoofed address of the server;
  - client sends SYN/ACK to the server (which will not reply);

# Mitnick's TCP Session Hijack Attack (1994)



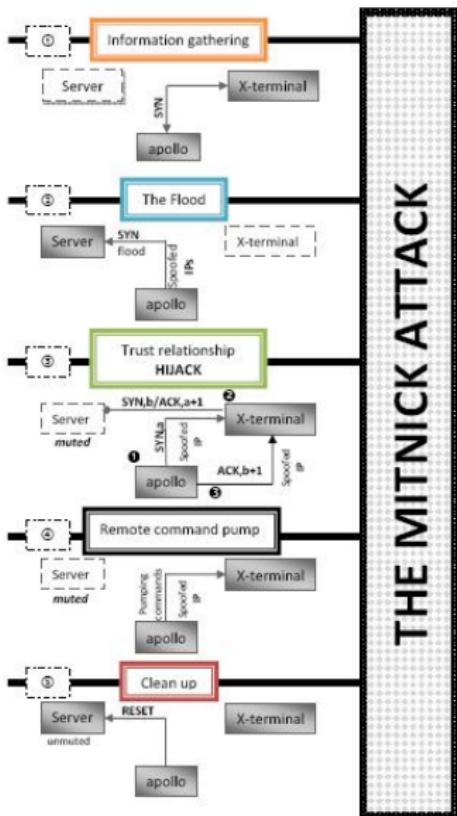
1. Information gathering:
  - determine trusted relationships (IP, TCP ISN);
  - determine the TCP ISN
2. Server flooding:
  - fill the memory of server the with half-open SYN requests from spoofed IPs.
3. Session hijacking:
  - send SYN to the client with the spoofed address of the server;
  - client sends SYN/ACK to the server (which will not reply);
  - send spoofed ACK to the client and complete the three-way handshake.

# Mitnick's TCP Session Hijack Attack (1994)



1. Information gathering:
  - determine trusted relationships (IP, TCP ISN);
  - determine the TCP ISN
2. Server flooding:
  - fill the memory of server the with half-open SYN requests from spoofed IPs.
3. Session hijacking:
  - send SYN to the client with the spoofed address of the server;
  - client sends SYN/ACK to the server (which will not reply);
  - send spoofed ACK to the client and complete the three-way handshake.
4. Create the backdoor.

# Mitnick's TCP Session Hijack Attack (1994)



1. Information gathering:
  - determine trusted relationships (IP, TCP ISN);
  - determine the TCP ISN
2. Server flooding:
  - fill the memory of server the with half-open SYN requests from spoofed IPs.
3. Session hijacking:
  - send SYN to the client with the spoofed address of the server;
  - client sends SYN/ACK to the server (which will not reply);
  - send spoofed ACK to the client and complete the three-way handshake.
4. Create the backdoor.
5. Clean up:
  - send RST to server to cancel all the SYN requests.

# Zalewski's analysis of ISNs (2001, 2002)

The security of the protocol relies on the *difficulty* to guess the TCP sequence number. Any attacker that can generate (*and send!*)  $2^{32}$  packets would be able to guess the sequence number.

Unfortunately, many OS had a flow in the way they used to generate sequence numbers<sup>5,6</sup>

Zalewski's assumptions:

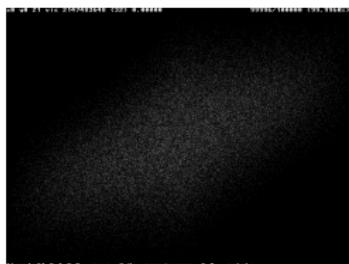
- the attacker collected around 50,000 SN values;
- the attacker will perform at most 5,000 attempts to guess the SN  
⇒ this is achievable in a few seconds.



Test date: April 2001  
OS: Windows 98 SE  
Attack feasibility: 100%



Test date: April 2001  
OS: MacOS 9  
Attack feasibility: 89%



Test date: October 2002  
OS: MacOS X  
Attack feasibility: 0%

<sup>5</sup> <https://lcamtuf.coredump.cx/oldtcp/tcpseq.html>

<sup>6</sup> <https://lcamtuf.coredump.cx/newtcp/>

# TCP Hijacking in 2016

In the original TCP Linux implementation, TCP segments that fall within a certain range can be accepted - this opens the field to blind RST attacks and data injection attacks.

*RFC 5961* introduces a challenge ACKs, which are sent to the source IP upon receiving a SYN packet for resetting an ongoing a TCP connection. Challenge ACKs are rate limited, and monitored using a single global variable for all TCP connections.

An attacker can leverage this information to:

- **check if two arbitrary hosts are communicating** using a TCP connection (and discover the port numbers on each host);
- **infer the sequence number and the ACK number** on each host.

Leveraging the above knowledge, the attacker can tear down connections or inject arbitrary data. Attack success rate: 88-97%; time required to carry out the attack: 40-60 seconds.

More details about the attack in the paper:

<https://www.usenix.org/conference/usenixsecurity16/technical-sessions/presentation/cao>

# Port Scanning (UDP and TCP)

The attacker needs to know which services are active on a host.

The process of enumerating the listening services is called **portscan**.

*Sweeping*: when the attacker scans multiple hosts on a single port.

# Port Scanning (UDP and TCP)

The attacker needs to know which services are active on a host.

The process of enumerating the listening services is called **portscan**.

*Sweeping*: when the attacker scans multiple hosts on a single port.

UDP port scan: send an empty UDP

datagram; the response is:

- **any response**  
⇒ port is open
- **no response**  
⇒ port is open *or filtered* (e.g.  
*firewall*)
- **ICMP port unreachable error**  
⇒ port is closed

# Port Scanning (UDP and TCP)

The attacker needs to know which services are active on a host.

The process of enumerating the listening services is called **portscan**.

*Sweeping*: when the attacker scans multiple hosts on a single port.

UDP port scan: send an empty UDP datagram; the response is:

- **any response**  
⇒ port is open
- **no response**  
⇒ port is open *or filtered* (e.g. *firewall*)
- **ICMP port unreachable error**  
⇒ port is closed

TCP SYN scan: send a segment with the SYN flag; the response is:

- **SYN + ACK**  
⇒ port is open
- **RST**  
⇒ the port is closed
- **no response**  
⇒ port is filtered

# Port Scanning (UDP and TCP)

The attacker needs to know which services are active on a host.

The process of enumerating the listening services is called **portscan**.

*Sweeping*: when the attacker scans multiple hosts on a single port.

UDP port scan: send an empty UDP datagram; the response is:

- **any response**  
⇒ port is open
- **no response**  
⇒ port is open *or filtered* (e.g. *firewall*)
- **ICMP port unreachable error**  
⇒ port is closed

TCP SYN scan: send a segment with the SYN flag; the response is:

- **SYN + ACK**  
⇒ port is open
- **RST**  
⇒ the port is closed
- **no response**  
⇒ port is filtered

Popular open-source tools:



# Port Scanning (UDP and TCP)

The  
The  
Swe

UDP  
data

- a
- r
- f
- l

Pop

In case you thought that this is just “*hackers business*”

British govt is scanning all Internet devices hosted in UK

By Sergiu Gatlan

November 4, 2022



03:22 PM



3



Sources:

<https://www.bleepingcomputer.com/news/security/british-govt-is-scanning-all-internet-devices-hosted-in-uk/>

<https://www.ncsc.gov.uk/information/ncsc-scanning-information>

<https://www.ncsc.gov.uk/blog-post/scanning-the-internet-for-fun-and-profit>

# Port Scanning (TCP) - Alternative Techniques

## From RFC 793:

- port is closed && segment without RST flag arrives  $\Rightarrow$  reply with RST
- port is open && a segment without SYN/RST/ACK flag arrives  $\Rightarrow$  drop the segment



NULL, FIN, PSH, URG scans response contains:

- **RST**: closed
- **No response**: open *or filtered* (e.g., firewall)
- **ICMP unreachable error**: filtered

## References:

<https://www.rfc-editor.org/rfc/rfc793> (page 65) and  
<https://nmap.org/book/scan-methods-ack-scan.html>

# Port Scanning (TCP) - Alternative Techniques

From RFC 793:

- port is closed && segment without RST flag arrives ⇒ reply with RST
- port is open && a segment without SYN/RST/ACK flag arrives ⇒ drop the segment



NULL, FIN, PSH, URG scans response contains:

- **RST**: closed
- **No response**: open *or filtered* (e.g., firewall)
- **ICMP unreachable error**: filtered



ACK scan response contains:

- **RST**: not filtered
- **No response**: filtered
- **ICMP unreachable error**: filtered

⇒ is used for testing firewall rules

References:

<https://www.rfc-editor.org/rfc/rfc793> (page 65) and  
<https://nmap.org/book/scan-methods-ack-scan.html>

# Port Scanning (TCP) - Alternative Techniques

From RFC 793:

- port is closed && segment without RST flag arrives ⇒ reply with RST
- port is open && a segment without SYN/RST/ACK flag arrives ⇒ drop the segment



NULL, FIN, PSH, URG scans response contains:

- **RST**: closed
- **No response**: open *or filtered* (e.g., firewall)
- **ICMP unreachable error**: filtered



ACK scan response contains:

- **RST**: not filtered
- **No response**: filtered
- **ICMP unreachable error**: filtered

⇒ is used for testing firewall rules



connect() scan:

- **Invasive** ⇒ **Why do we need it?**

References:

<https://www.rfc-editor.org/rfc/rfc793> (page 65) and  
<https://nmap.org/book/scan-methods-ack-scan.html>

# Port Scanning (TCP) - Alternative Techniques

From RFC 793:

- port is closed && segment without RST flag arrives ⇒ reply with RST
- port is open && a segment without SYN/RST/ACK flag arrives ⇒ drop the segment



NULL, FIN, PSH, URG scans response contains:

- **RST**: closed
- **No response**: open *or filtered* (e.g., firewall)
- **ICMP unreachable error**: filtered



ACK scan response contains:

- **RST**: not filtered
- **No response**: filtered
- **ICMP unreachable error**: filtered

⇒ is used for testing firewall rules



connect() scan:

- **Invasive** ⇒ **Why do we need it?**
- **Allows to grab banners**

References:

<https://www.rfc-editor.org/rfc/rfc793> (page 65) and  
<https://nmap.org/book/scan-methods-ack-scan.html>

# OS Fingerprinting

## Why?

- Can detect old, outdated or unpatched services.
- Allows to focus on particular vulnerabilities and develop exploits.
- Provides information on the hosts/elements of a network.
- Can identify “suspicious” machines which are not expected to be in the network.

## How?

Operating systems implement stacks in *slightly different* ways.

If we send **specially crafted packets** we can infer the OS that is running on a host.

### Examples:

1. IP TTL: Linux sets the initial TTL to 64, Windows uses 128.
2. Windows and Cisco IOS do not reply with a RST to FIN packets (as in RFC 793).
3. Infer OS from the pattern used in the TCP sequence number.

Two ways for performing OS fingerprinting:

1. **Active:** send specific packets (e.g., nmap).
2. **Passive:** passively observe traffic and infer the OS (e.g., p0f).

References: <https://nmap.org/book/osdetect.html>

# Denial of Service (DoS)

## What is DoS?

A type of **cyber attack** in which a malicious actor aims to render a computer or other device unavailable to its intended users by **interrupting the normal functioning of the device**.



# Denial of Service (DoS)

## What is DoS?

A type of **cyber attack** in which a malicious actor aims to render a computer or other device unavailable to its intended users by **interrupting the normal functioning of the device**.



How?

# Denial of Service (DoS)

## What is DoS?

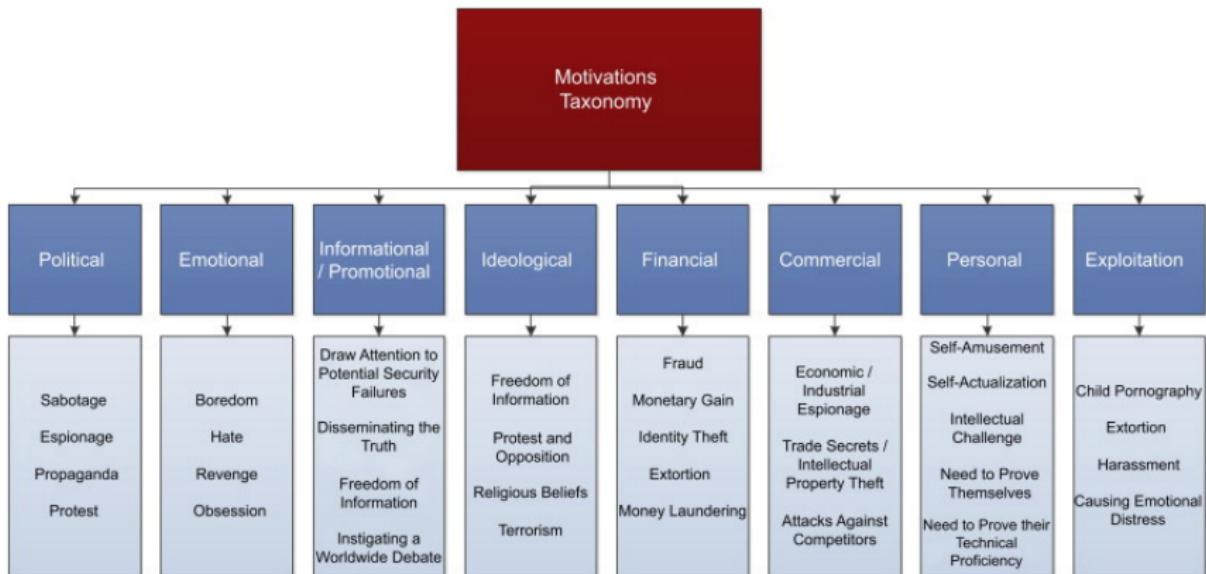
A type of **cyber attack** in which a malicious actor aims to render a computer or other device unavailable to its intended users by **interrupting the normal functioning of the device.**



The attacker sends huge volume of traffic over the network to **to starve victim's resources** in terms of:

- Network bandwidth ⇒ flooding.
- Network routing ⇒ sinkhole, router hijacking.
- CPU processing.
- Memory.
- Protocol state ⇒ SYN flooding (e.g., Mitnick's attack)

# DoS - Motivations



Source: <https://www.perimeter81.com/blog/network/the-psychology-behind-ddos-attacks>

# DoS - Attack Properties

- Type of resource being starved.
- Type of traffic sent:
  - UDP, TCP, SYN, ...



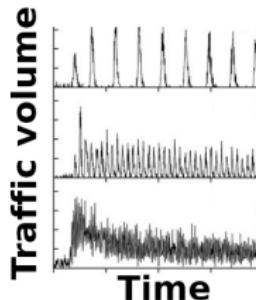
# DoS - Attack Properties

- Type of resource being starved.
- Type of traffic sent:
  - UDP, TCP, SYN, ...
- Spoofing used.
- Source(s):
  - **Centralized** → single or few attackers.
  - **Distributed (DDoS)** → many attackers (e.g., botnet).
- Third-party involvement.
  - Reflector amplification.

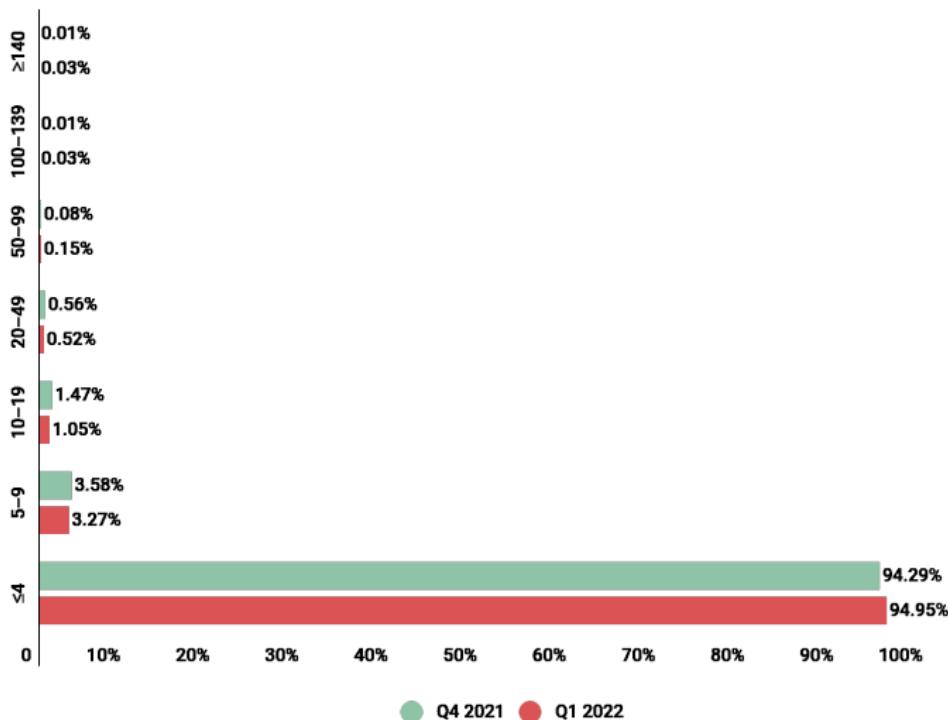


# DoS - Attack Properties

- Type of resource being starved.
- Type of traffic sent:
  - UDP, TCP, SYN, ...
- Spoofing used.
- Source(s):
  - **Centralized** → single or few attackers.
  - **Distributed (DDoS)** → many attackers (e.g., botnet).
- Third-party involvement.
  - Reflector amplification.
- Duration:
  - Constant, periodic or temporary.



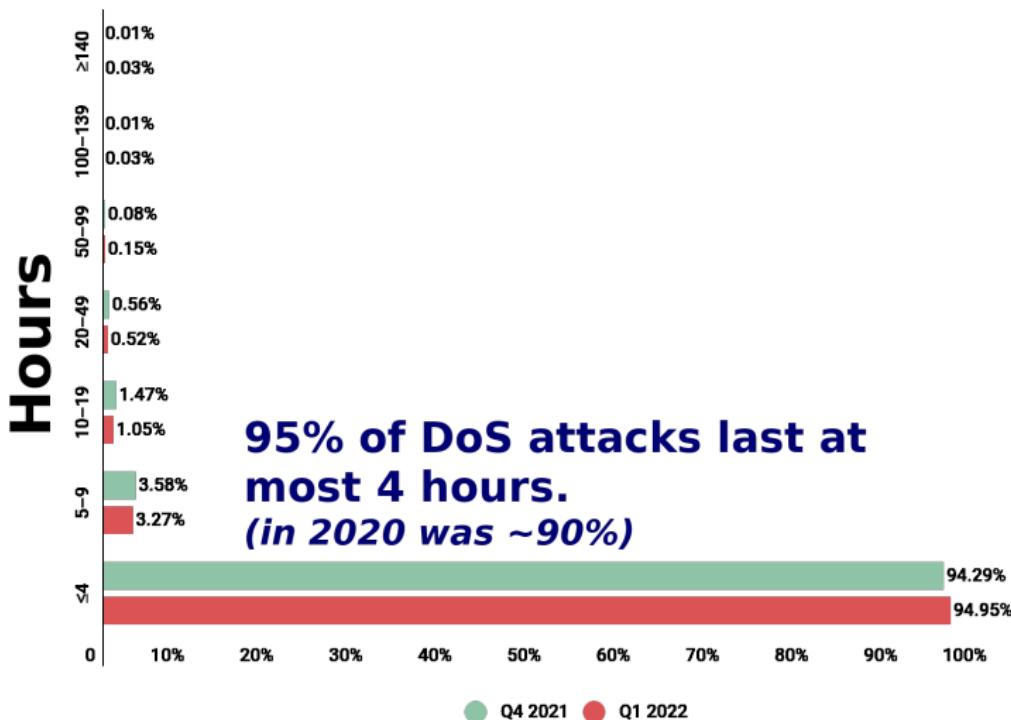
# DoS - Duration



Sources:

<https://securelist.com/ddos-attacks-in-q1-2022/106358/>

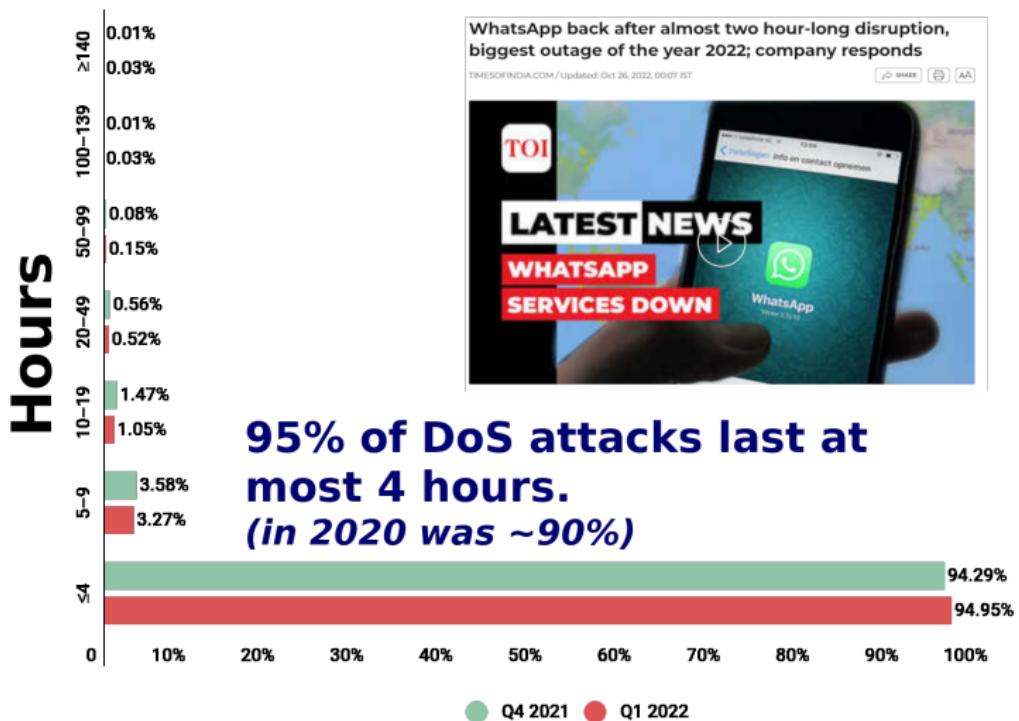
# DoS - Duration



Sources:

<https://securelist.com/ddos-attacks-in-q1-2022/106358/>

# DoS - Duration

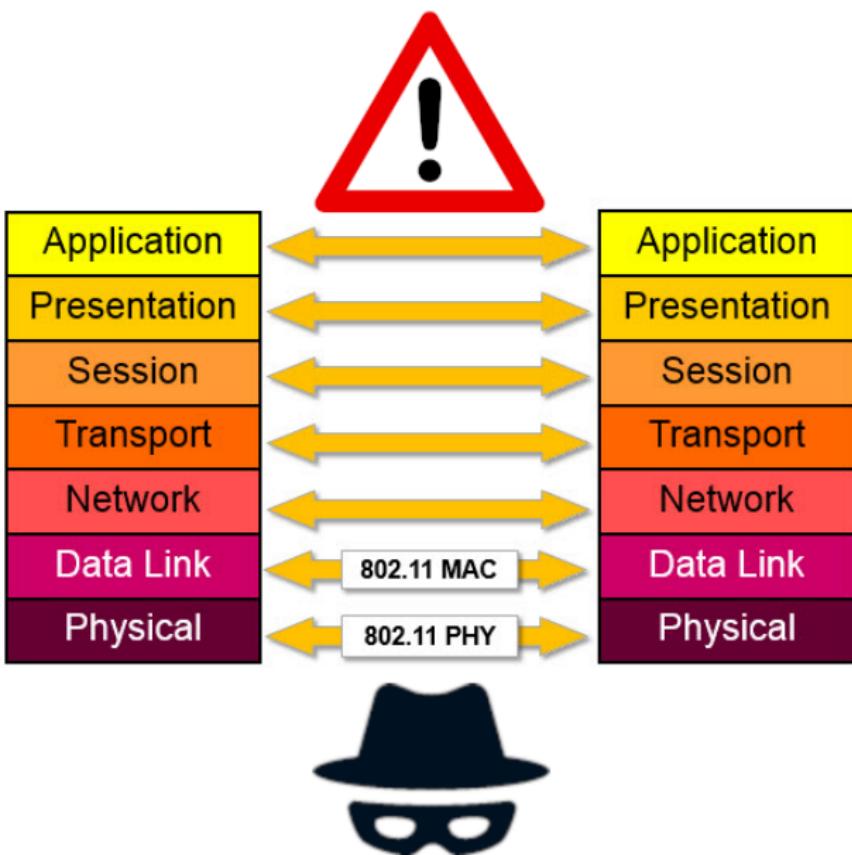


Sources:

<https://securelist.com/ddos-attacks-in-q1-2022/106358/>

<https://timesofindia.indiatimes.com/gadgets-news/whatsapp-down-for-some-users/articleshow/95075496.cms>

# DoS can Target any Layer



# DoS Attacks on 802.11

## 1. Physical layer jamming



# DoS Attacks on 802.11

## 1. Physical layer jamming

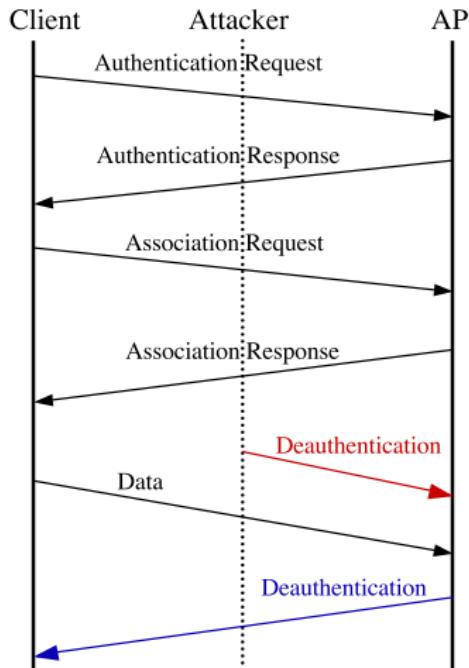
## 2. Link Layer

### □ Deauthentication

- The deauthentication frame is not authenticated
- one attacker packet every 6 victim packets

### □ Channel reservation

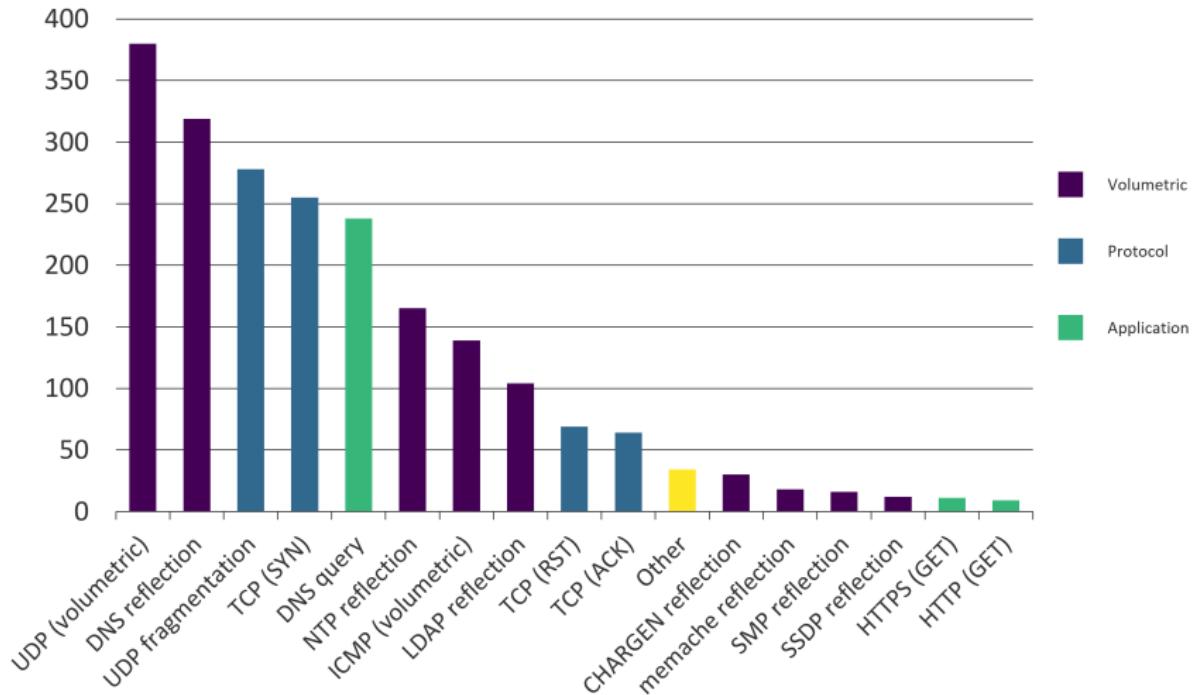
- Client can reserve channels for up to  $2^{15}$  seconds  
⇒ the attacker can block a channel



References:

[https://www.usenix.org/legacy/events/sec03/tech/full\\_papers/bellardo/bellardo.pdf](https://www.usenix.org/legacy/events/sec03/tech/full_papers/bellardo/bellardo.pdf)

# DoS Attack Types (2021)



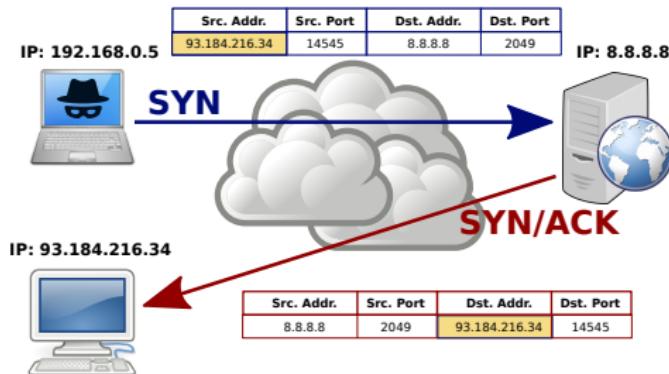
Source: <https://www.f5.com/labs/articles/threat-intelligence/2022-application-protection-report-ddos-attack-trends>

# Spoofing

## Issue:

Missing authentication for IP addresses  
⇒ the attacker can *fake the source address*.

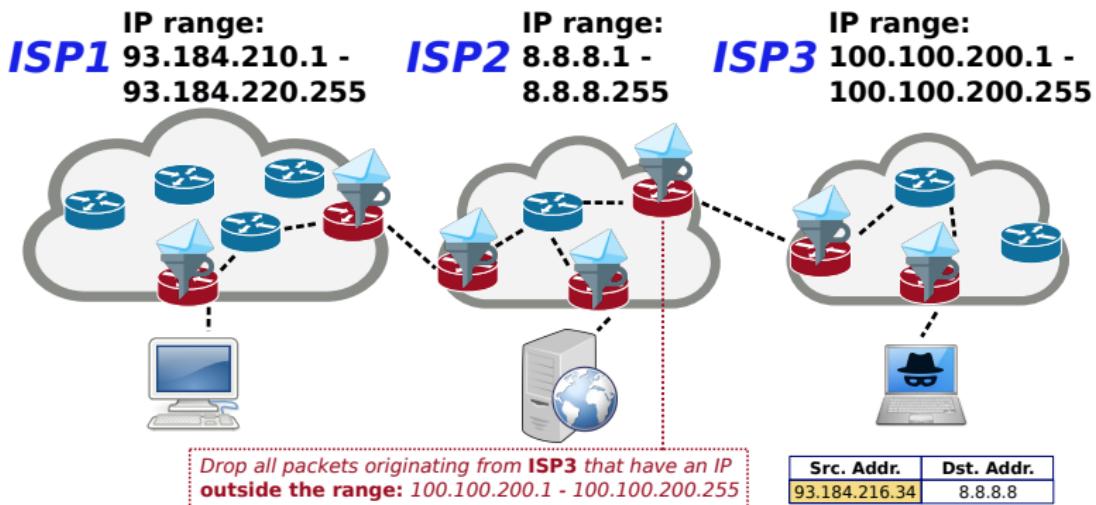
- “Shoot and forget” approach:
  - UDP packets
  - TCP packets without connection (e.g., SYN, ACK, RST)
- Tracing the origin of the attack is challenging!



# Ingress Filtering

Solution:

ISPs filter incoming packets with spoofed IP addresses RFC 2827 and 3704.



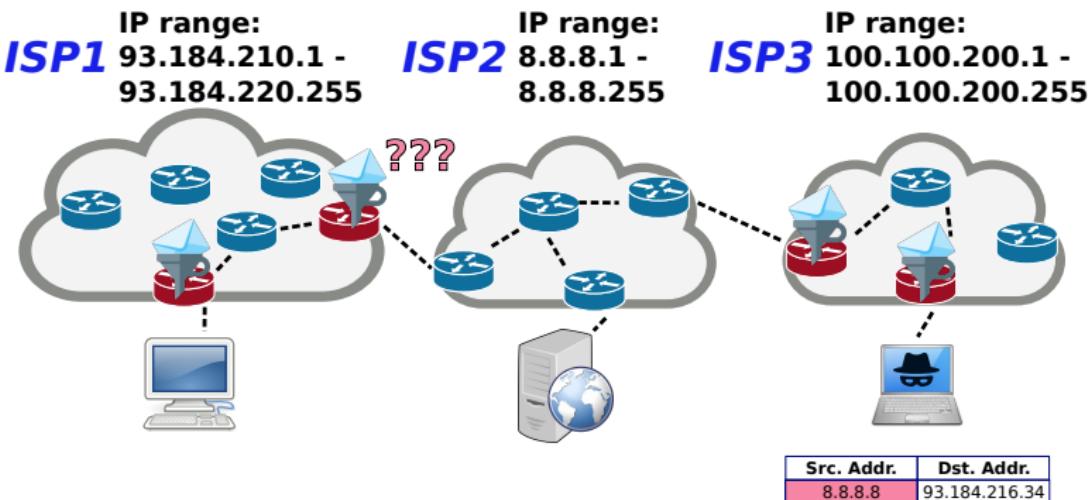
# Ingress Filtering

## Solution:

ISPs filter incoming packets with spoofed IP addresses RFC 2827 and 3704.

## Challenges:

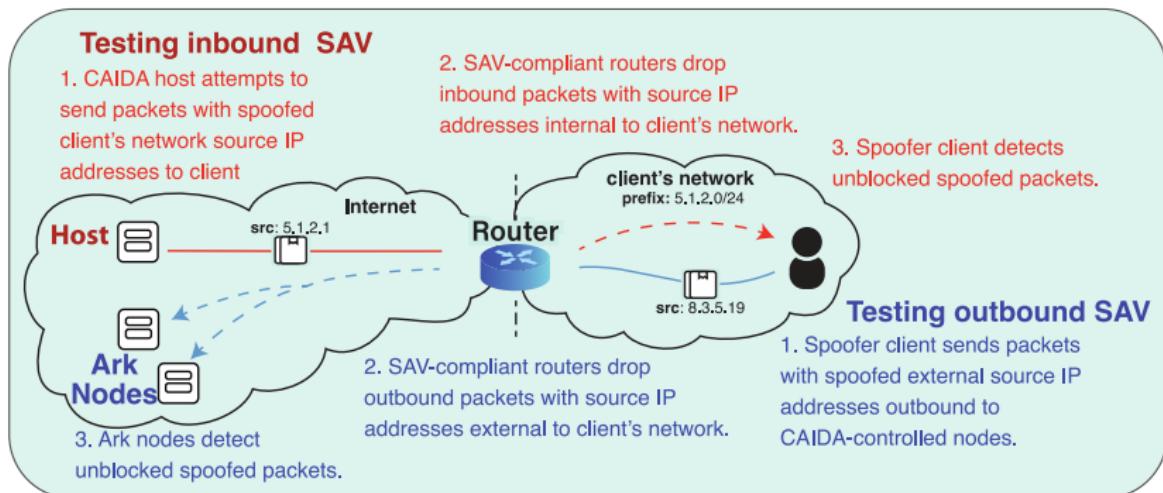
1. The majority/all ISPs must deploy filtering in order to work.
2. Benefits other networks → little incentive for deployment.



# Deployment of Filters

Several projects focused on measuring the deployment of ingress filters.

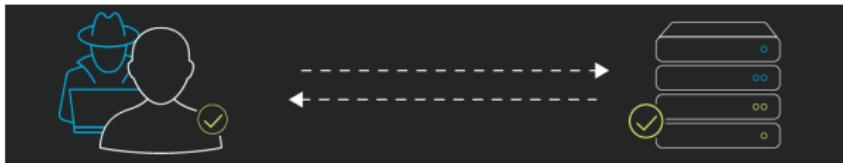
- Spoof project (started in 2005 at MIT → now CAIDA)
  1. Based on crowd-sourcing → volunteers run tests on their networks;
  2. spoofed packets from servers to clients (ingress filter);
  3. spoofed packets are sent from clients to servers (egress filter).



Source: <https://www.caida.org/projects/spoof/>

# Spoof Project

- Recent tests: [https://spoof.caida.org/recent\\_tests.php](https://spoof.caida.org/recent_tests.php)
- Overall statistics: <https://spoof.caida.org/summary.php>



- IPv4 results for 2022:
  - 18.5% of Autonomous Systems are spoofable (20.8% in 2022);
    - 8.3% only “partly” (11.5% in 2022);
  - 17.3% of IPv4 block are spoofable (18.9% in 2022)
- some inconsistencies among ASes when comparing IPv4 and IPv6:
  - 18.5% for IPv4; 25.3% for IPv6 blocks.
- NATs behave differently<sup>7</sup>:
  - Some may block spoofed traffic.
  - Some do not rewrite addresses outside the NAT's internal prefix.
  - Some rewrite (source address ↔ public IP) and pass spoofed packets.

<sup>7</sup> Spoof project on NATs: <https://www.caida.org/projects/spoofer/faq/>

## Backscatter:

The attacker randomly chooses a spoofed source address and the selected sources will receive answers from the victim.  $\Rightarrow$  can be leveraged to study DoS!

How: using a darknet aka *Network Telescope*.

- Monitors unused address space and receives:
  - scans;
  - DoS backscatter.
- If *source* (for **spoofing**) or *destination* (for **scanning**) are chosen randomly... the darknet will receive traffic proportional to its size.
  - CAIDA telescope ( $\approx$ 16M IPs)  $\Rightarrow$  2,000-3,000 attacks/week<sup>8</sup>.
- Greynet: mixes used and unused address space.
  - Arbor Networks' Active Threat Level Analysis System (ATLAS) is formed by many ISP greynets<sup>9</sup>.



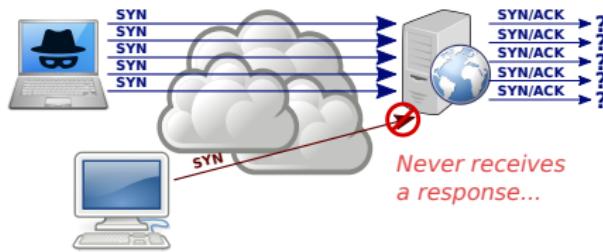
<sup>8</sup> [https://www.caida.org/catalog/papers/2006\\_backscatter\\_dos/backscatter\\_dos.pdf](https://www.caida.org/catalog/papers/2006_backscatter_dos/backscatter_dos.pdf)

<sup>9</sup> <https://www.dlt.com/resources/ddos-attack-protection-arbor-network-s-atlas>

# SYN Flooding

Issue: SYN received → store all the info (i.e.,  $state^{10}$ ) for processing segments.

Attacker's goal: fill victim's SYN queue so that it cannot accept other connections.



## Features:

- Protocol state attack:
  - targets TCP state control mechanism;
  - low rate ⇒ no need to flood the network.
- Source address can be spoofed.

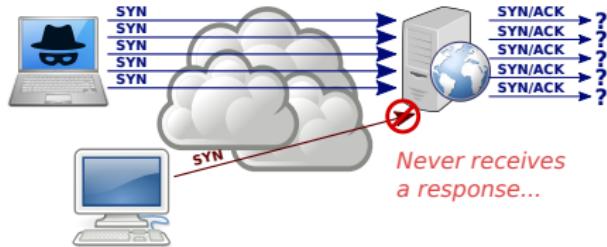
---

<sup>10</sup> RFC with all the details about the data structure used to keep the information for each connection:  
<https://datatracker.ietf.org/doc/html/rfc2140>

# SYN Flooding

Issue: SYN received → store all the info (i.e., *state*<sup>10</sup>) for processing segments.

Attacker's goal: fill victim's SYN queue so that it cannot accept other connections.



## Features:

- Protocol state attack:
  - targets TCP state control mechanism;
  - low rate ⇒ no need to flood the network.
- Source address can be spoofed.

## Defenses:

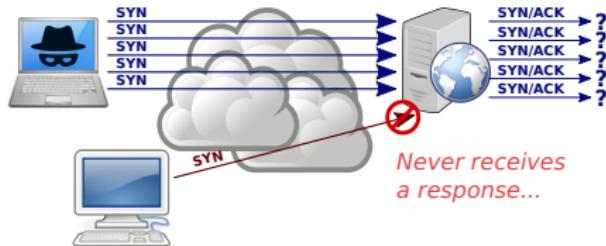
---

<sup>10</sup> RFC with all the details about the data structure used to keep the information for each connection:  
<https://datatracker.ietf.org/doc/html/rfc2140>

# SYN Flooding

Issue: SYN received → store all the info (i.e., *state*<sup>10</sup>) for processing segments.

Attacker's goal: fill victim's SYN queue so that it cannot accept other connections.



## Features:

- Protocol state attack:
  - targets TCP state control mechanism;
  - low rate ⇒ no need to flood the network.
- Source address can be spoofed.

## Defenses:

- Increase the size of the queue

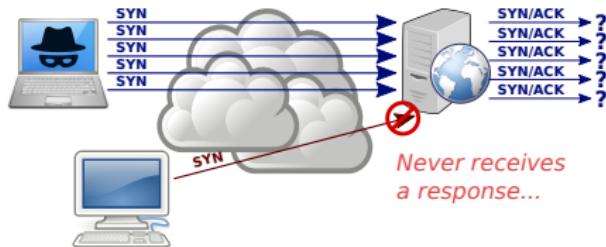
---

<sup>10</sup> RFC with all the details about the data structure used to keep the information for each connection:  
<https://datatracker.ietf.org/doc/html/rfc2140>

# SYN Flooding

Issue: SYN received → store all the info (i.e., *state*<sup>10</sup>) for processing segments.

Attacker's goal: fill victim's SYN queue so that it cannot accept other connections.



## Features:

- Protocol state attack:
  - targets TCP state control mechanism;
  - low rate ⇒ no need to flood the network.
- Source address can be spoofed.

## Defenses:

- Increase the size of the queue ⇒ the attacker sends more packets.

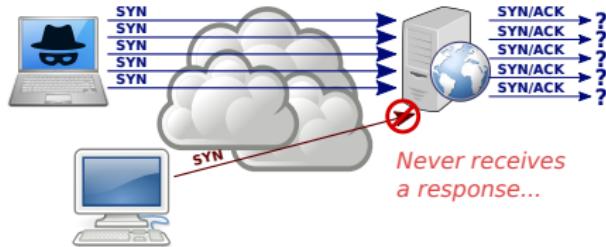
---

<sup>10</sup> RFC with all the details about the data structure used to keep the information for each connection:  
<https://datatracker.ietf.org/doc/html/rfc2140>

# SYN Flooding

Issue: SYN received → store all the info (i.e., *state*<sup>10</sup>) for processing segments.

Attacker's goal: fill victim's SYN queue so that it cannot accept other connections.



## Features:

- Protocol state attack:
  - targets TCP state control mechanism;
  - low rate ⇒ no need to flood the network.
- Source address can be spoofed.

## Defenses:

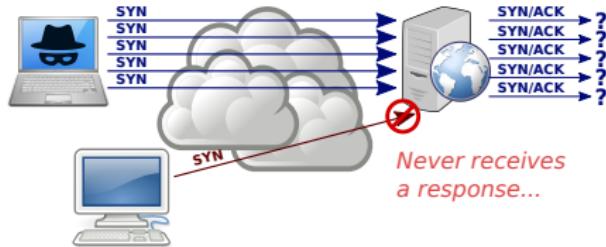
- Increase the size of the queue ⇒ the attacker sends more packets.
- Decrease timeout

<sup>10</sup> RFC with all the details about the data structure used to keep the information for each connection:  
<https://datatracker.ietf.org/doc/html/rfc2140>

# SYN Flooding

Issue: SYN received → store all the info (i.e., *state*<sup>10</sup>) for processing segments.

Attacker's goal: fill victim's SYN queue so that it cannot accept other connections.



## Features:

- Protocol state attack:
  - targets TCP state control mechanism;
  - low rate ⇒ no need to flood the network.
- Source address can be spoofed.

## Defenses:

- Increase the size of the queue ⇒ the attacker sends more packets.
- Decrease timeout ⇒ the attacker increases the rate.

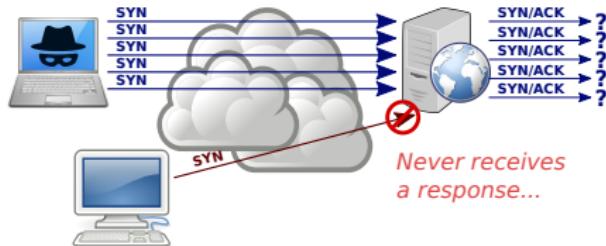
---

<sup>10</sup> RFC with all the details about the data structure used to keep the information for each connection:  
<https://datatracker.ietf.org/doc/html/rfc2140>

# SYN Flooding

Issue: SYN received → store all the info (i.e., *state*<sup>10</sup>) for processing segments.

Attacker's goal: fill victim's SYN queue so that it cannot accept other connections.



## Features:

- Protocol state attack:
  - targets TCP state control mechanism;
  - low rate ⇒ no need to flood the network.
- Source address can be spoofed.

## Defenses: Poor defenses:

- Increase the size of the queue ⇒ the attacker sends more packets.
- Decrease timeout ⇒ the attacker increases the rate.

<sup>10</sup> RFC with all the details about the data structure used to keep the information for each connection:  
<https://datatracker.ietf.org/doc/html/rfc2140>