



# Software Exploitation

Computer Security Fall 2024

Georgios (George) Portokalidis

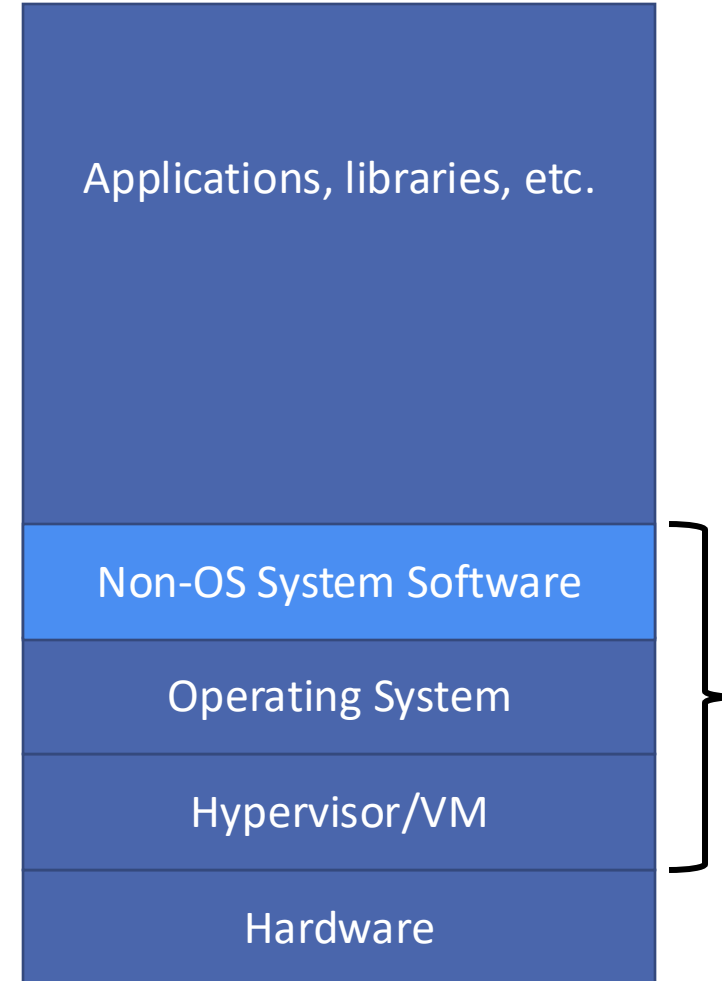


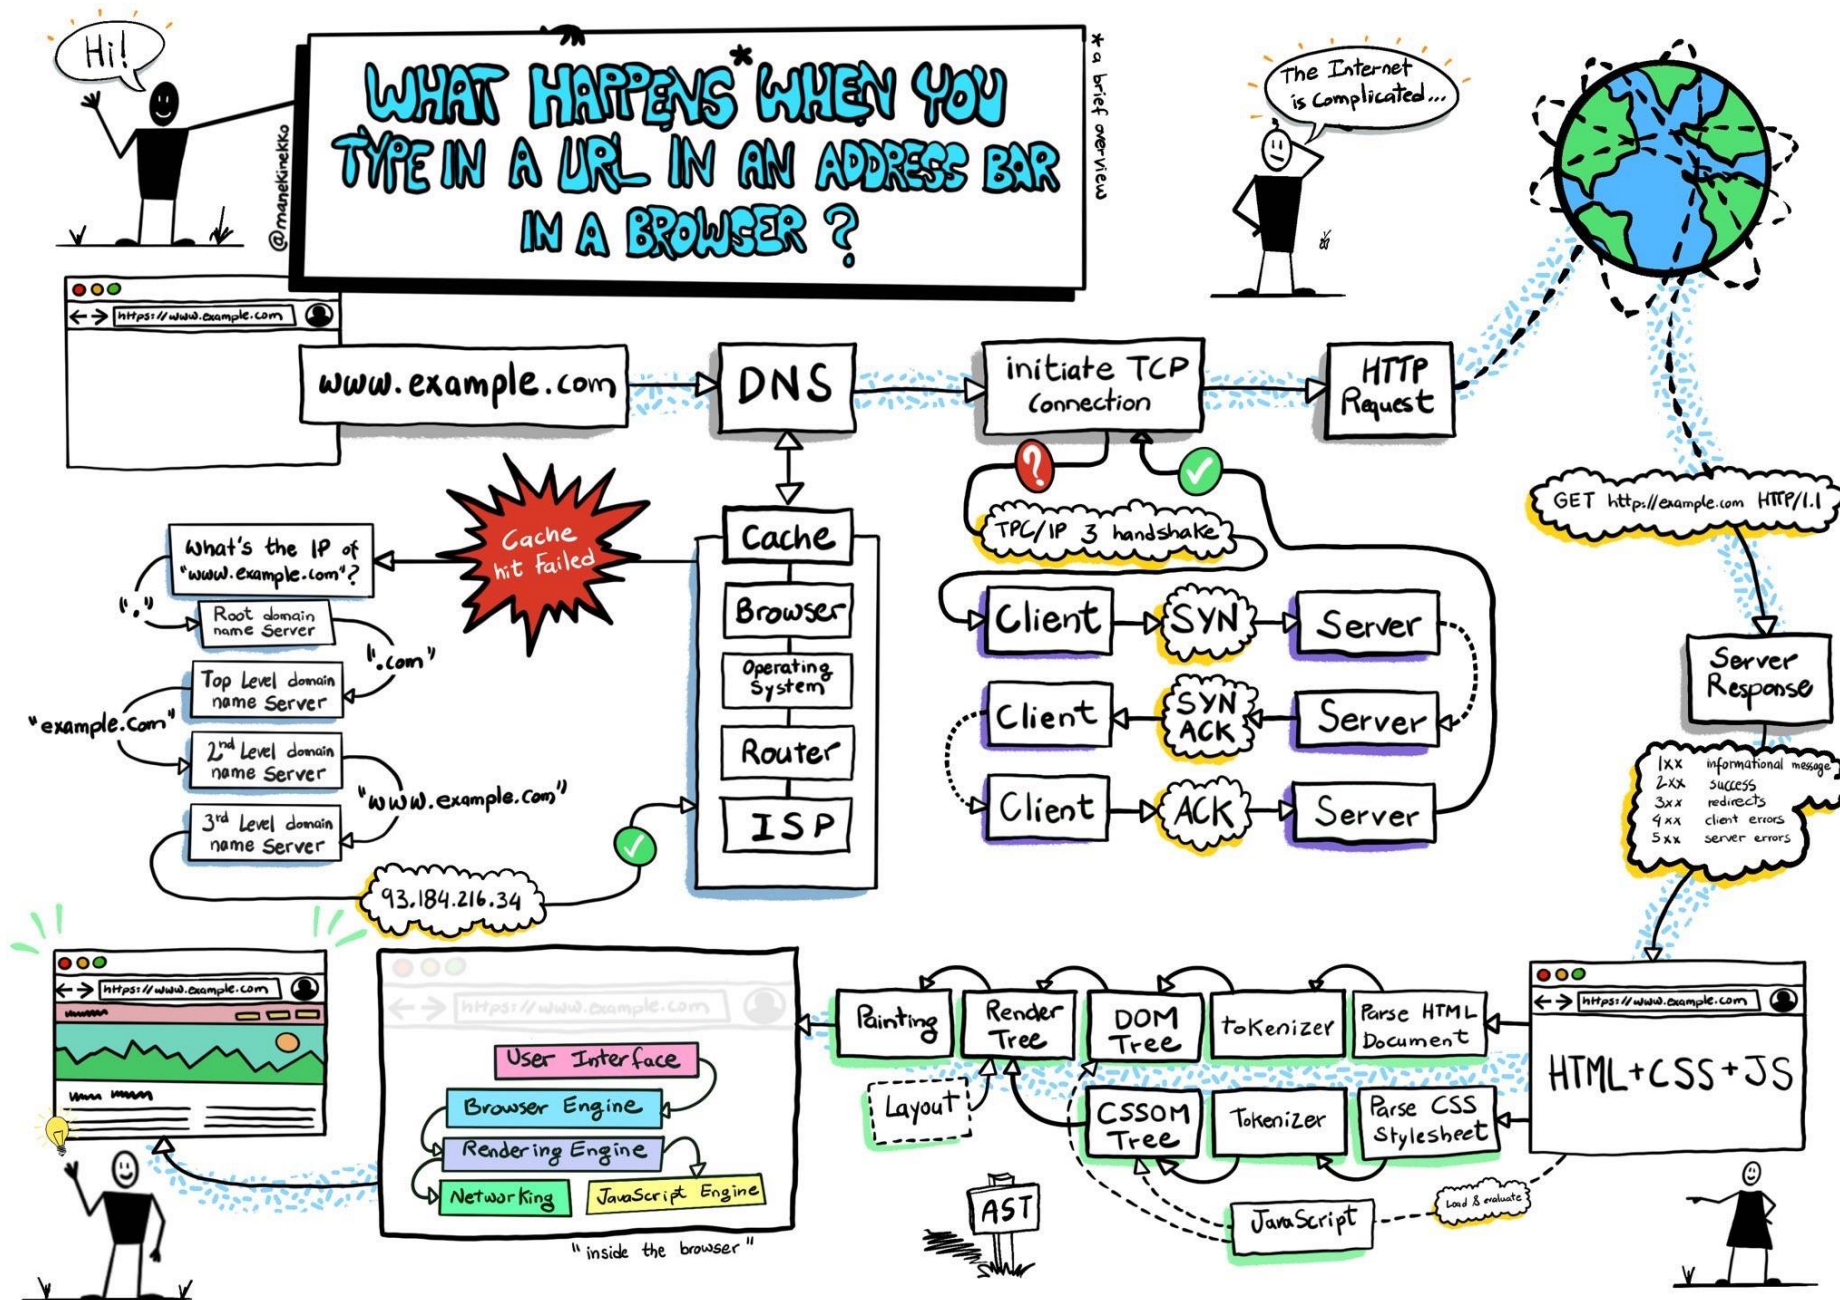
# Part 1. Introduction



# Which Software?

- This section of the course will focus on **systems software**
- **Systems software** is designed to provide a platform for other software
- Examples:
  - Operating systems
  - Hardware drivers
  - Language runtimes
  - Virtual machines
  - Frameworks
  - Performance critical software





C and C++ power the majority of systems software

# C and C++

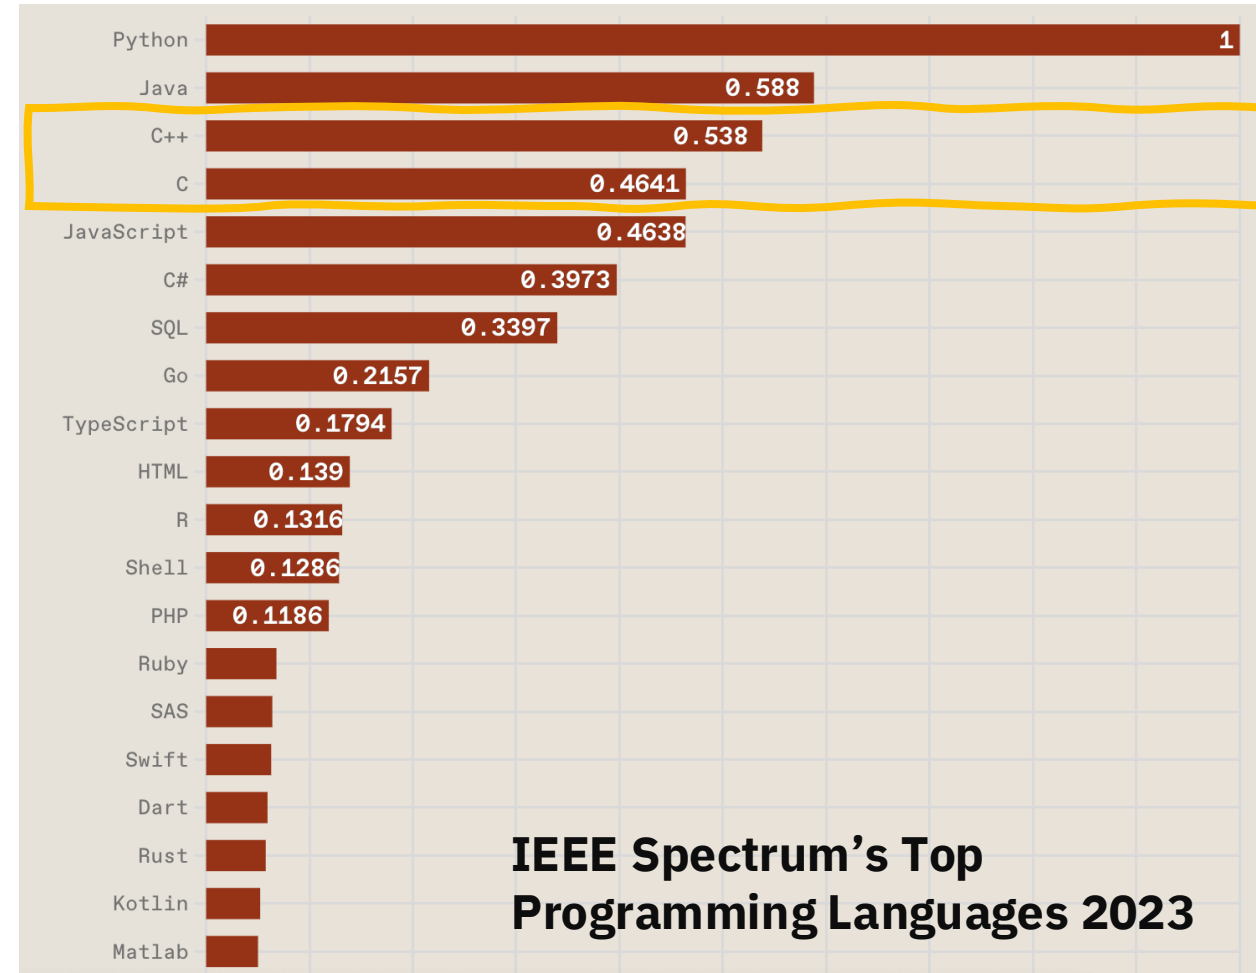
- Versatile and efficient! Pereira et al. [SLE'17]
  - It is fast
  - Uses the least energy
  - Has a small memory footprint

**Table 4.** Normalized global results for Energy, Time, and Memory

Total					
	Energy		Time		Mb
(c) C	1.00	(c) C	1.00	(c) Pascal	1.00
(c) Rust	1.03	(c) Rust	1.04	(c) Go	1.05
(c) C++	1.34	(c) C++	1.56	(c) C	1.17
(c) Ada	1.70	(c) Ada	1.85	(c) Fortran	1.24
(v) Java	1.98	(v) Java	1.89	(c) C++	1.34
(c) Pascal	2.14	(c) Chapel	2.14	(c) Ada	1.47
(c) Chapel	2.18	(c) Go	2.83	(c) Rust	1.54
(v) Lisp	2.27	(c) Pascal	3.02	(v) Lisp	1.92
(c) Ocaml	2.40	(c) Ocaml	3.09	(c) Haskell	2.45
(c) Fortran	2.52	(v) C#	3.14	(i) PHP	2.57
(c) Swift	2.79	(v) Lisp	3.40	(c) Swift	2.71
(c) Haskell	3.10	(c) Haskell	3.55	(i) Python	2.80
(v) C#	3.14	(c) Swift	4.20	(c) Ocaml	2.82
(c) Go	3.23	(c) Fortran	4.20	(v) C#	2.85
(i) Dart	3.83	(v) F#	6.30	(i) Hack	3.34
(v) F#	4.13	(i) JavaScript	6.52	(v) Racket	3.52
(i) JavaScript	4.45	(i) Dart	6.67	(i) Ruby	3.97
(v) Racket	7.91	(v) Racket	11.27	(c) Chapel	4.00
(i) TypeScript	21.50	(i) Hack	26.99	(v) F#	4.25
(i) Hack	24.02	(i) PHP	27.64	(i) JavaScript	4.59
(i) PHP	29.30	(v) Erlang	36.71	(i) TypeScript	4.69
(v) Erlang	42.23	(i) Jruby	43.44	(v) Java	6.01
(i) Lua	45.98	(i) TypeScript	46.20	(i) Perl	6.62
(i) Jruby	46.54	(i) Ruby	59.34	(i) Lua	6.72
(i) Ruby	69.91	(i) Perl	65.79	(v) Erlang	7.20
(i) Python	75.88	(i) Python	71.90	(i) Dart	8.64
(i) Perl	79.58	(i) Lua	82.91	(i) Jruby	19.84

# C and C++

- Versatile and efficient! Pereira et al. [SLE'17]
  - It is fast
  - Uses the least energy
  - Has a small memory footprint
- Popular



<https://spectrum.ieee.org/the-top-programming-languages-2023>



# C and C++


- Versatile and efficient! Pereira et al. [SLE'17]
  - It is fast
  - Uses the least energy
  - Has a small memory footprint
- Popular
- Error prone!
  - No memory safety
    - Developer handles
  - Leads to vulnerabilities
    - Powerful exploits can lead to complete takeover of the system



**Matthew Green** @matthew\_d\_green · 2 Nov 2015

Just a reminder: everything your beautiful 'safe' language depends on is still written in C, and also we all die alone.

## Vulnerabilities matching “overflow” Total Matches By Year



<https://nvd.nist.gov/vuln/search>



# Memory Safety

- What happens when line 6 executes?



```
1. void foo()  
2. {  
3.     int a;  
4.     char buffer[4];  
5.     ...  
6.     buffer[4] = 'A';  
7.     ...  
8. }
```

**This is classified as  
"undefined behavior"**

**Whatever you guessed  
may be correct**



# Memory Management

```
int mytest(char *str)
{
    char *buf = malloc(16); 
    /* Do something with buf */
    free(buf); 
    return 0;
}
```

**Developers are  
responsible to allocate  
and de-allocate memory**

# Memory Management

```
int mytest(char *str)
{
    char *buf = malloc(16);

    /* Do something with buf */

    free(buf);
    buf[1] = '\\0'; ←
    return 0;
}
```

**Developers are  
responsible to allocate  
and de-allocate memory**

**What happens if buf is  
used after free?**

# Memory Corruption

---

“Memory corruption occurs in a computer program when the contents of a memory location are unintentionally modified due to programming errors; this is termed **violating memory safety**.

When the corrupted memory contents are used later in that program, it leads either to program crash or to **strange and bizarre program behavior**. “

--wikipedia

# Common Vulnerabilities

---

- Overflows: Writing beyond the end of a buffer
- Underflows: Writing beyond the beginning of a buffer
- Format string vulnerabilities: Evaluating input string as format string
- Uninitialized memory: Using pointer before initialization
- Null pointer dereferences: Using NULL pointers
- Use-after-free: Using memory after it has been freed
- Type confusion: Assume a variable/object has the wrong type

# Why Are These Bugs (Serious) Vulnerabilities?

---

- **Arbitrary code execution (ACE)**, put simply, is a vulnerability that allows attackers to inject their own malicious code onto a target system without user awareness or permission
- A **Remote code execution (RCE)** attack is one where an attacker can run malicious code on an organization's computers or network

# Generic Attack Types

---

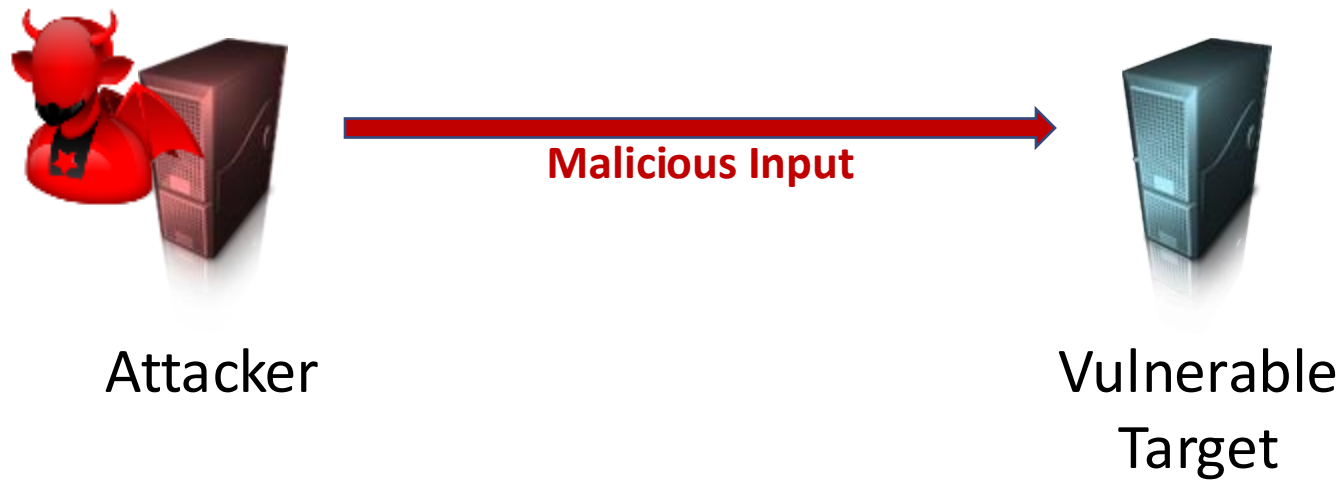
Local Attacks

Remote Attacks

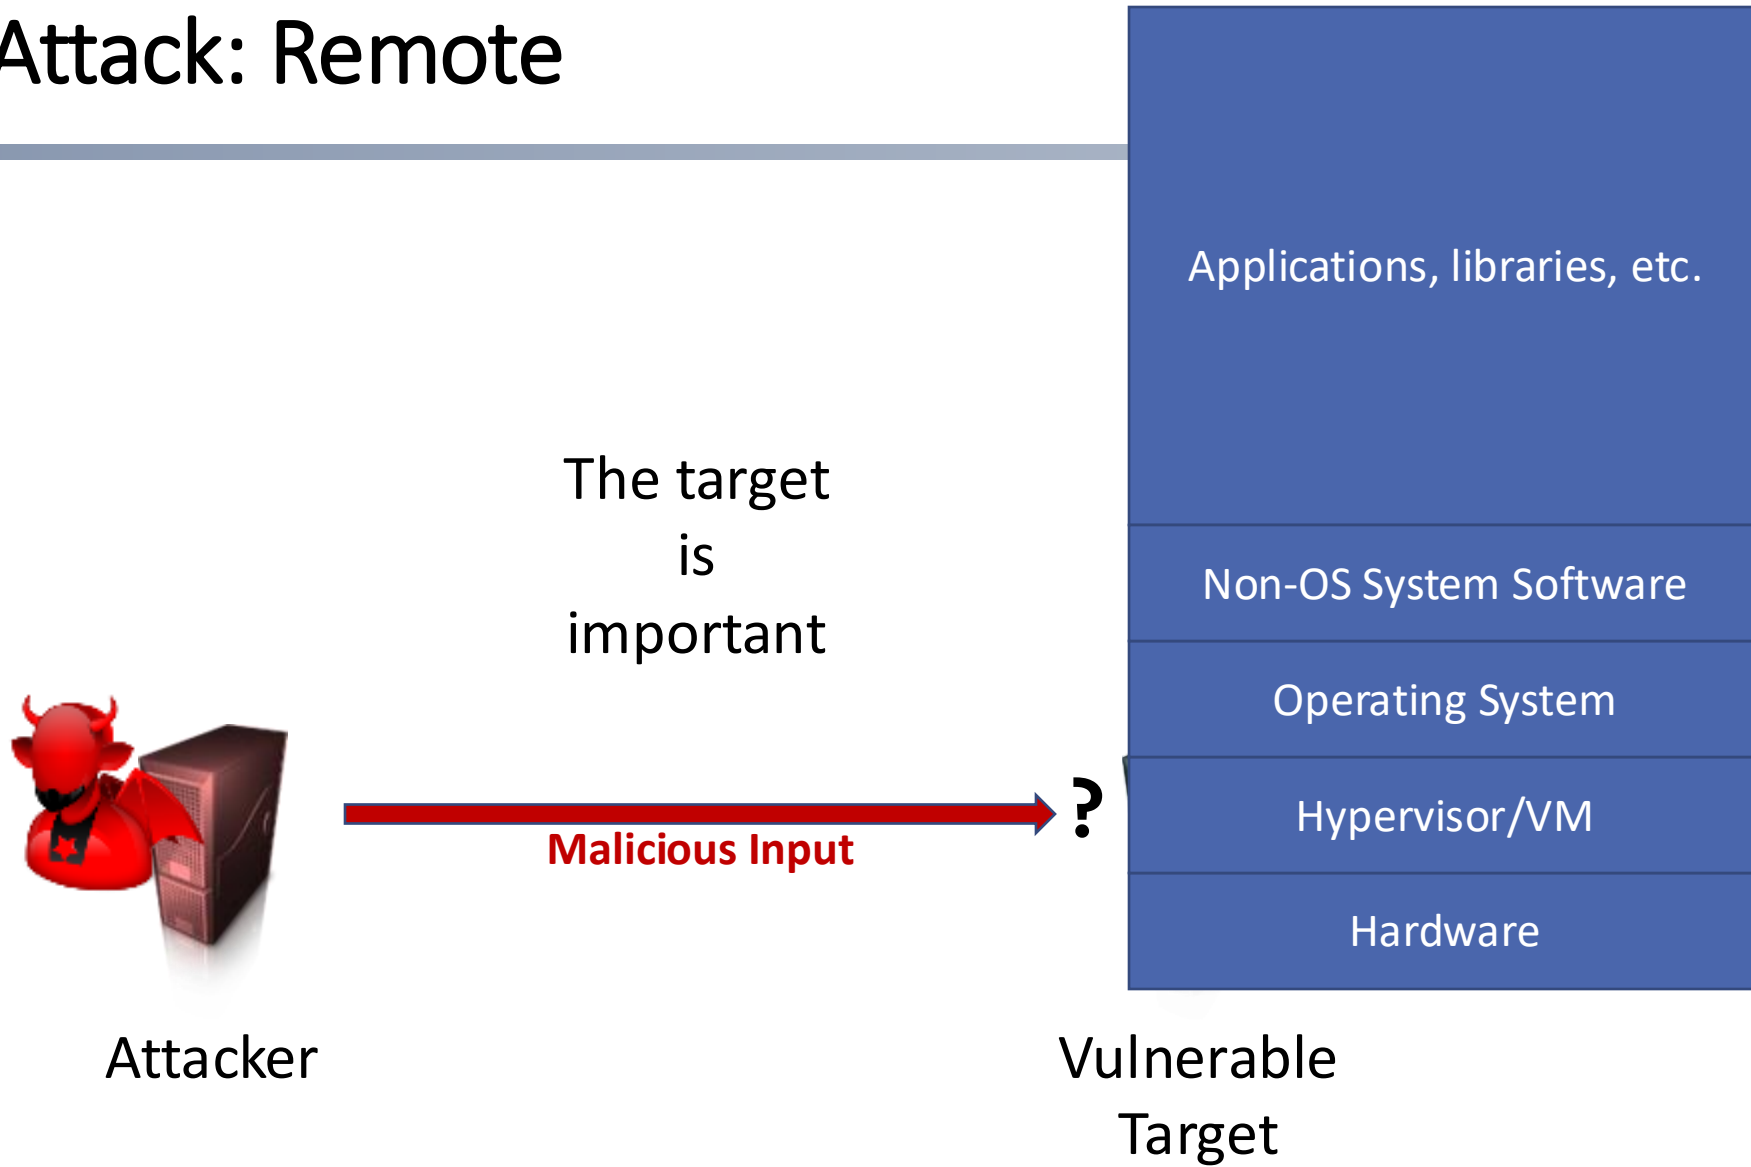


# Types of Attack: Remote

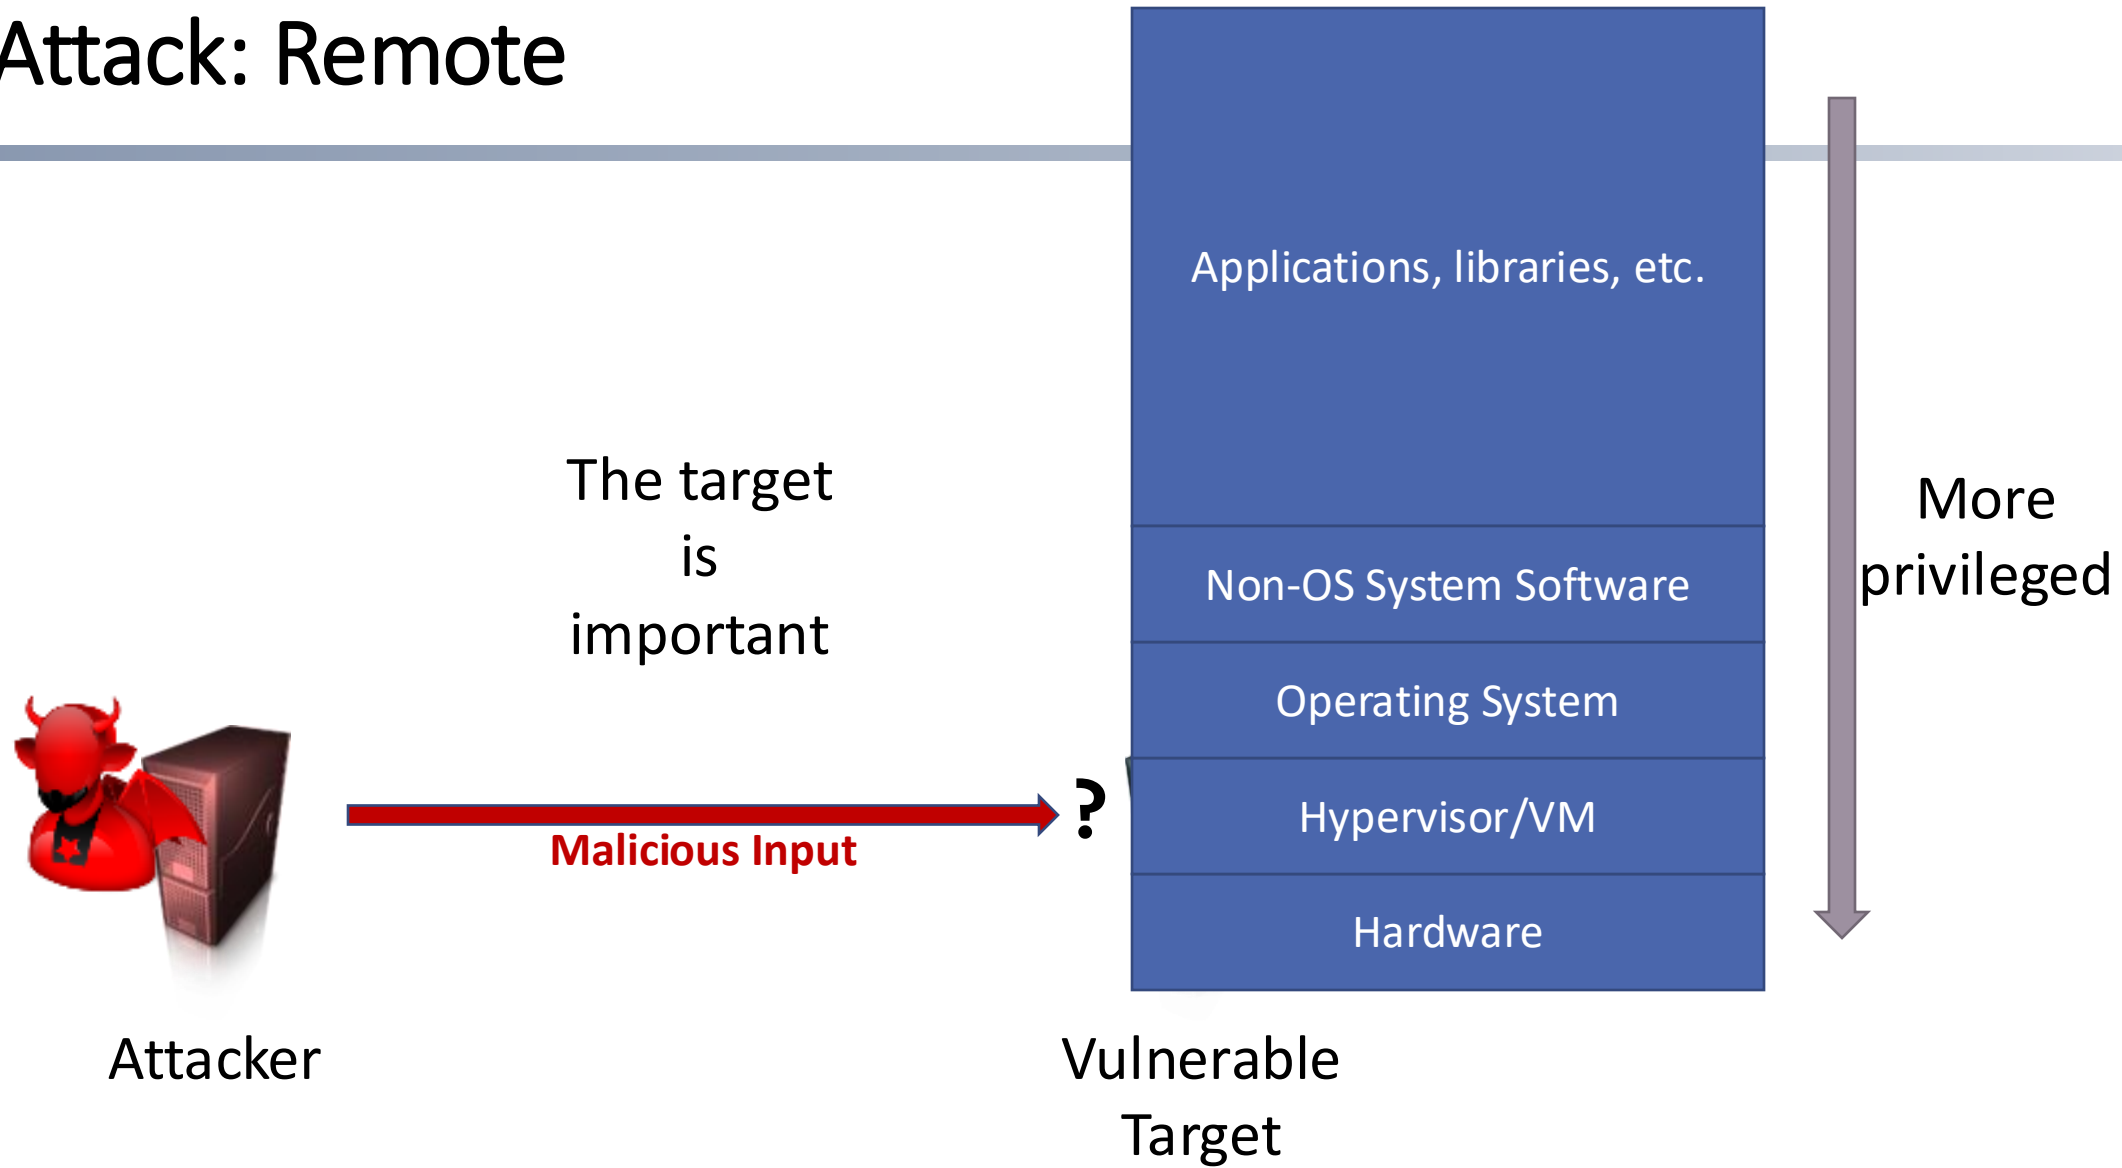
---




# Types of Attack: Remote



# Types of Attack: Remote

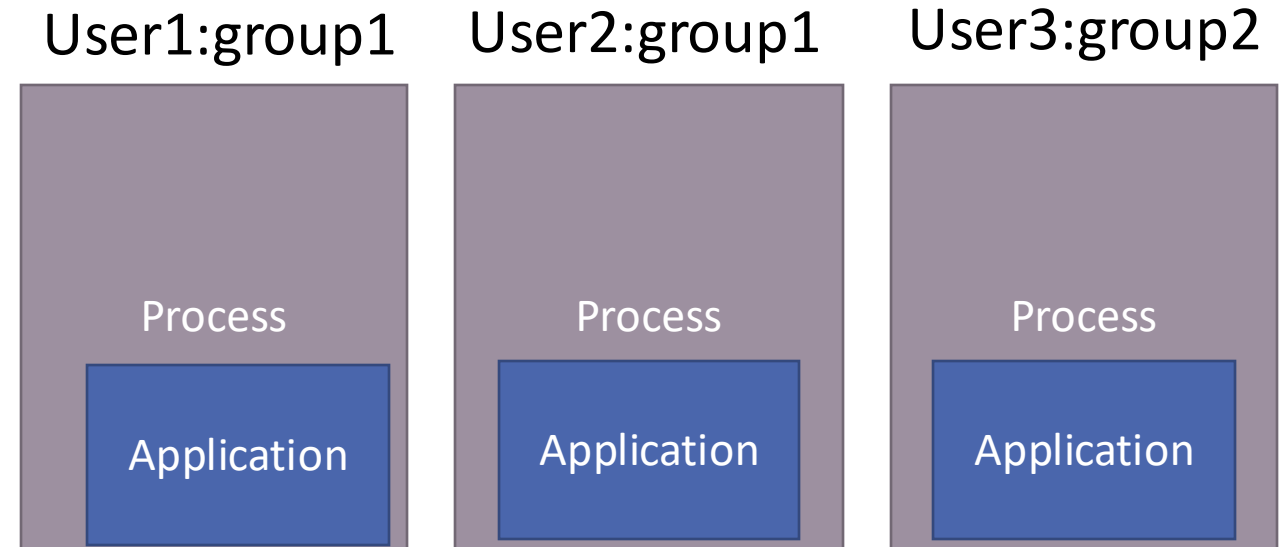


# Types of Attack: Remote



# Processes and Permissions

- Example shows UNIX-style permissions

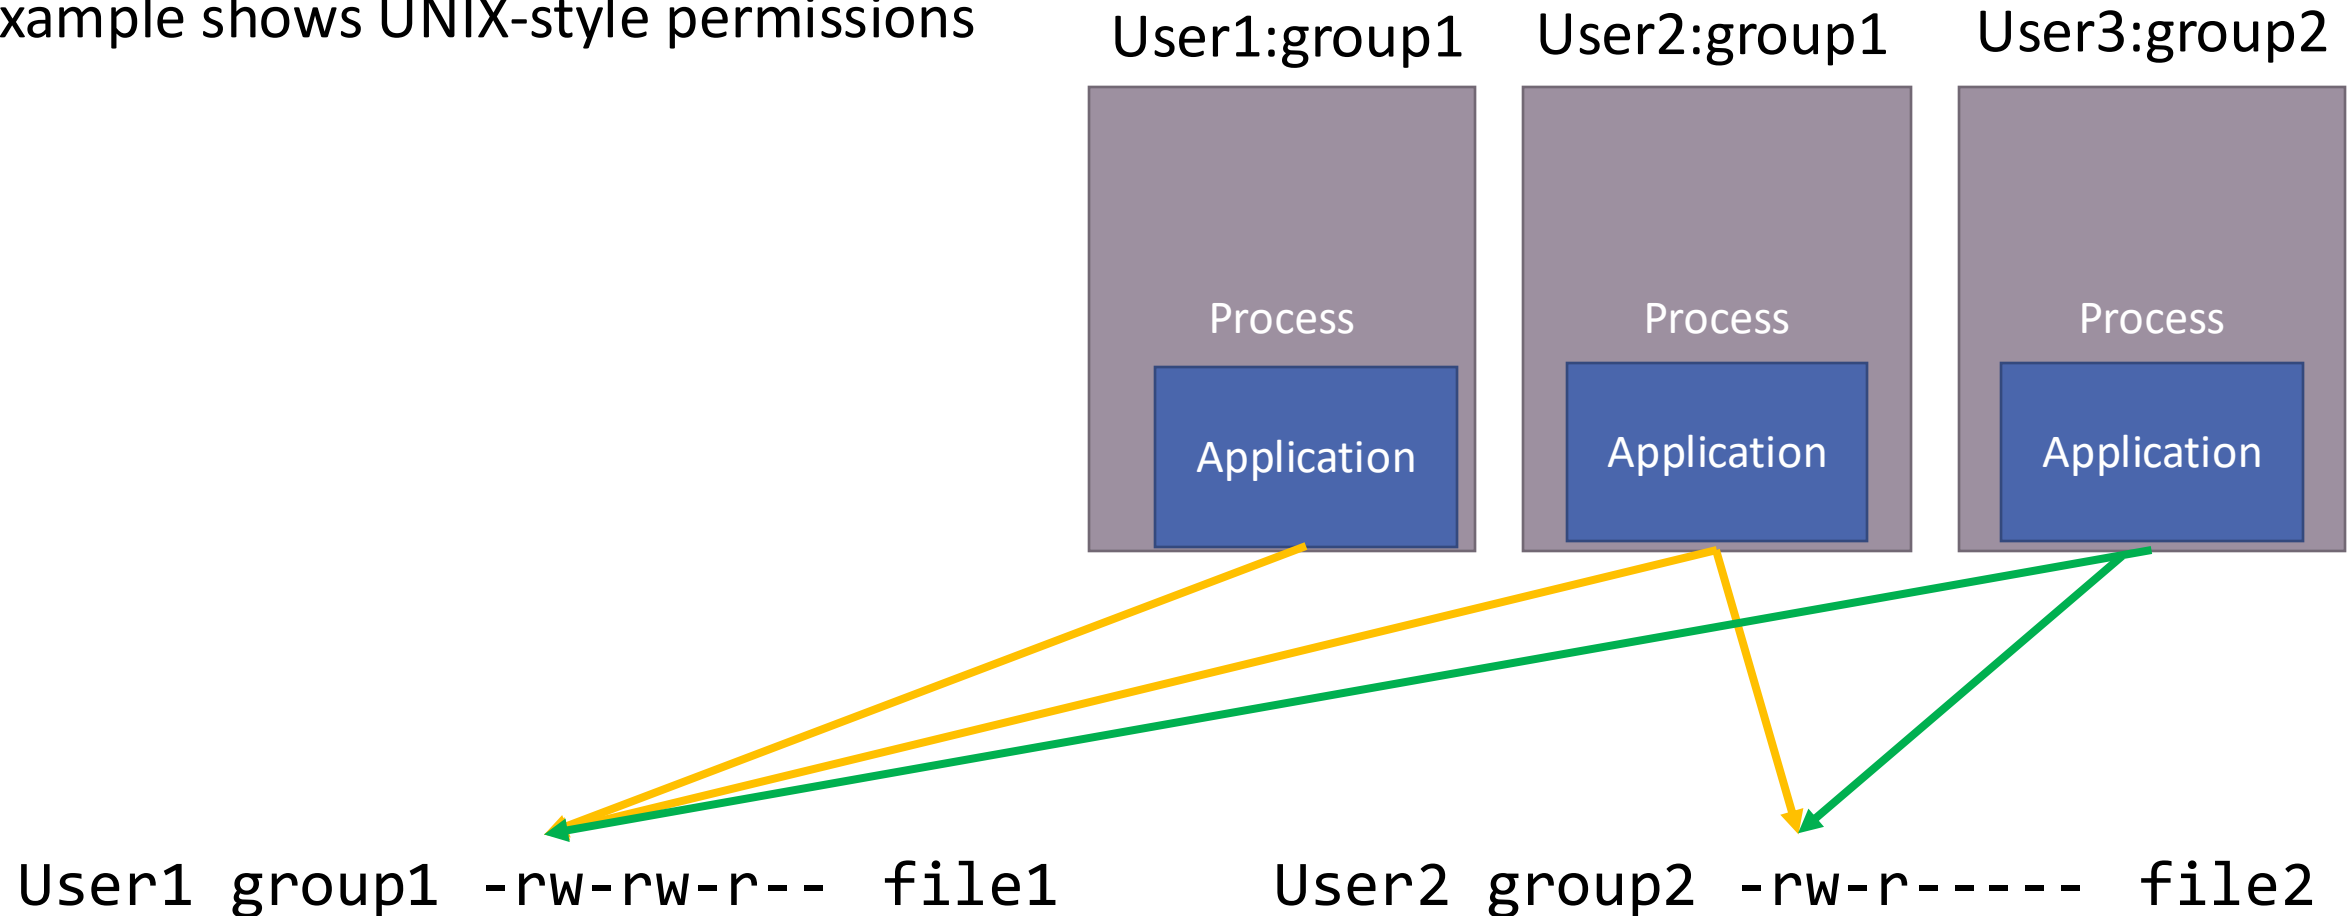


User1 group1 -rw-rw-r-- file1

User2 group2 -rw-r----- file2

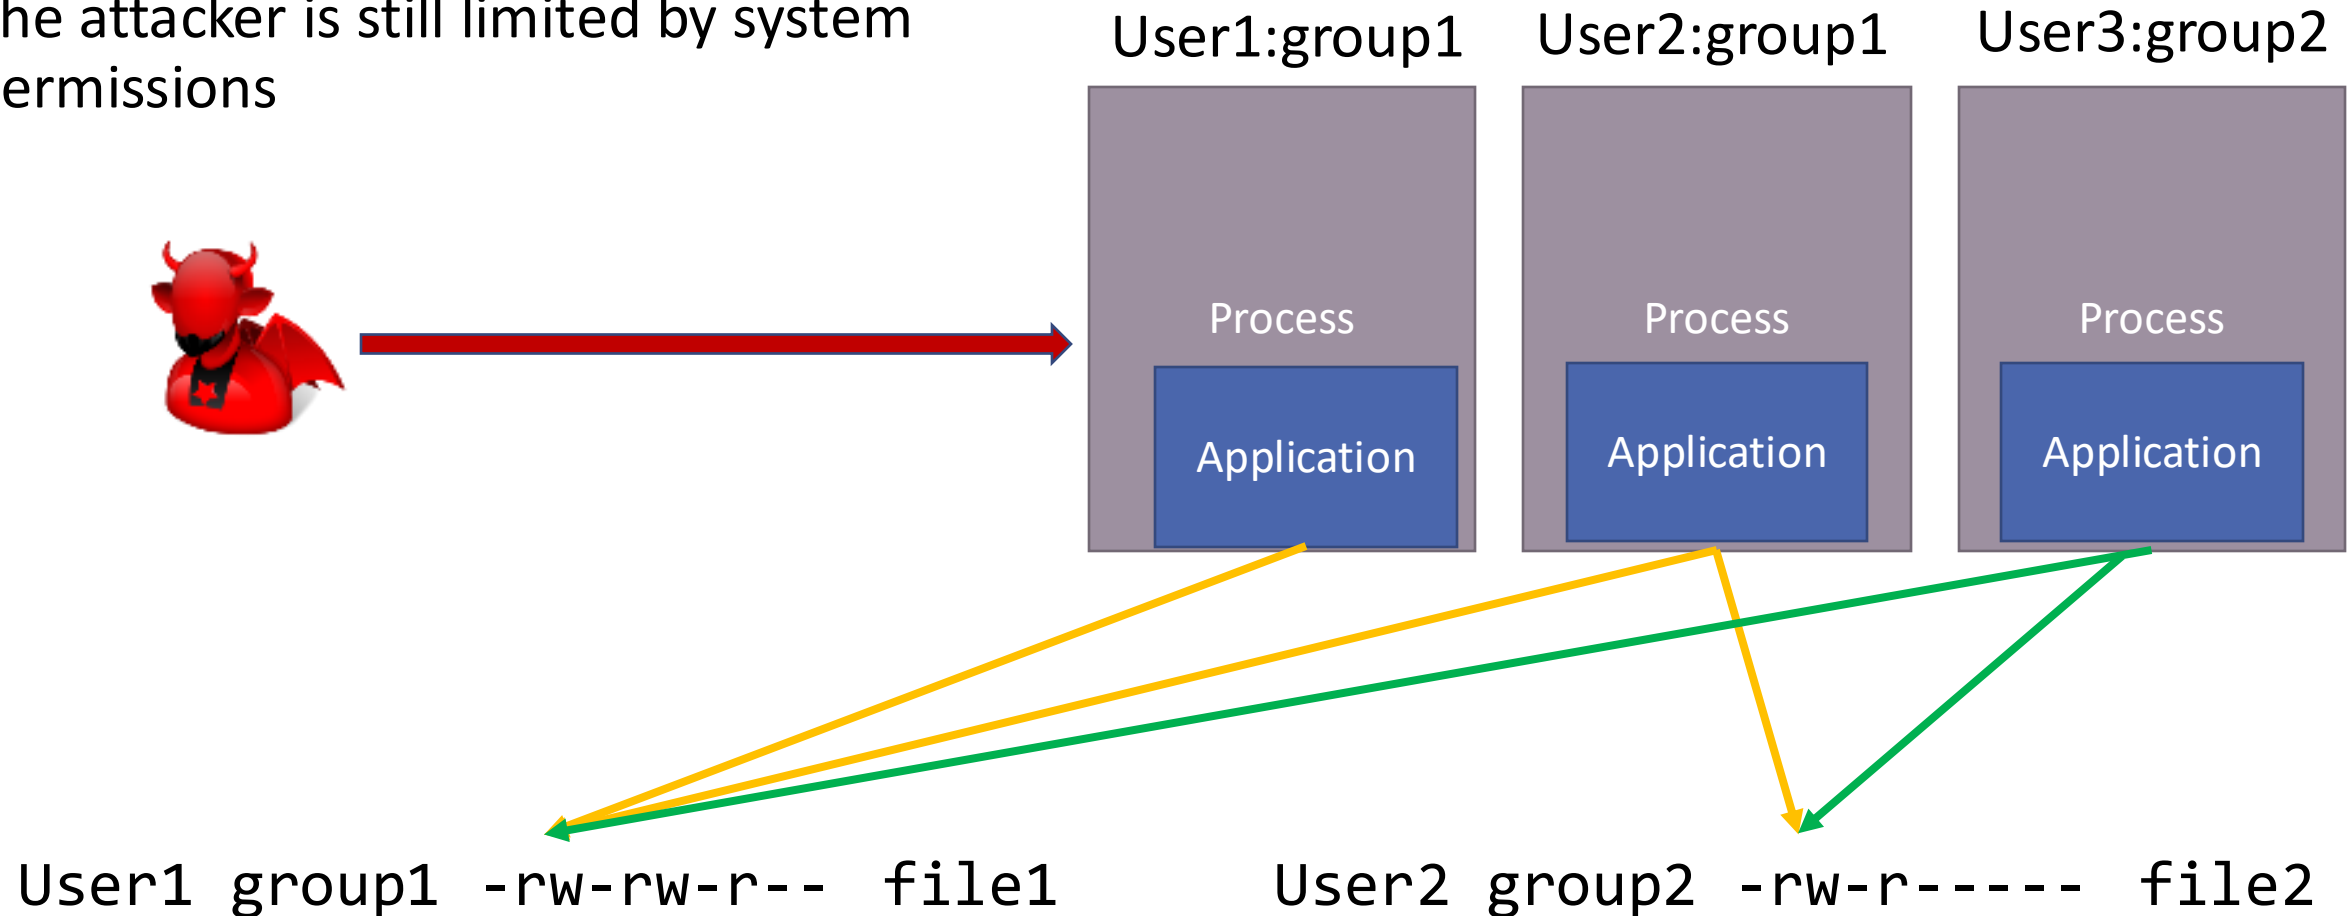
# Processes and Permissions

- Example shows UNIX-style permissions



# Processes and Permissions

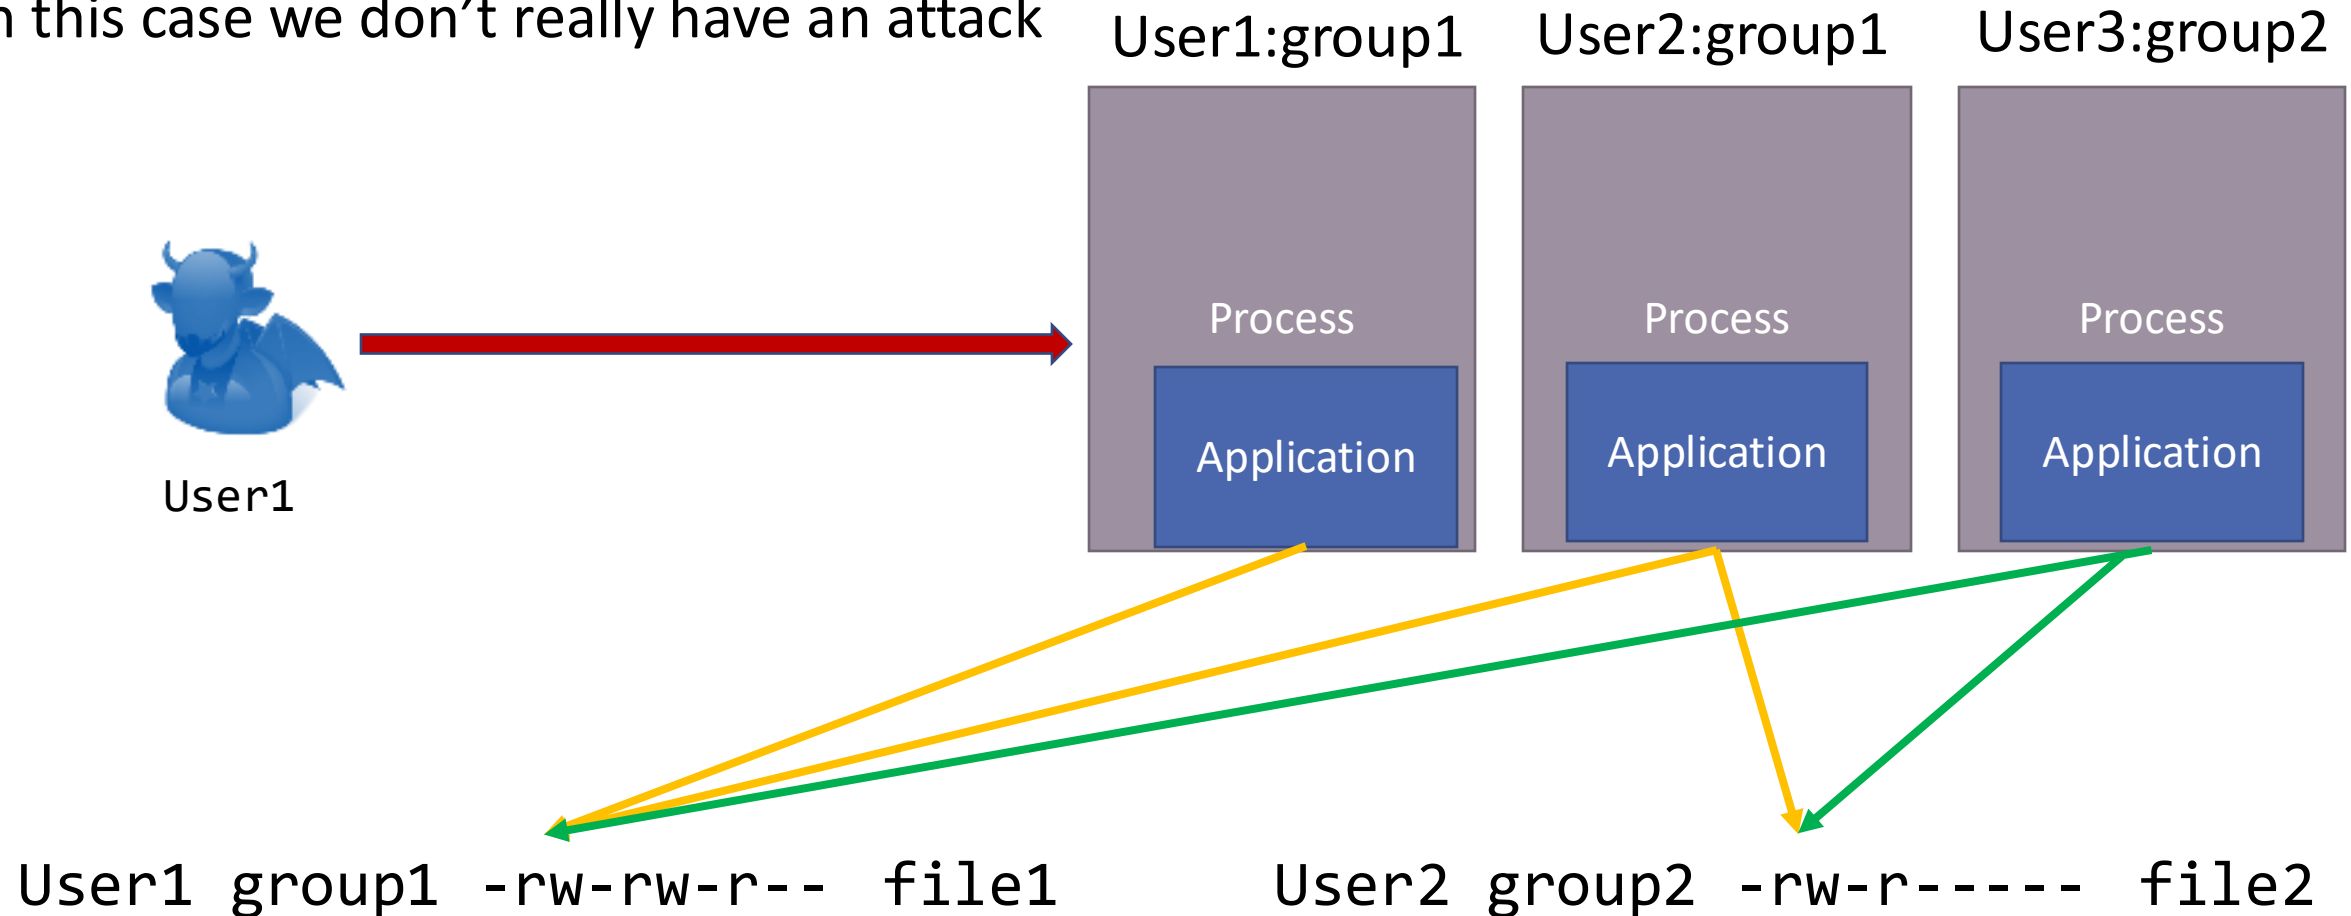
- The attacker is still limited by system permissions





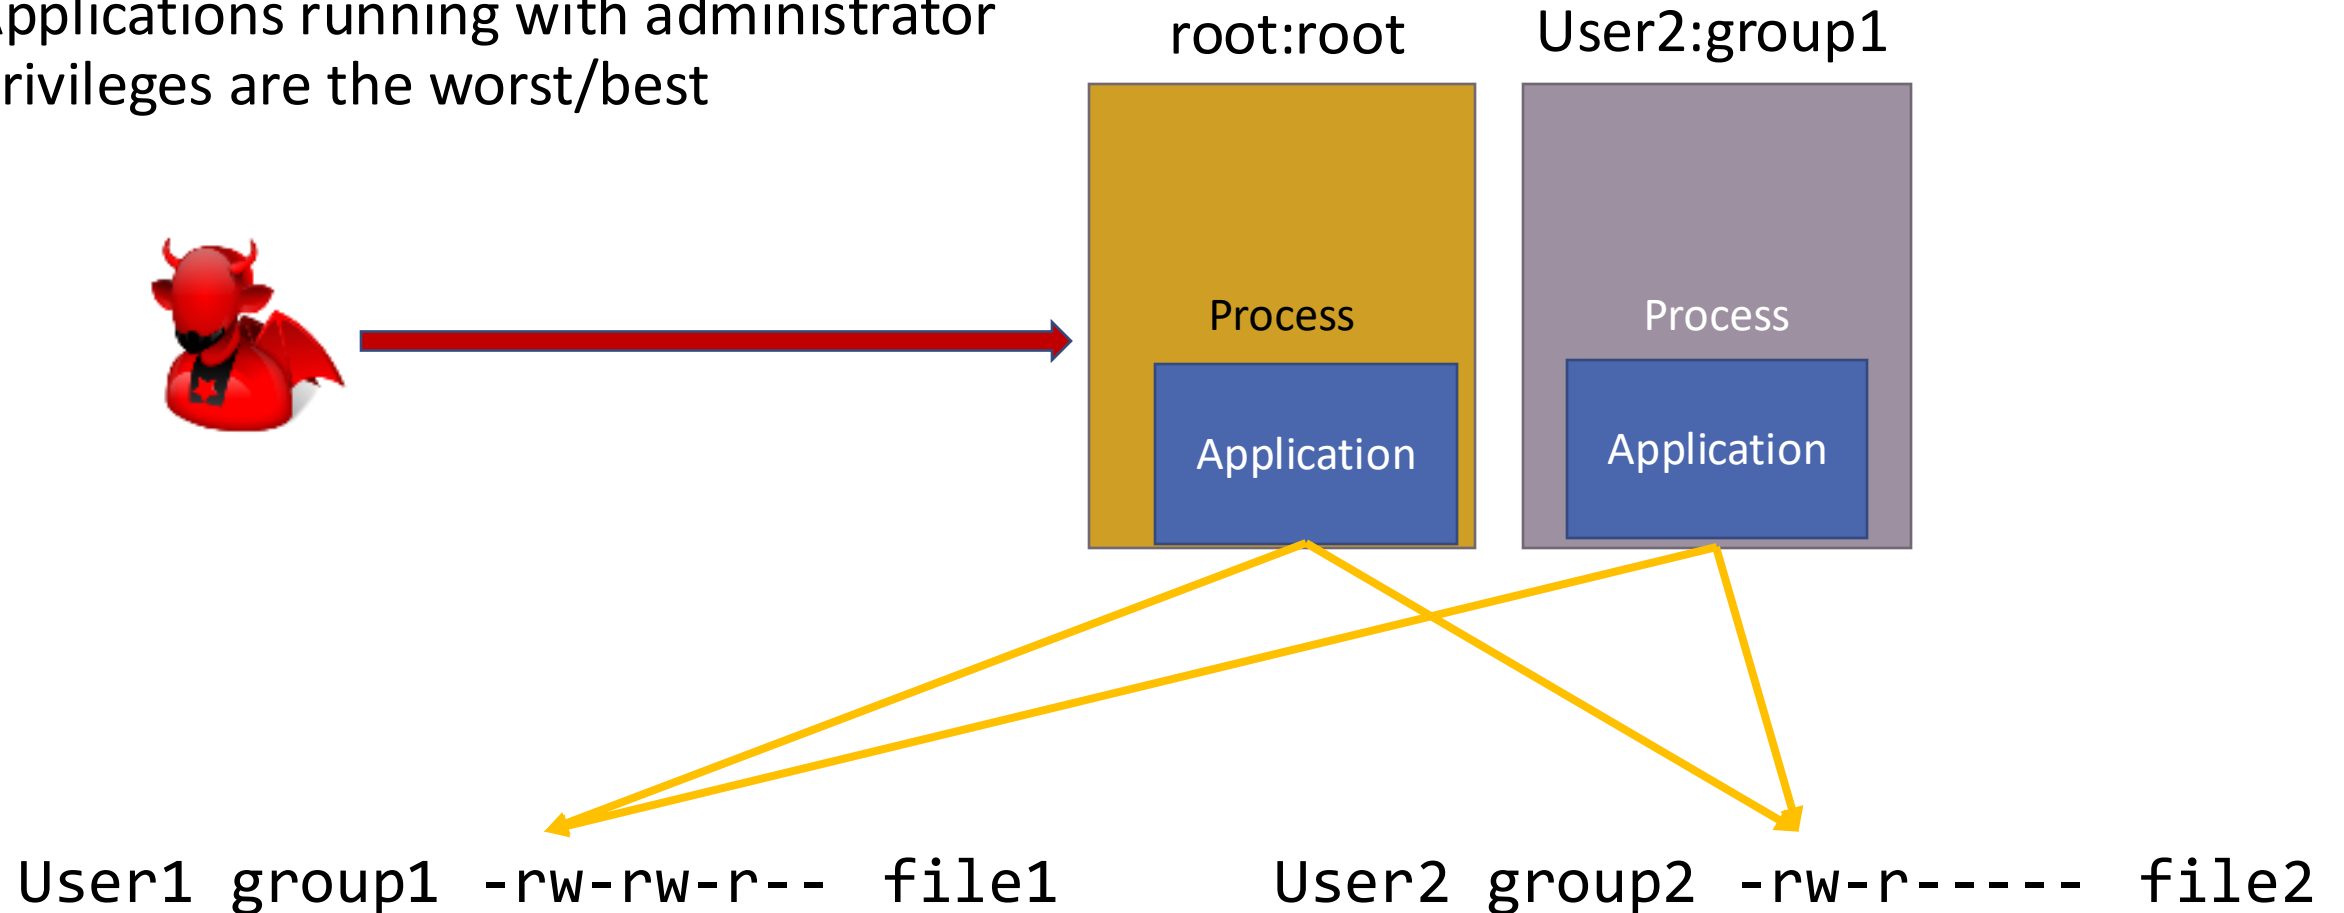
# Processes and Permissions

- In this case we don't really have an attack



# Processes and Permissions

- Applications running with administrator privileges are the worst/best




# Processes and Permissions


- Applications designed carefully avoid running as administrator or only run part of their code as one



User2:group1



root:root



User1 group1 -rw-rw-r-- file1

User2 group2 -rw-r----- file2

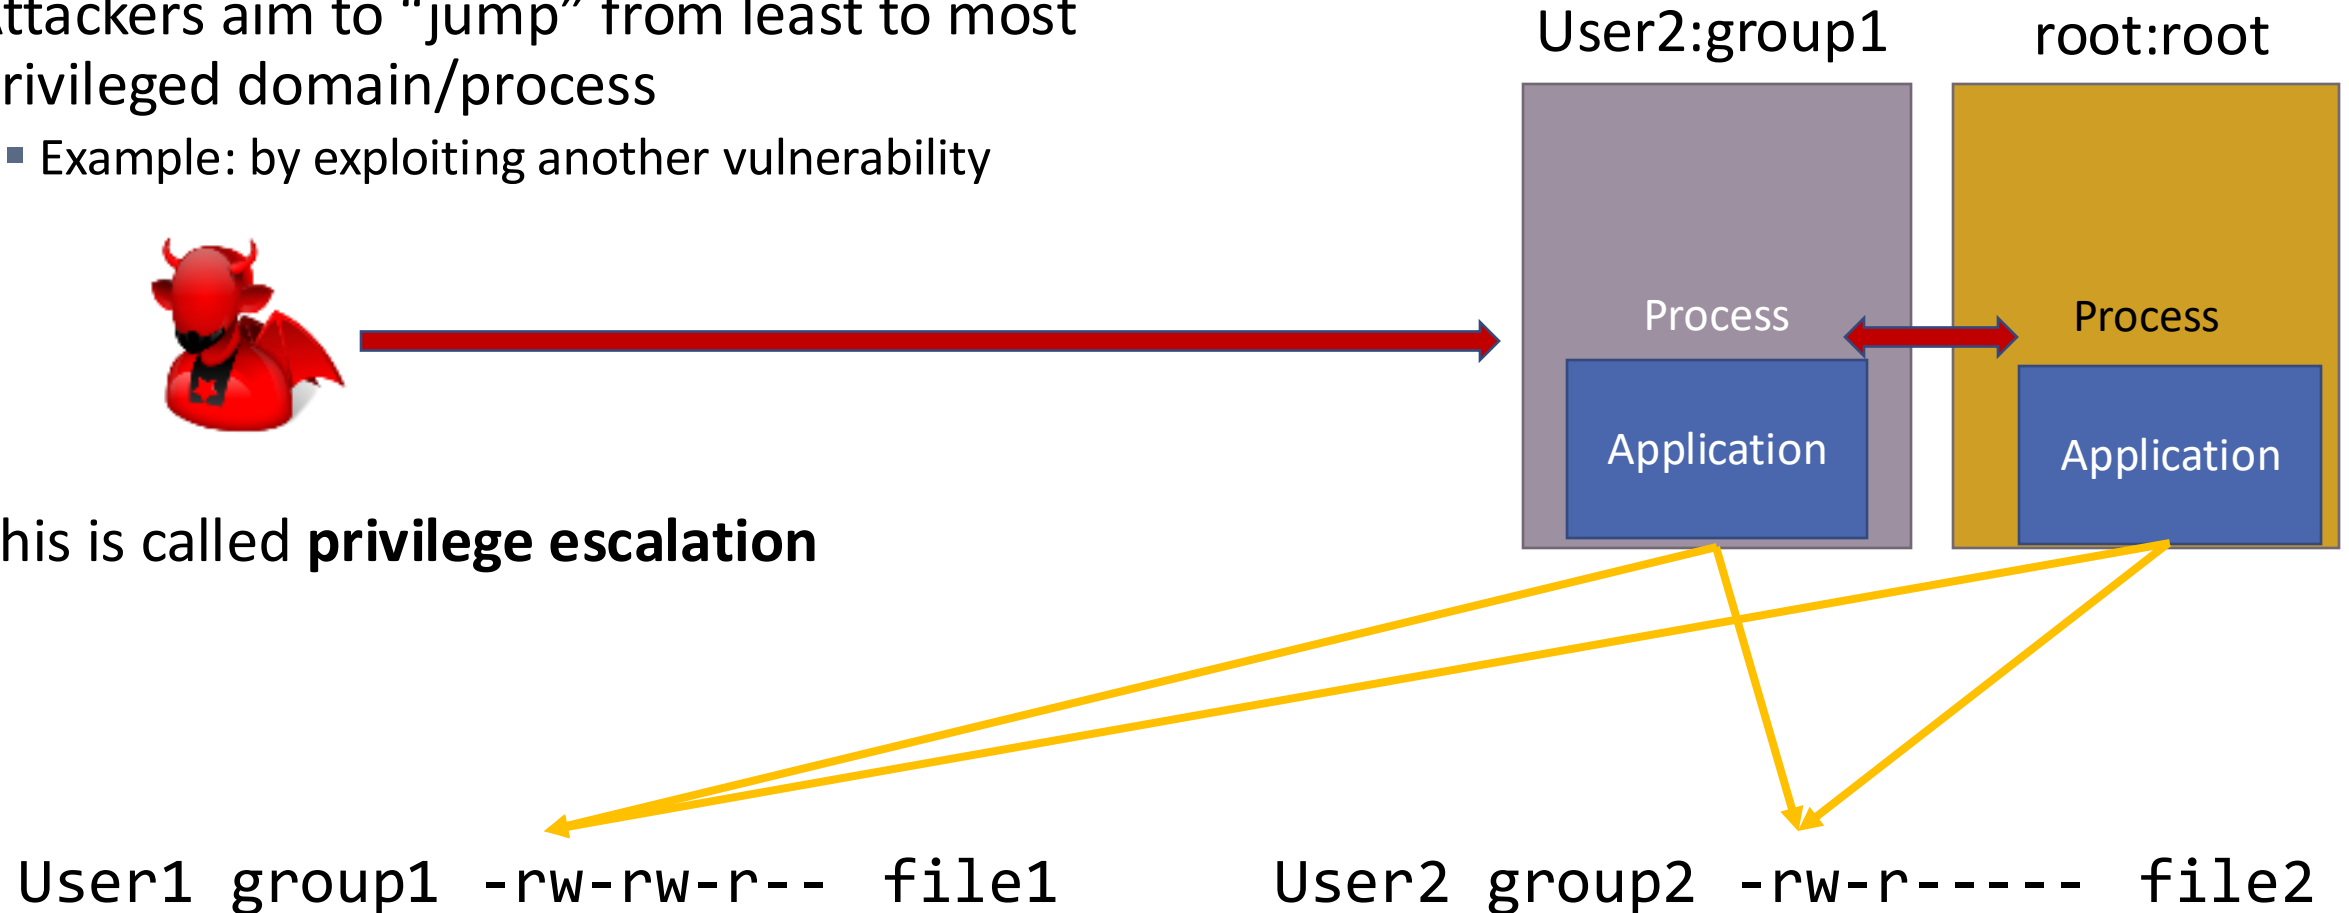


# Processes and Permissions

- Attackers aim to “jump” from least to most privileged domain/process
  - Example: by exploiting another vulnerability



- This is called **privilege escalation**

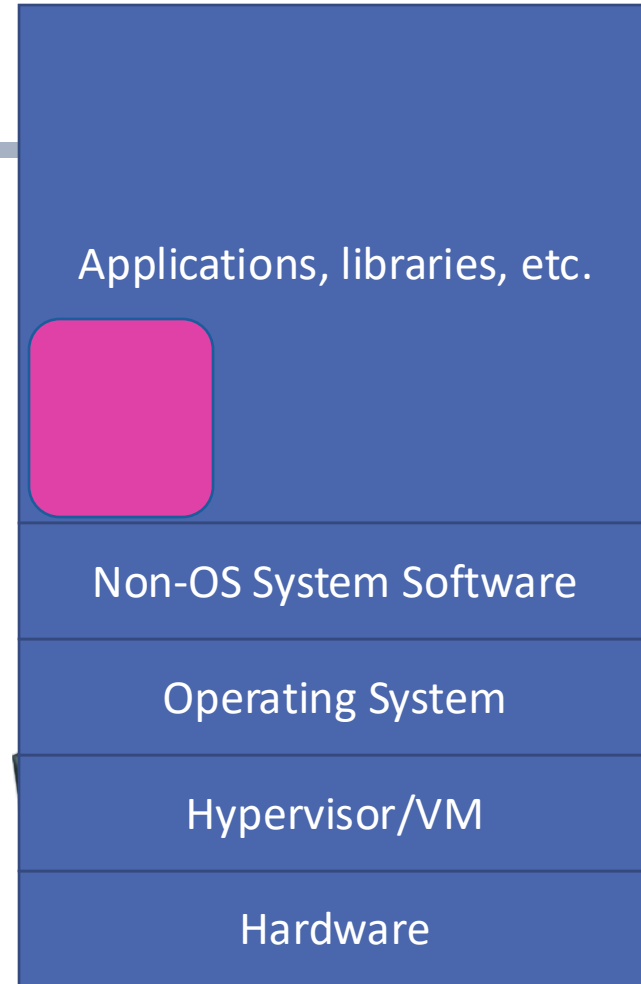


# Typical Privilege Escalation



Attacker

1. Compromise software running in user process



More  
privileged

Vulnerable  
Target

# Typical Privilege Escalation

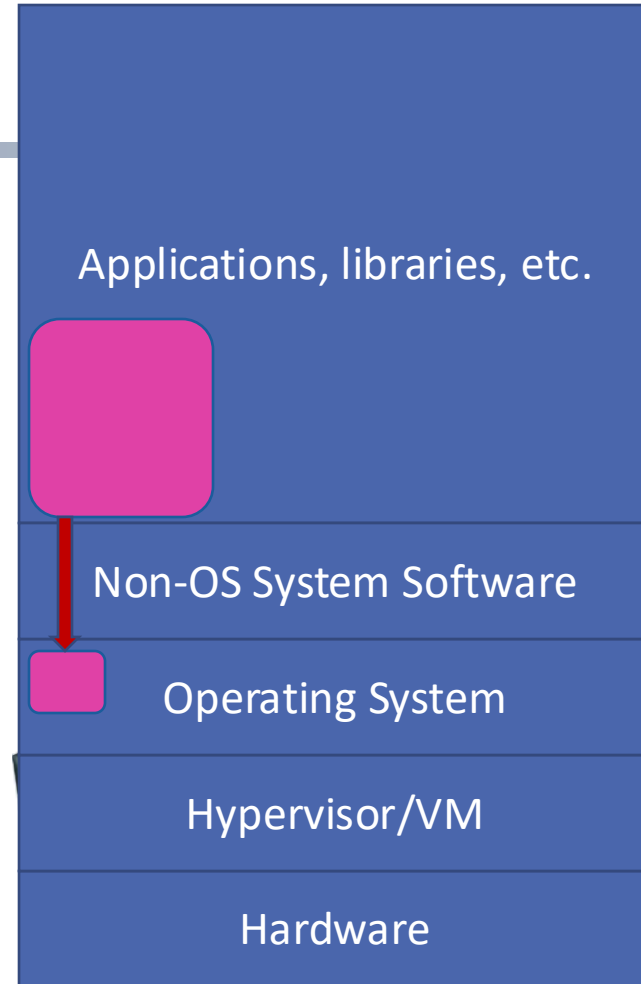


Attacker

1. Compromise software running in user process



2. Compromise operating system




More privileged

Vulnerable Target

# Types of Attack: Local


- The attacker already has access to the system but is limited by the access control in place
  - Example: non-admin user in a UNIX-based system






# Types of Attack: Local

- The attacker already has access to the system but is limited by the access control in place
  - Example: non-admin user in a UNIX-based system
- Goal: elevate privileges
  - Exploit a vulnerability in software running in process with privileges
  - Exploit a vulnerability in the layers below
  - Special case: SUID binaries



# (SUID) Set User ID on Execution

- Example: Enable a user to change their own password



\$ passwd

File contains user meta data (name, shell, etc.)

```
-rw-r--r-- 1 root root ... /etc/passwd
```

```
-rw-r----- 1 root shadow ... /etc/shadow
```

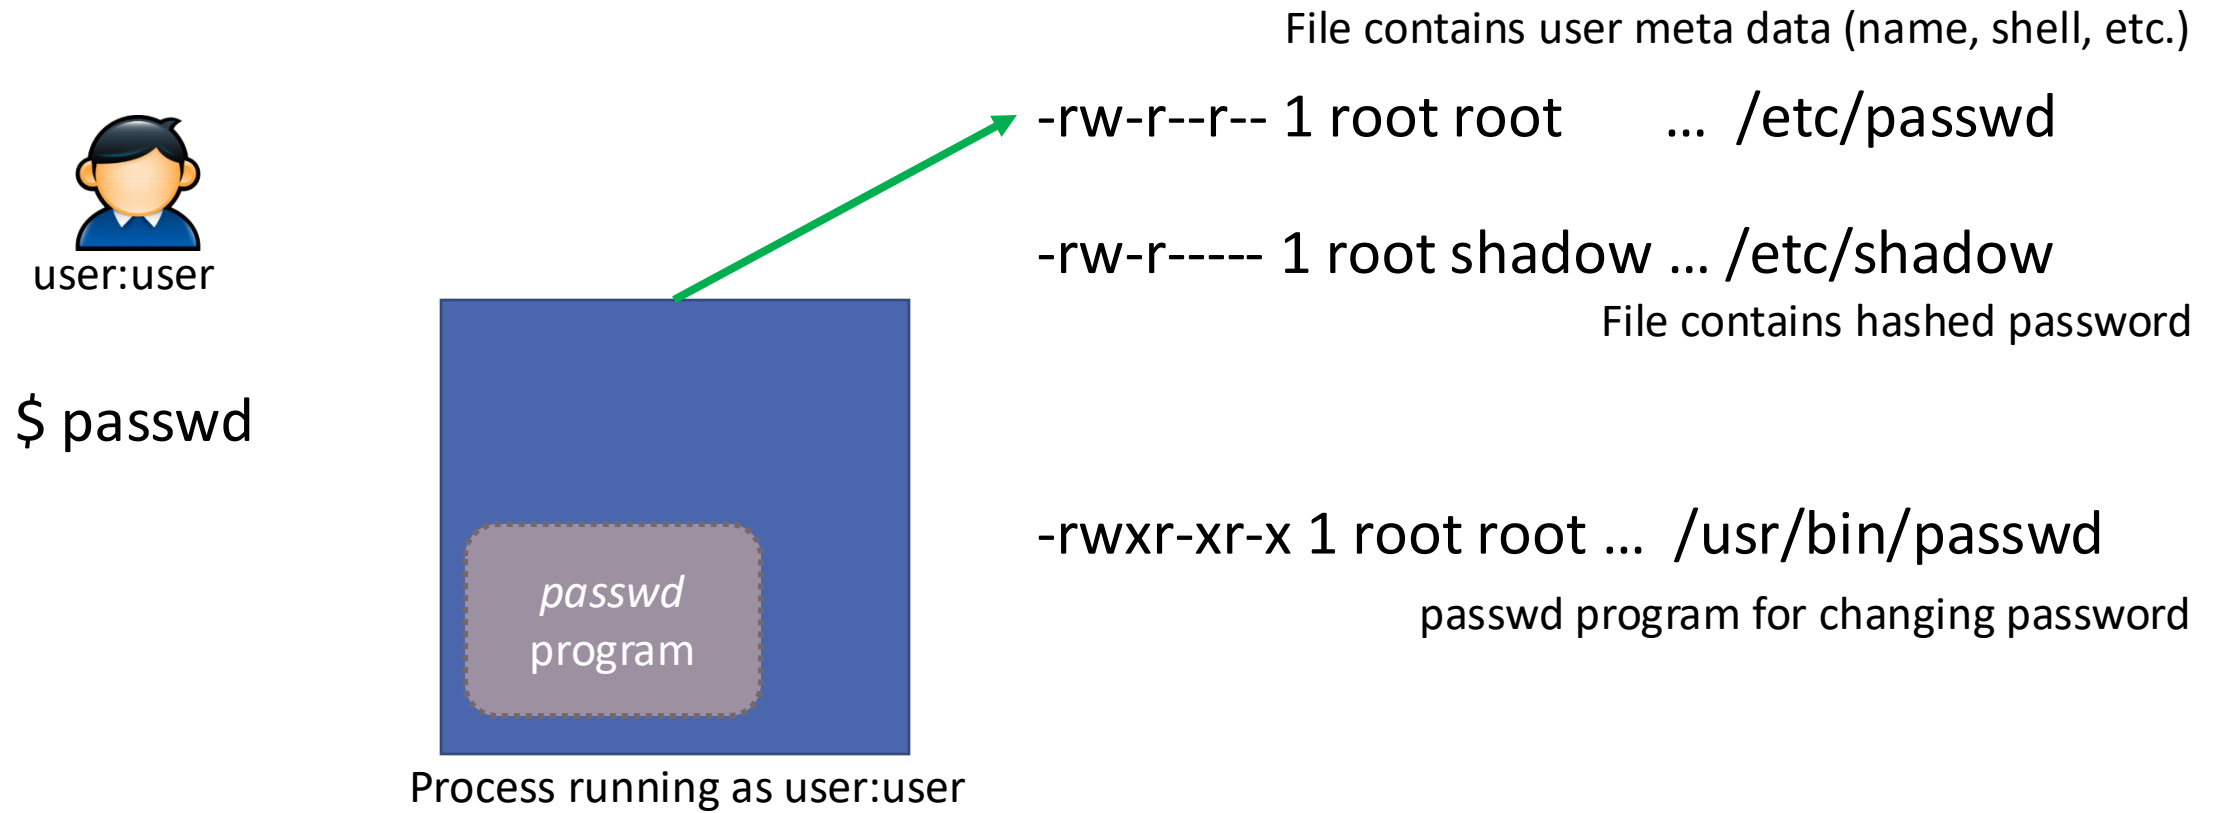
File contains hashed password

```
-rwxr-xr-x 1 root root ... /usr/bin/passwd
```

passwd program for changing password

# (SUID) Set User ID on Execution


- Example: Enable a user to change their own password



# (SUID) Set User ID on Execution

- Example: Enable a user to change their own password

- **SUID programs run as the owner**



\$ passwd



Process running as **root**:user

File contains user meta data (name, shell, etc.)

-rw-r--r-- 1 root root ... /etc/passwd

-rw-r----- 1 root shadow ... /etc/shadow

File contains hashed password

-rwx**s**-xr-x 1 root root ... /usr/bin/passwd

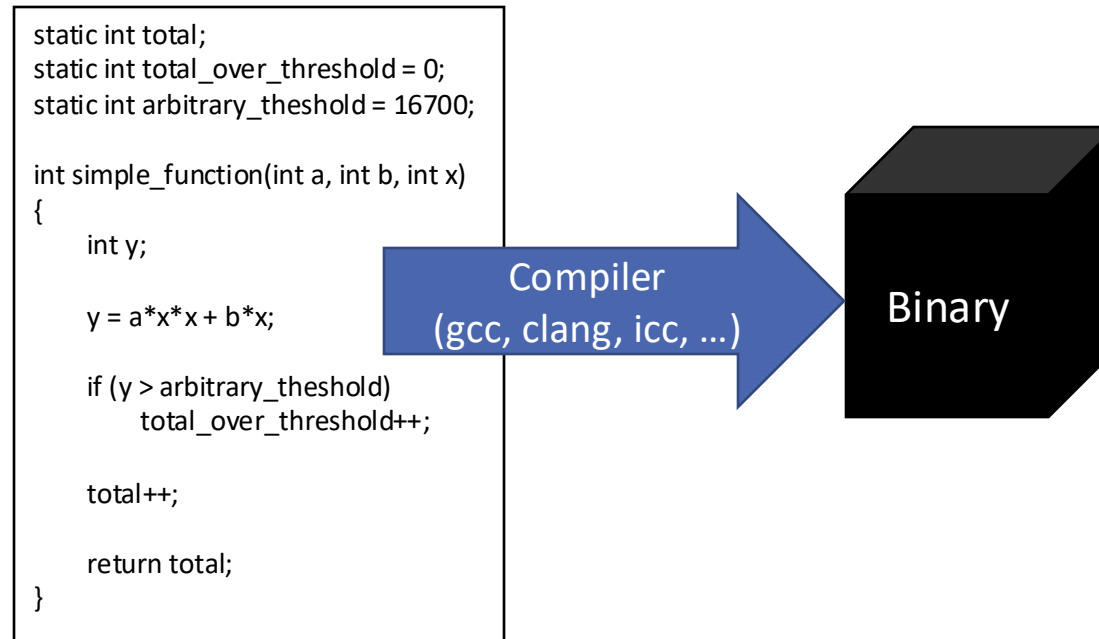
passwd program for changing password

# (SUID) Set User ID on Execution

- Example: Enable a user to change their own password
- **SUID programs run as the owner**
- Useful for performing actions that would otherwise require super-user privileges
- The program is **trusted** to perform only the actions advertised
- Also, SGID programs run as the group
- You can find SUID programs in your system using `find`  
*\$ sudo find /usr/bin -perm -u=s*
- **How can these programs be misused by attackers?**


# Beyond Source Code

- Computer can only execute machine (binary) code
- C/C++ programs are compiled to binary code
- Binary (executable) programs are loaded to memory and executed by the operating system (OS)
  - Using a process



# Computer Architecture Basics

- PC: Program counter
  - Address of next instruction
- Memory
  - Byte addressable array
  - Code and user data
  - Stack to support procedures
- Register file
  - Heavily used program data
- Condition codes
  - Store status information about most recent arithmetic or logical operation
  - Used for conditional branching





# x86 Integer Registers

- General purpose registers
  - On 32-bit architectures EAX, EBX, ECX, EDX, EDI, ESI, ESP, EBP
- The instruction pointer (IP)
  - Also referred to as program counter (PC)
  - EIP on 32-bit
- FLAGS register
  - Used for control flow operations, etc.
  - EFLAGS

register encoding		high 8-bit	low 8-bit	16-bit	32-bit
0		AH (4)	AL	AX	EAX
3		BH (7)	BL	BX	EBX
1		CH (5)	CL	CX	ECX
2		DH (6)	DL	DX	EDX
6		SI		SI	ESI
7		DI		DI	EDI
5		BP		BP	EBP
4		SP		SP	ESP
	31	16	15	0	

	FLAGS		FLAGS	EFLAGS
	IP		IP	EIP
	31	0		

# x86-64 Integer Registers

<b>%rax</b>	<b>%eax</b>
<b>%rbx</b>	<b>%ebx</b>
<b>%rcx</b>	<b>%ecx</b>
<b>%rdx</b>	<b>%edx</b>
<b>%rsi</b>	<b>%esi</b>
<b>%rdi</b>	<b>%edi</b>
<b>%rsp</b>	<b>%esp</b>
<b>%rbp</b>	<b>%ebp</b>

<b>%r8</b>	<b>%r8d</b>
<b>%r9</b>	<b>%r9d</b>
<b>%r10</b>	<b>%r10d</b>
<b>%r11</b>	<b>%r11d</b>
<b>%r12</b>	<b>%r12d</b>
<b>%r13</b>	<b>%r13d</b>
<b>%r14</b>	<b>%r14d</b>
<b>%r15</b>	<b>%r15d</b>

# x86-64 Integer Registers

Can reference low-order bytes too

- d suffix for lower 32-bits (r8d)
- w suffix for lower 16-bits (r8w)
- b suffix for lower 8-bits (r8b)

**%r8**

**%r8d**

**%r9**

**%r9d**

**%r10**

**%r10d**

**%r11**

**%r11d**

**%r12**

**%r12d**

**%r13**

**%r13d**

**%r14**

**%r14d**

**%r15**

**%r15d**

# Typical Register Uses

- EAX: accumulator
- EBX : Pointer to data
- ECX: Counter for string operations and loops
- EDX: I/O Operations
- EDI: Destination for string operations
- ESP: Stack pointer
- EBP: Frame pointer

register encoding	high 8-bit		low 8-bit	16-bit	32-bit
0		AH (4)	AL	AX	EAX
3		BH (7)	BL	BX	EBX
1		CH (5)	CL	CX	ECX
2		DH (6)	DL	DX	EDX
6		SI		SI	ESI
7		DI		DI	EDI
5		BP		BP	EBP
4		SP		SP	ESP
	31	16	15	0	

	FLAGS		FLAGS	EFLAGS
	IP		IP	EIP
	31	0		

# Assembly Syntax

---

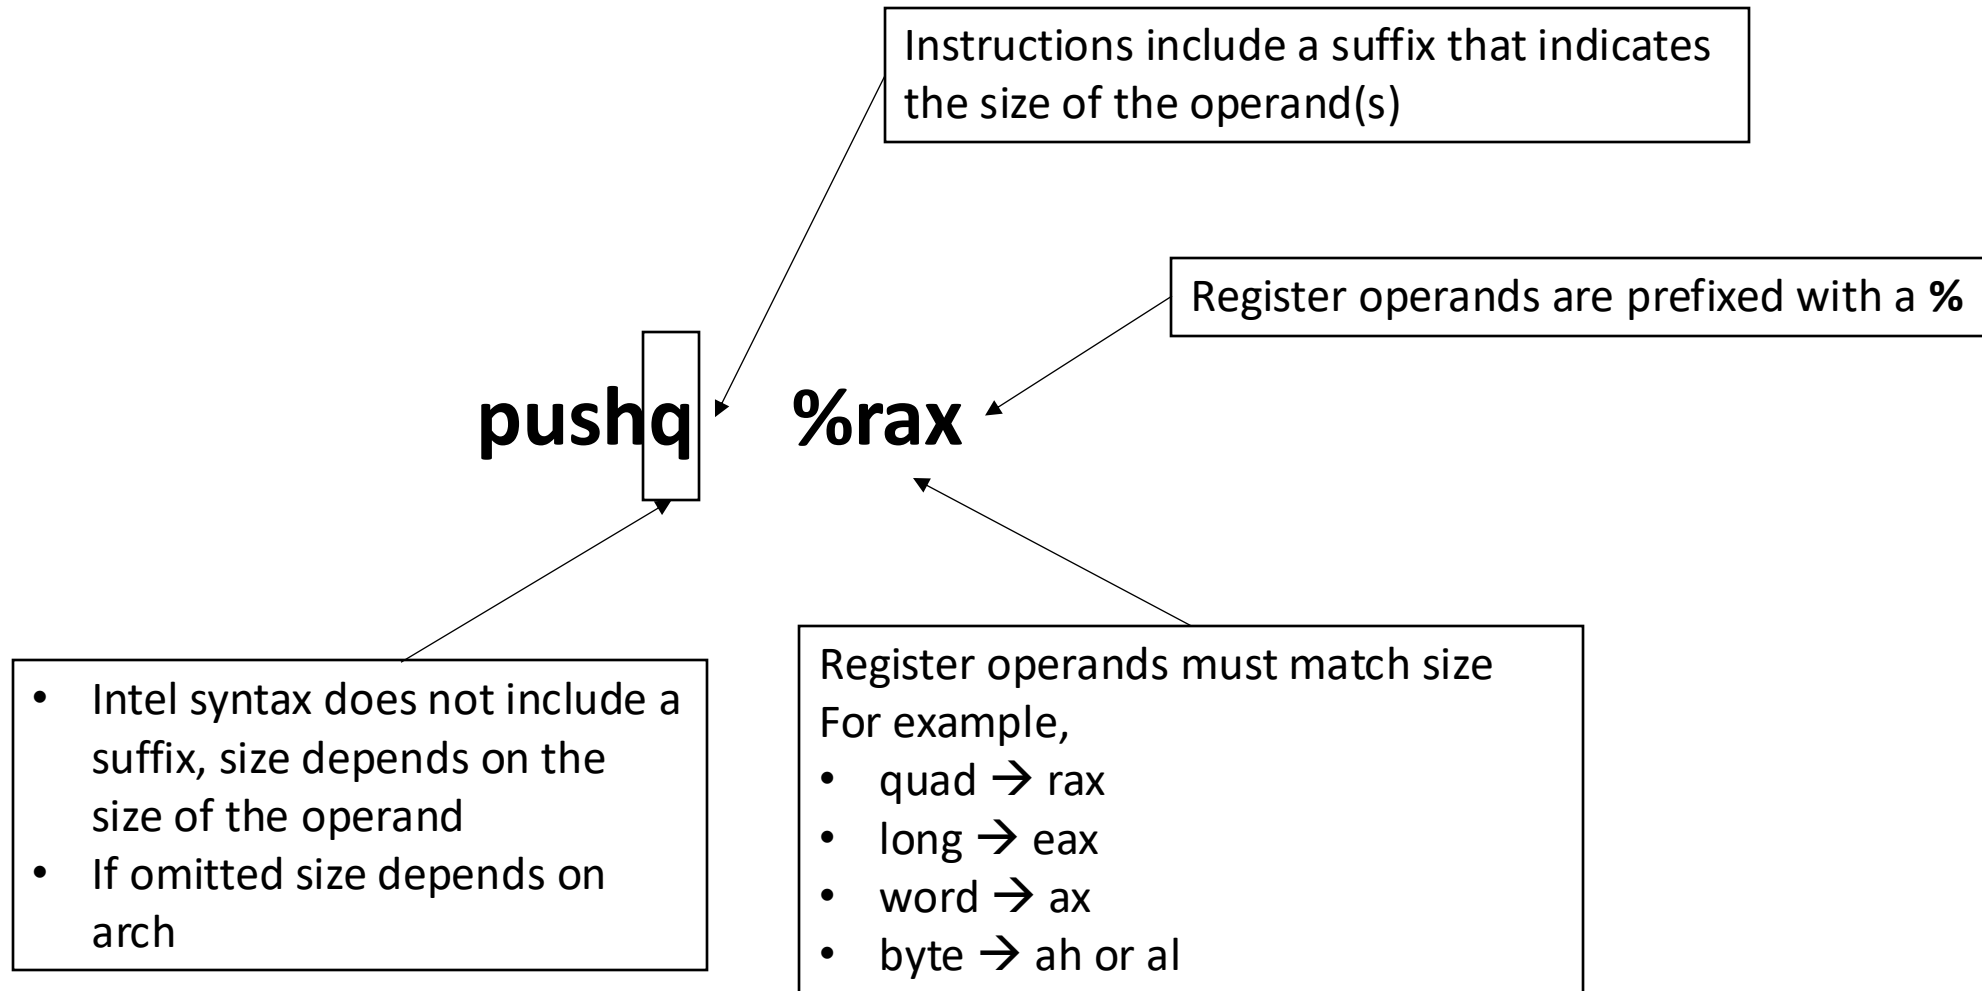
- Intel: OP dest, src
- AT&T: OP src, dest
- Unix systems prefer AT&T
  - We are going to use the same as the GNU assembler (gas syntax)

# Assembly Instructions

- **pushq**: push quad word to stack
- **movq**: Move quad word
- **imull**: Signed multiply long
- **addl**: Add long

```
pushq  %rbp
movq   %rsp, %rbp
movl   %edi, -20(%rbp)
movl   %esi, -24(%rbp)
movl   %edx, -28(%rbp)
movl   -20(%rbp), %eax
imull  -28(%rbp), %eax
movl   %eax, %edx
movl   -24(%rbp), %eax
addl   %edx, %eax
imull  -28(%rbp), %eax
```

# Operand Sizes



# Memory Operands

- Parentheses indicate a memory operand
- Each memory address can be defined as:  
**Base+Index\*Scale+Disp**
- In AT&T syntax: disp(base, index, scale)
  - disp, index, and scale are optional

```
pushq %rbp
movq  %rsp, %rbp
movl  %edi, -20(%rbp)
movl  %esi, -24(%rbp)
movl  %edx, -28(%rbp)
movl  -20(%rbp), %eax
imull -28(%rbp), %eax
movl  %eax, %edx
movl  -24(%rbp), %eax
addl  %edx, %eax
imull -28(%rbp), %eax
```





## Part 2. Smashing the Stack

---

# Buffer Overflows

- Writing outside the boundaries of a buffer
  - Buffers are arrays of bytes, integers, structs, etc.
  - Spatial violation
- Common programmer errors that lead to it ...
  - Insufficient input checks/wrong assumptions about input
  - Unchecked buffer size
  - Integer overflows



# Example

---

```
char buf[16];  
  
strcpy(buf, str);  
  
printf("%s\n", buf);  
  
return strlen(buf);
```

# BO Variations

---

```
int mytest(char *str)
{
    char buf[16];

    strcpy(buf, str);

    printf("%s\n", buf);

    return 0;
}
```

```
int mytest(char *str)
{
    char *buf = malloc(16);

    strcpy(buf, str);

    printf("%s\n", buf);

    return 0;
}
```

```
char buf[16];

int mytest(char *str)
{
    strcpy(buf, str);

    printf("%s\n", buf);

    return 0;
}
```

# BO Variations

## Stack Buffer Overflow

```
int mytest(char *str)
{
    char buf[16];

    strcpy(buf, str);

    printf("%s\n", buf);

    return 0;
}
```

## Heap Buffer Overflow

```
int mytest(char *str)
{
    char *buf = malloc(16);

    strcpy(buf, str);

    printf("%s\n", buf);

    return 0;
}
```

## Global Buffer Overflow

```
char buf[16];

int mytest(char *str)
{
    strcpy(buf, str);

    printf("%s\n", buf);

    return 0;
}
```

# Buffer Overflows

---

- Can happen when calling common string and buffer functions
  - strcpy(), strcat(), memcpy(), memset(), memmove(), etc.
- But not limited to those functions
  - Can also happen with functions as read(), fread(), gets(), fgets(), etc.

# Buffer Overflows

- Can happen when calling common string and buffer functions
  - strcpy(), strcat(), memcpy(), memset(), memmove(), etc.
- But not limited to those functions
  - Can also happen with functions as read(), fread(), gets(), fgets(), etc.
- Custom data copying code can also suffer

```
static void my_copy(char *dst, const char *src)
{
    char *dp, *sp;
    for (sp = src, dp = dst; *sp != '\0'; sp++, dp++) {
        *dp = *sp;
    }
    *dp = '\0';
}
```


# How Do Function Calls Work

---




# Stack Data Structure


- Stack plays a crucial role in supporting functions
  - Follows last-in first-out semantic



# Stack Operation Push




# Stack Operation Push




# Stack Operation Push


---




# Stack Operation Pop



# Stack Operation Pop




# Stack Operation Pop



# The Stack Pointer (SP)

The stack pointer points to the first element in the stack (the top).






# The Stack Pointer (SP)

The stack pointer points to the first element in the stack (the top).

Usually the RSP/ESP register is used to store the SP.




# The Stack Pointer (SP)

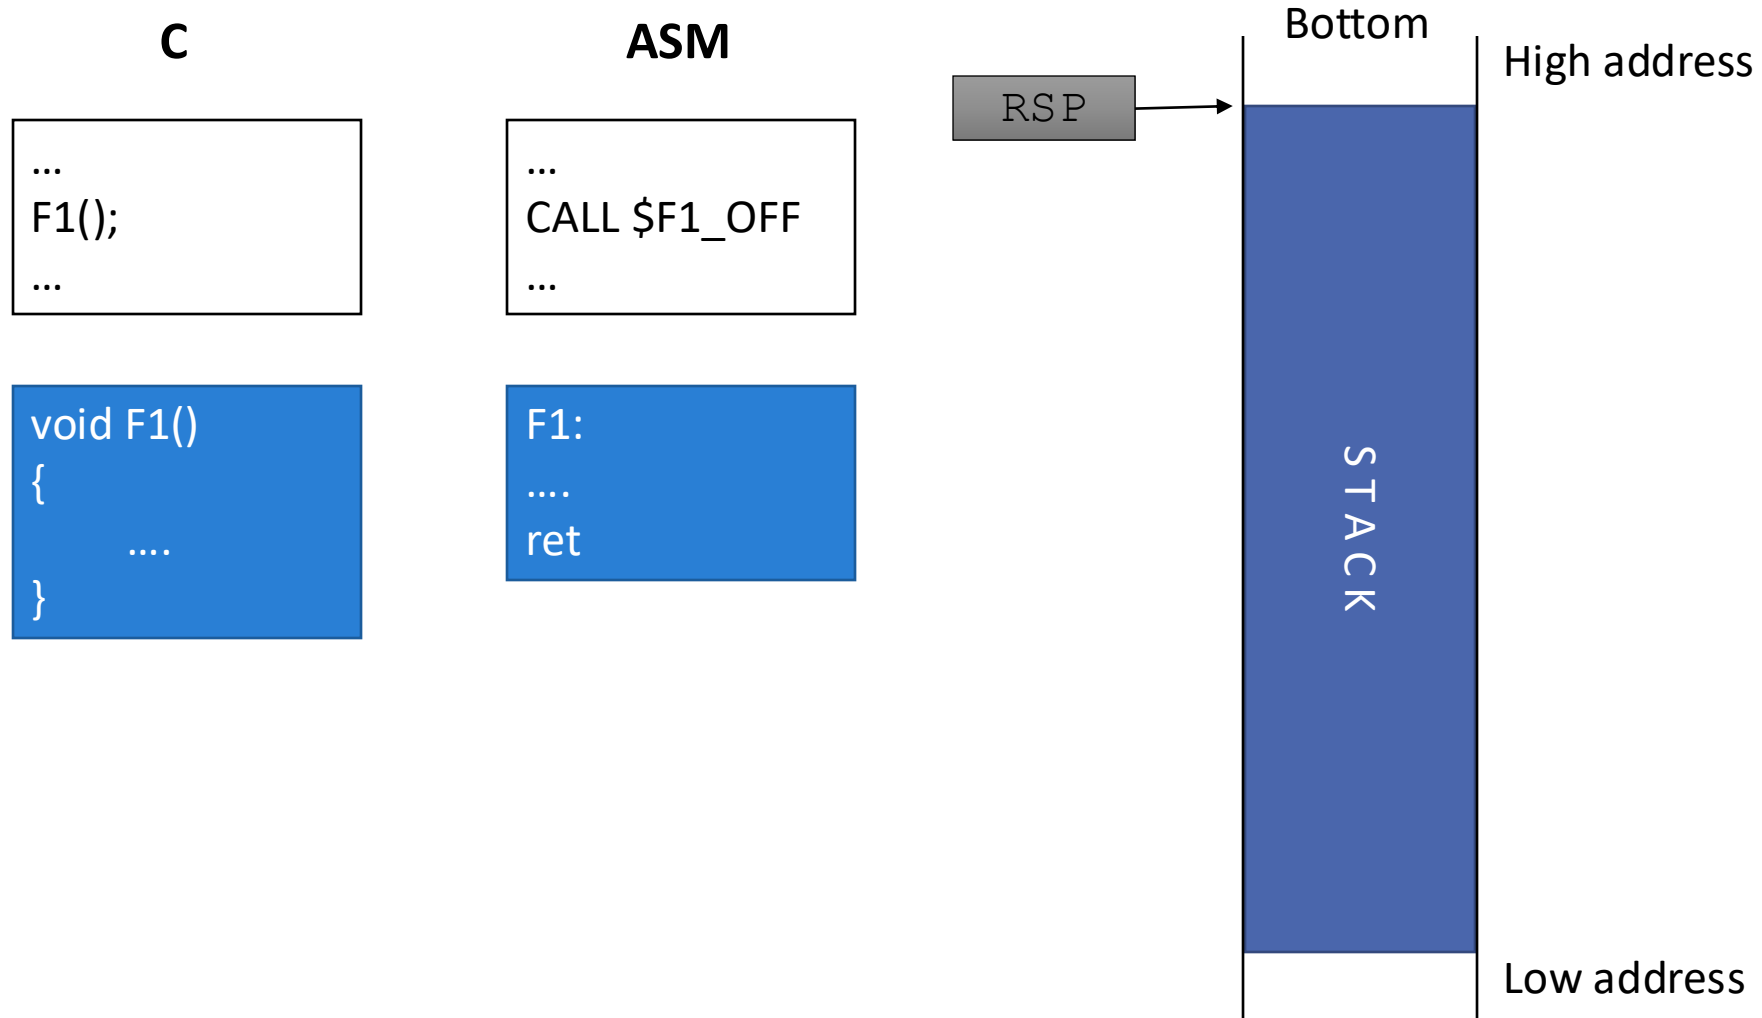
The stack pointer points to the first element in the stack (the top).

Usually the RSP/ESP register is used to store the SP.

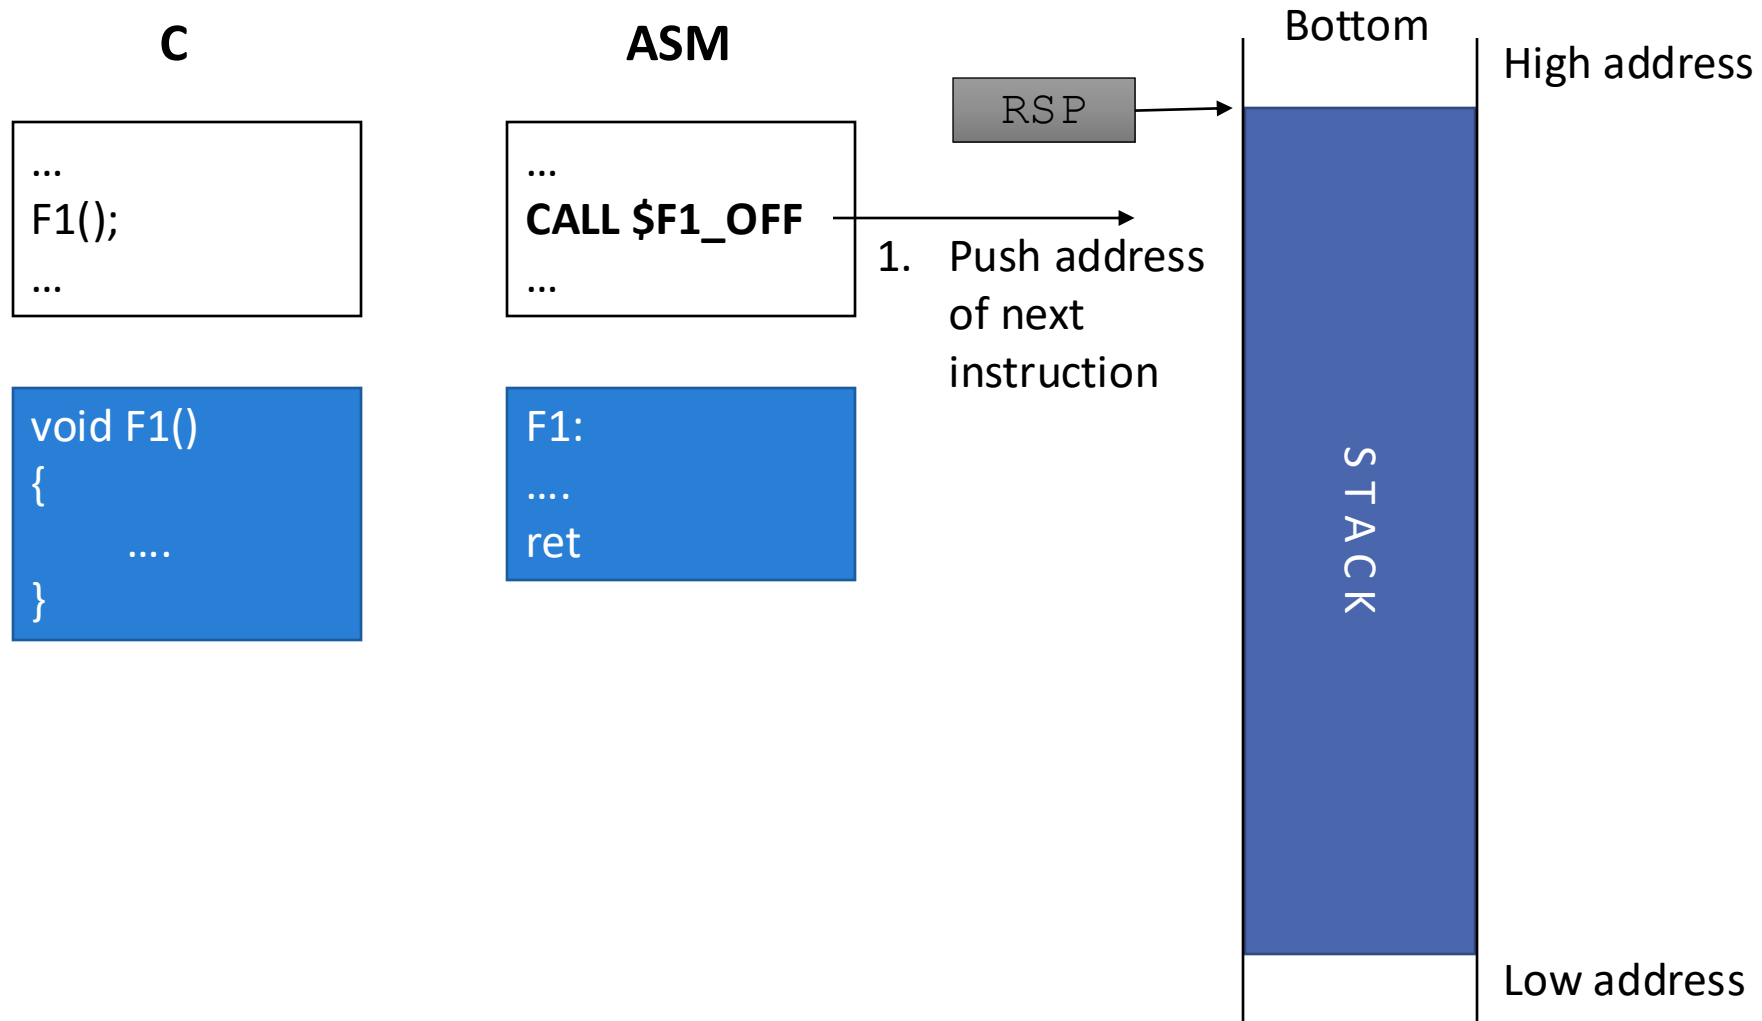
The stack grows towards lower addresses



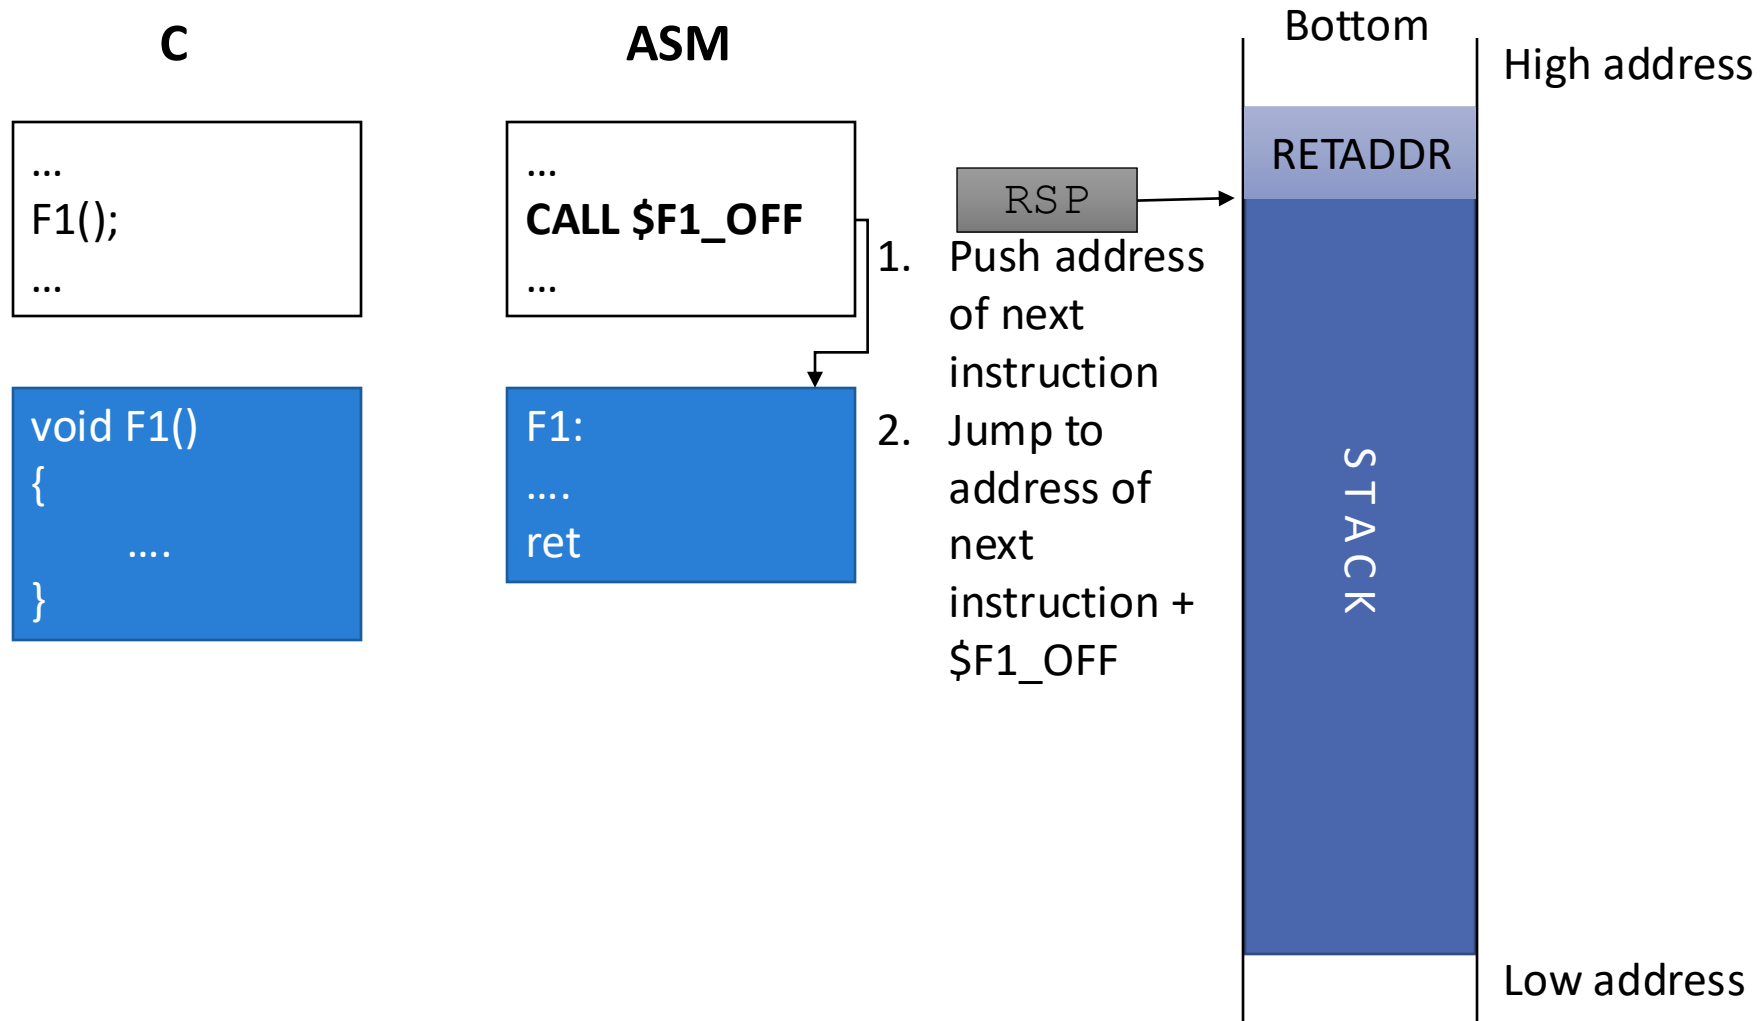
# Simple Function Call



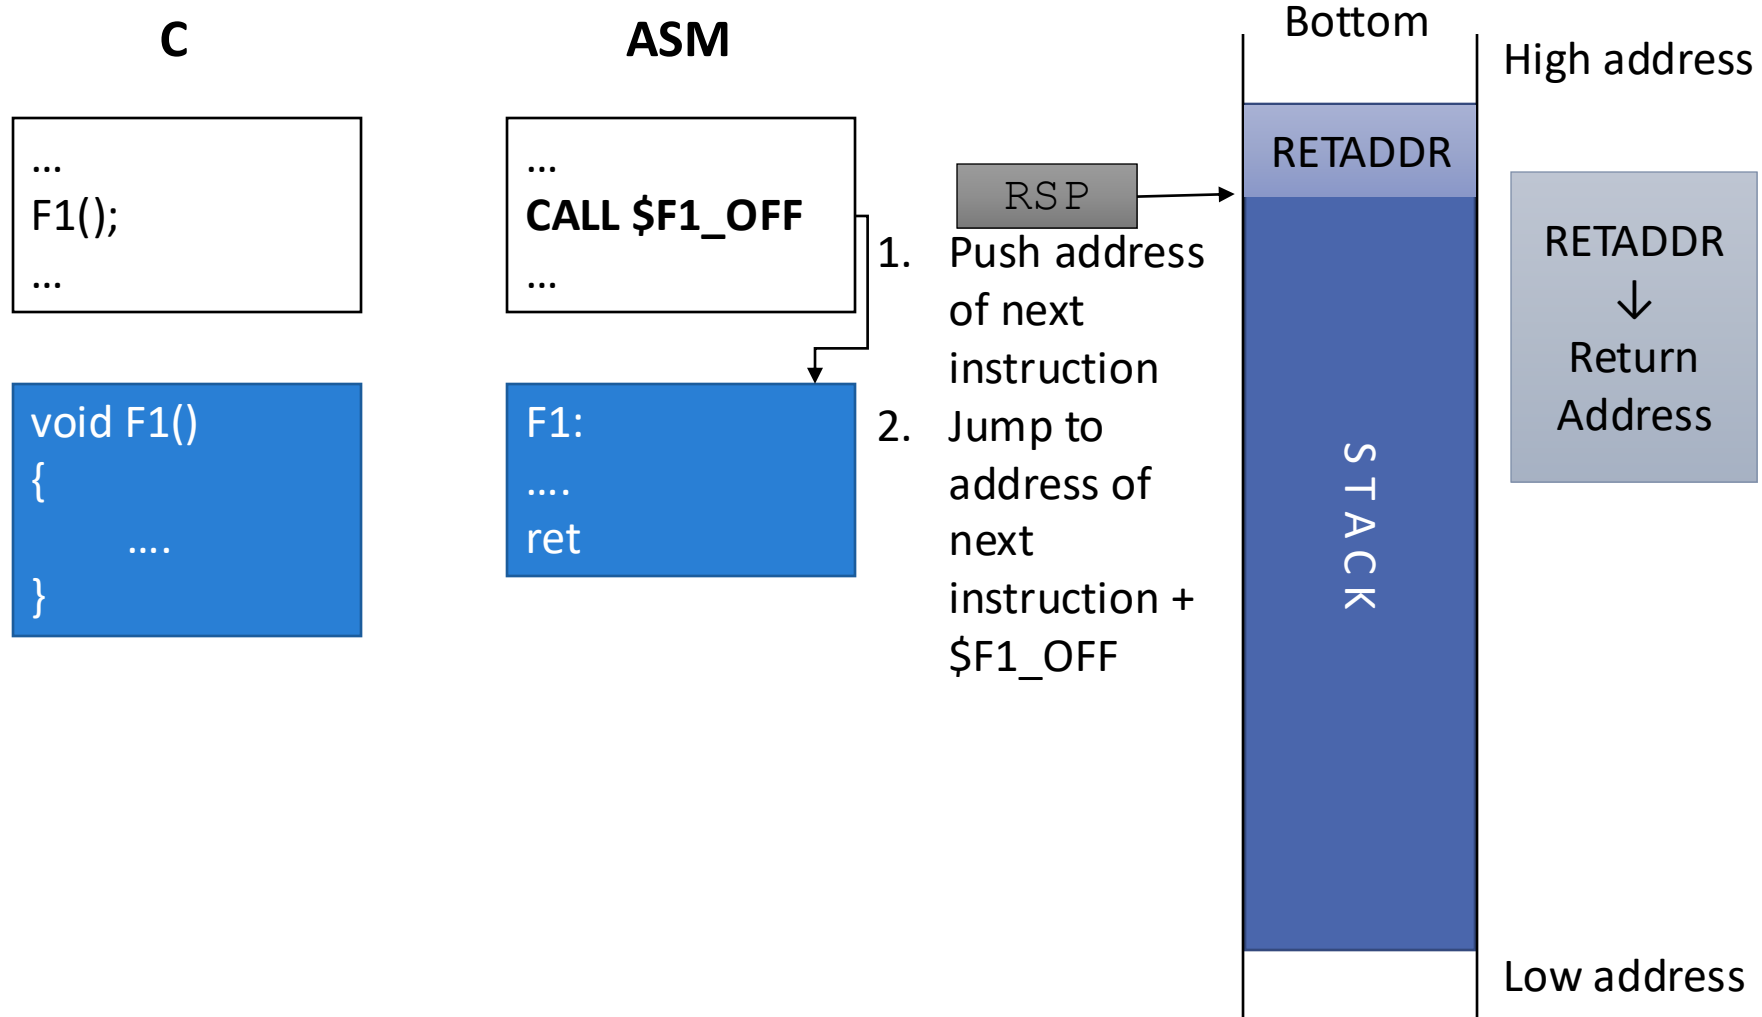
# Simple Function Call



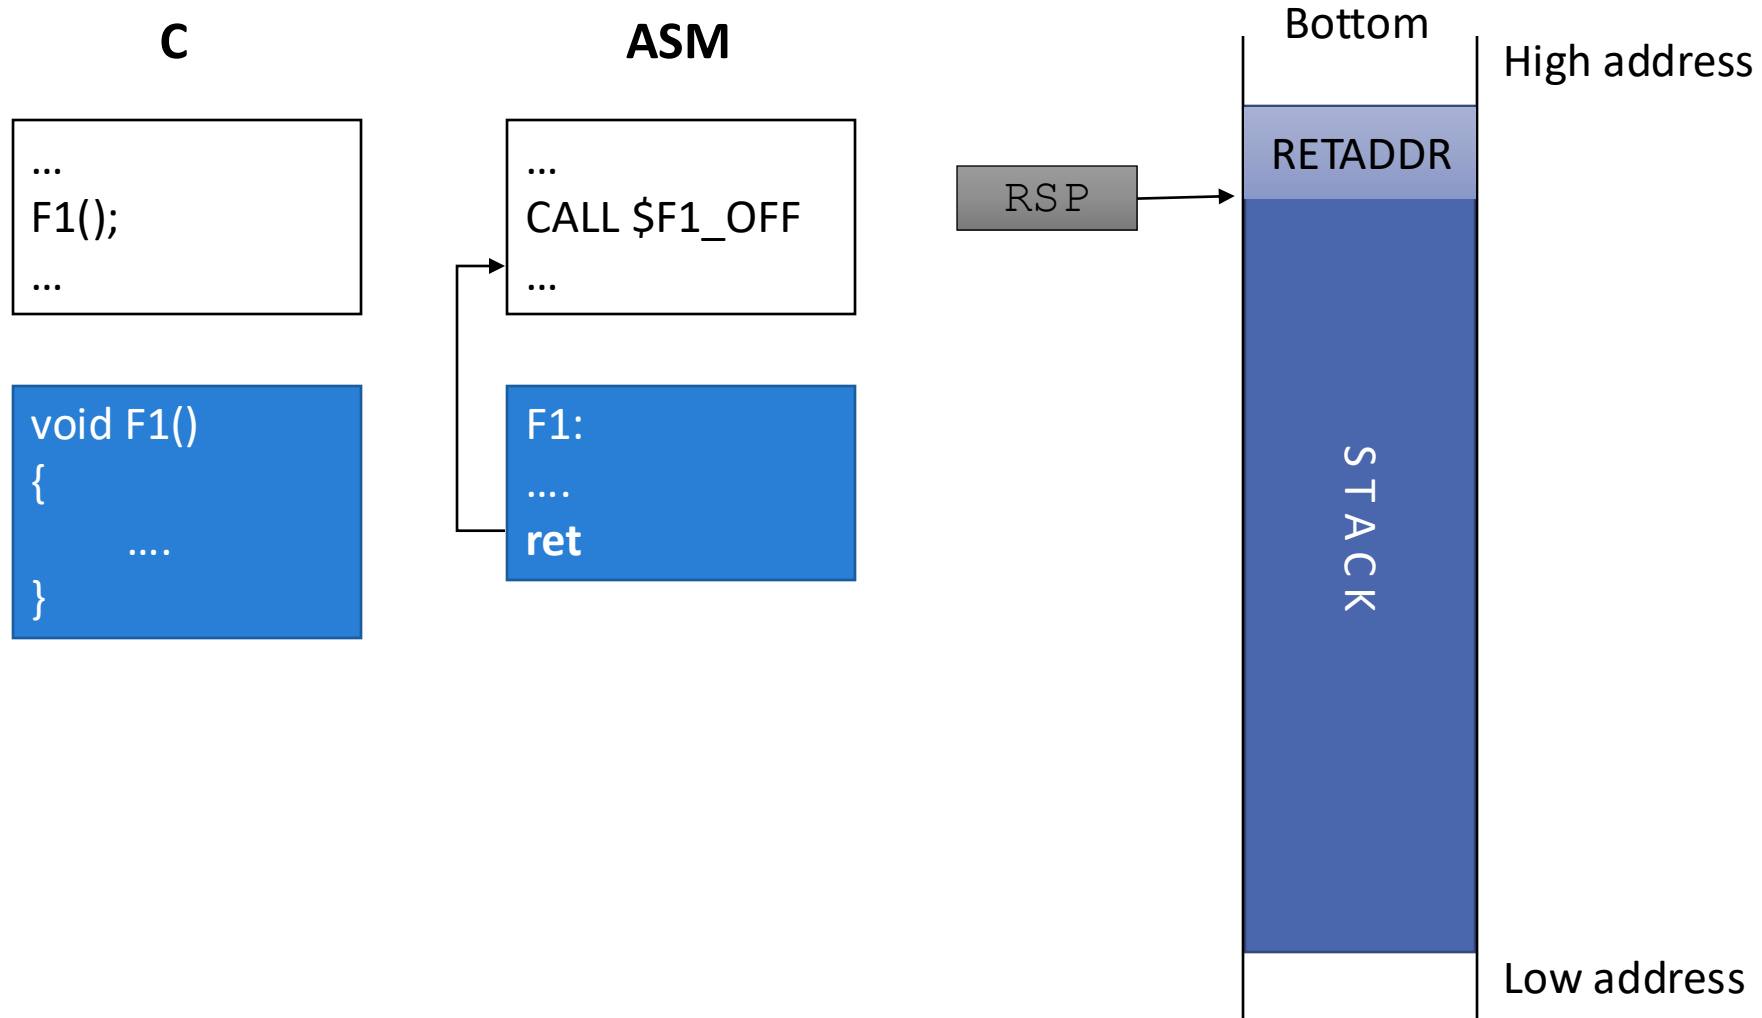
# Simple Function Call



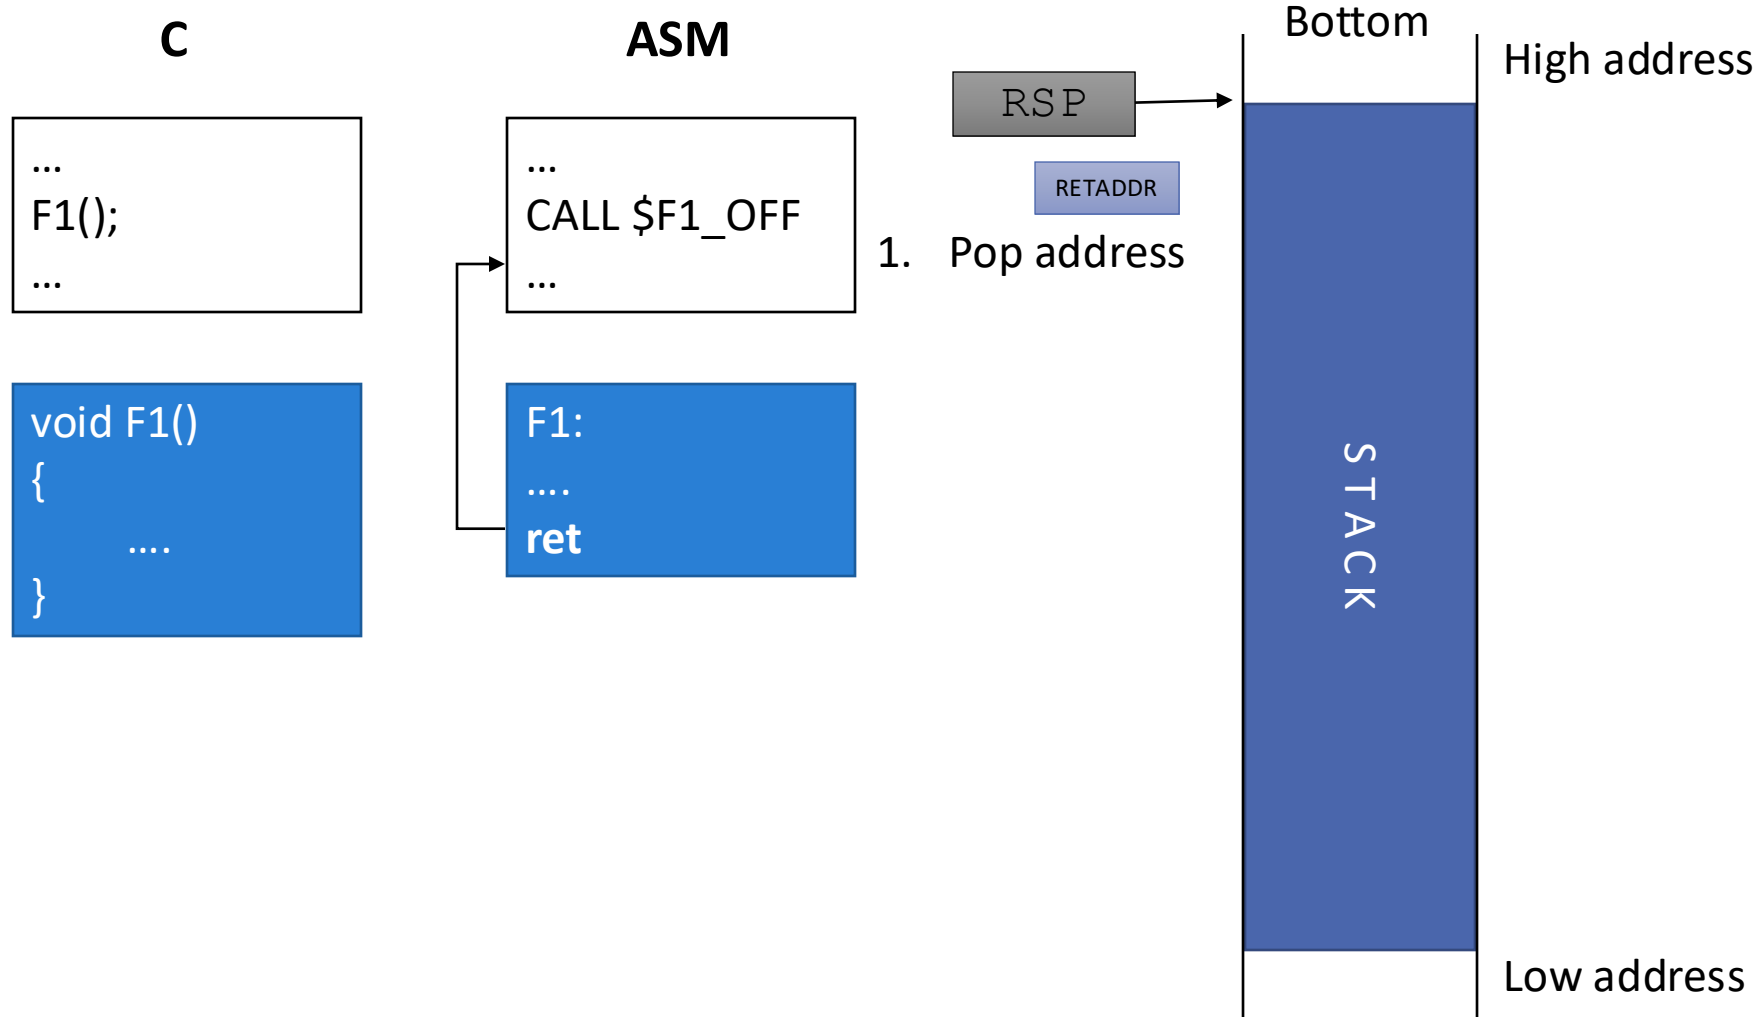
# Simple Function Call



# Simple Function Return

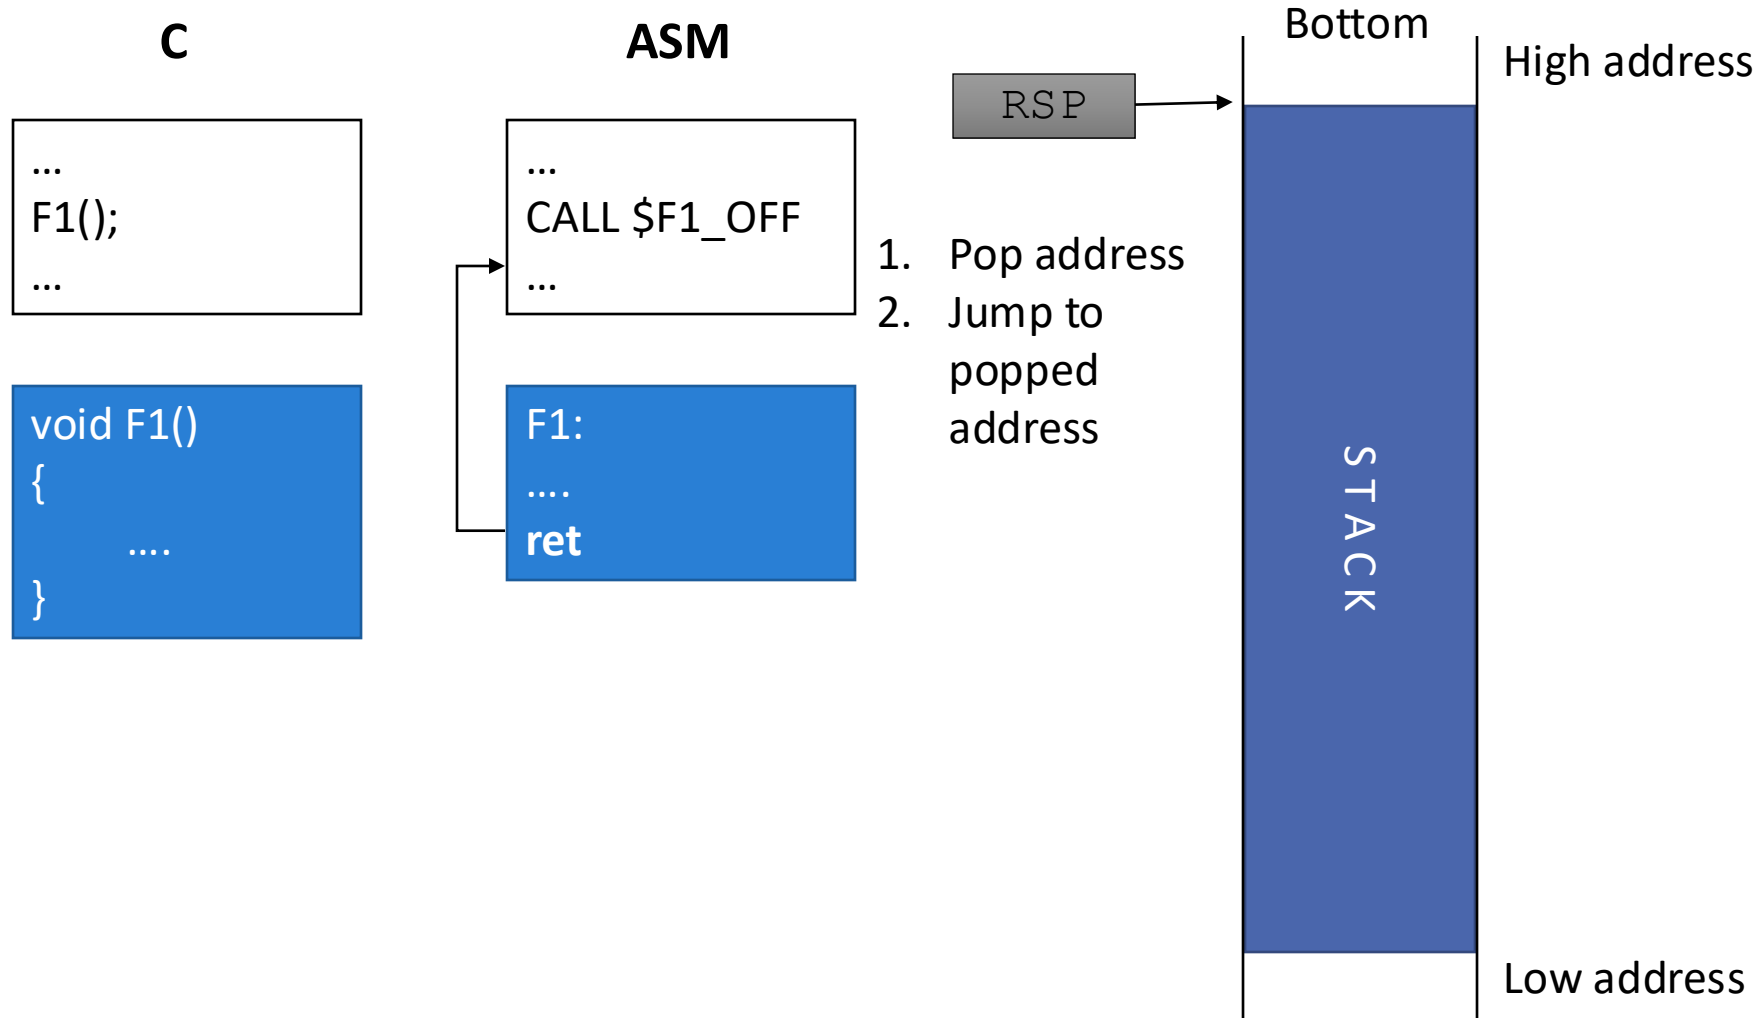


# Simple Function Return





# Simple Function Return



# Function Calls and Returns

## Calling a function (the callee)

- CALL instruction
  - Pushes `next_ins_addr` on stack and transfers control to address described by operand
- Most common syntax: CALL OFFSET
  - Target is `next_ins_addr + OFFSET`

## Returning to caller


- RET instruction
  - Pop return address from stack and transfers control to it

```
CALL tgt → push next_ins_addr; jmp tgt  
RET → pop retaddr; jmp retaddr
```

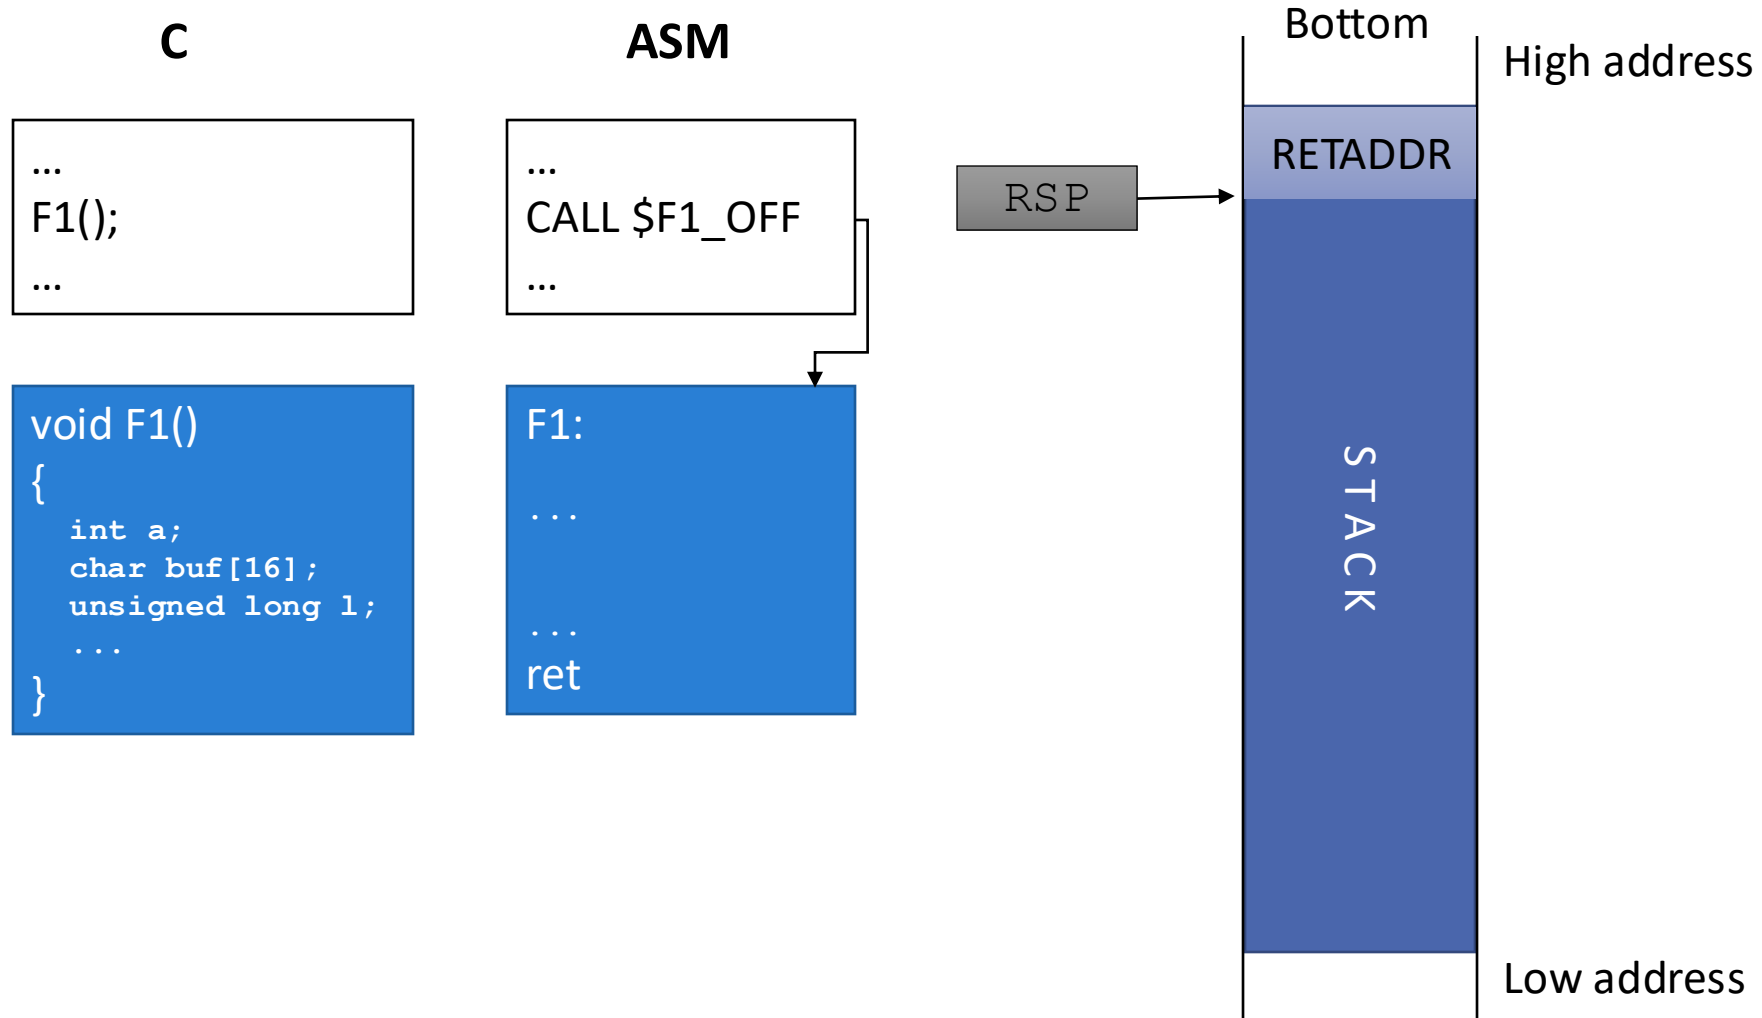
```
call and ret implicitly use the SP
```

# The Stack Is Used...

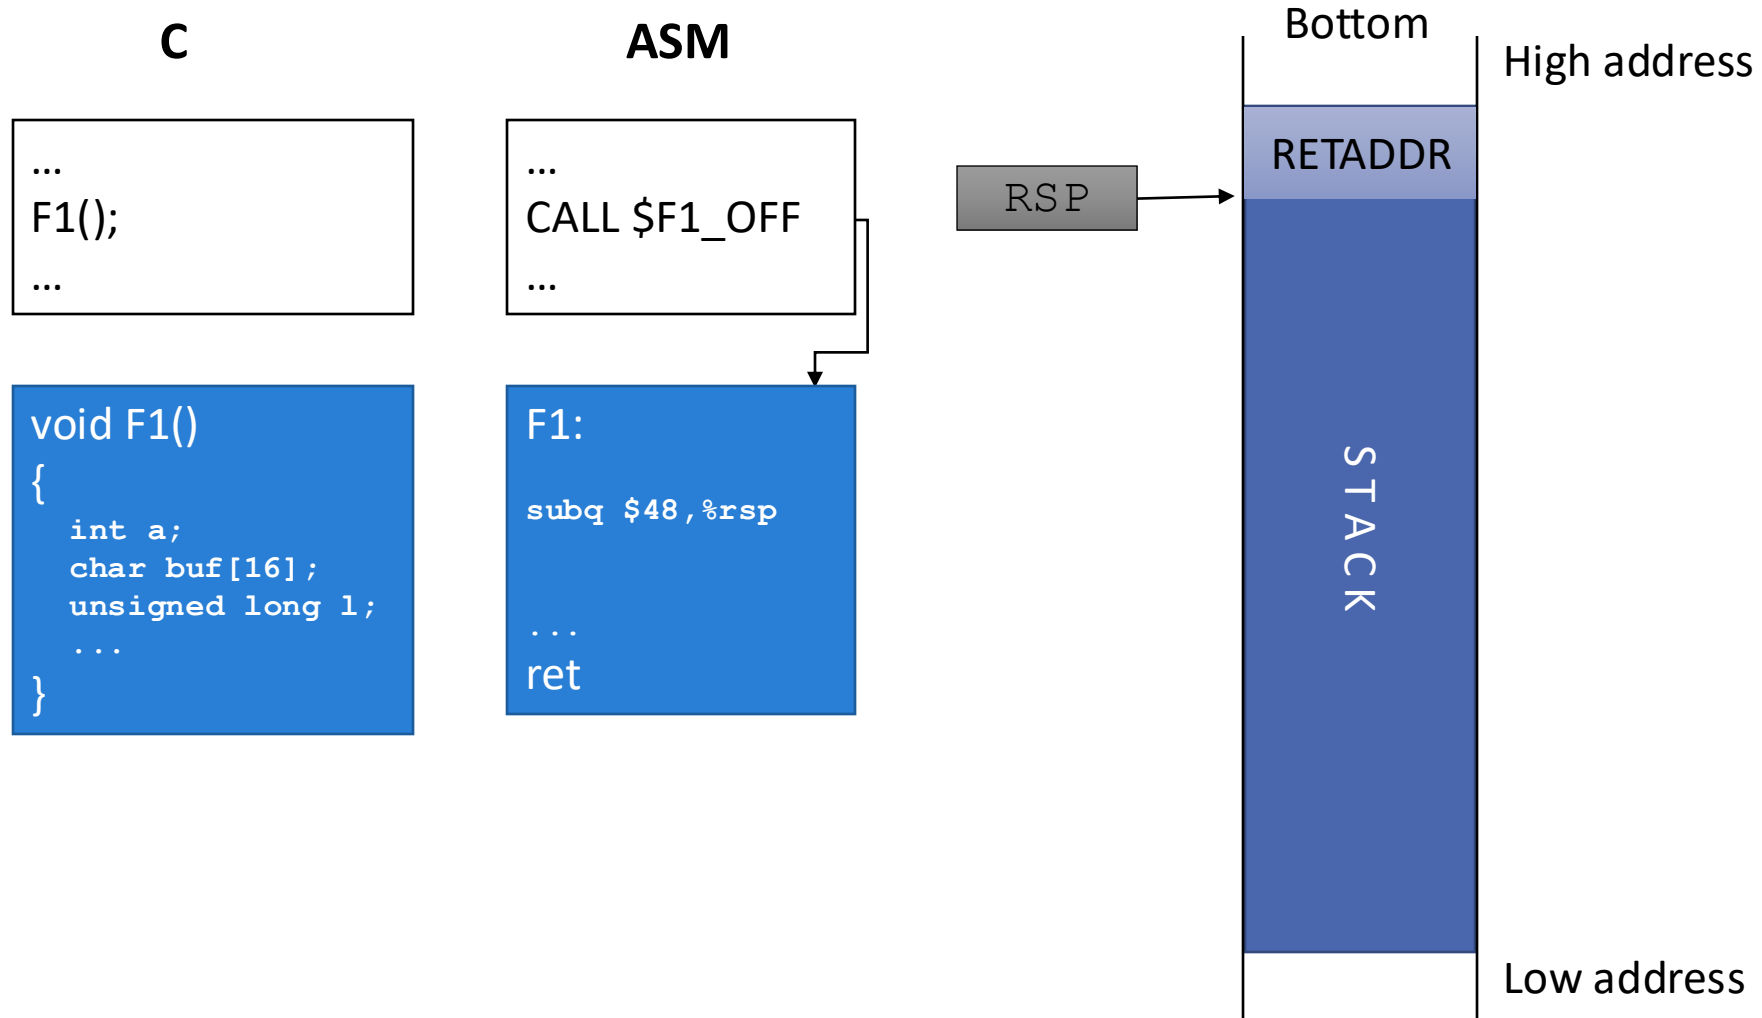
- ...to store the return address of caller functions
  - Code pointers!
- ...to store local variables
  - Aka stack variables



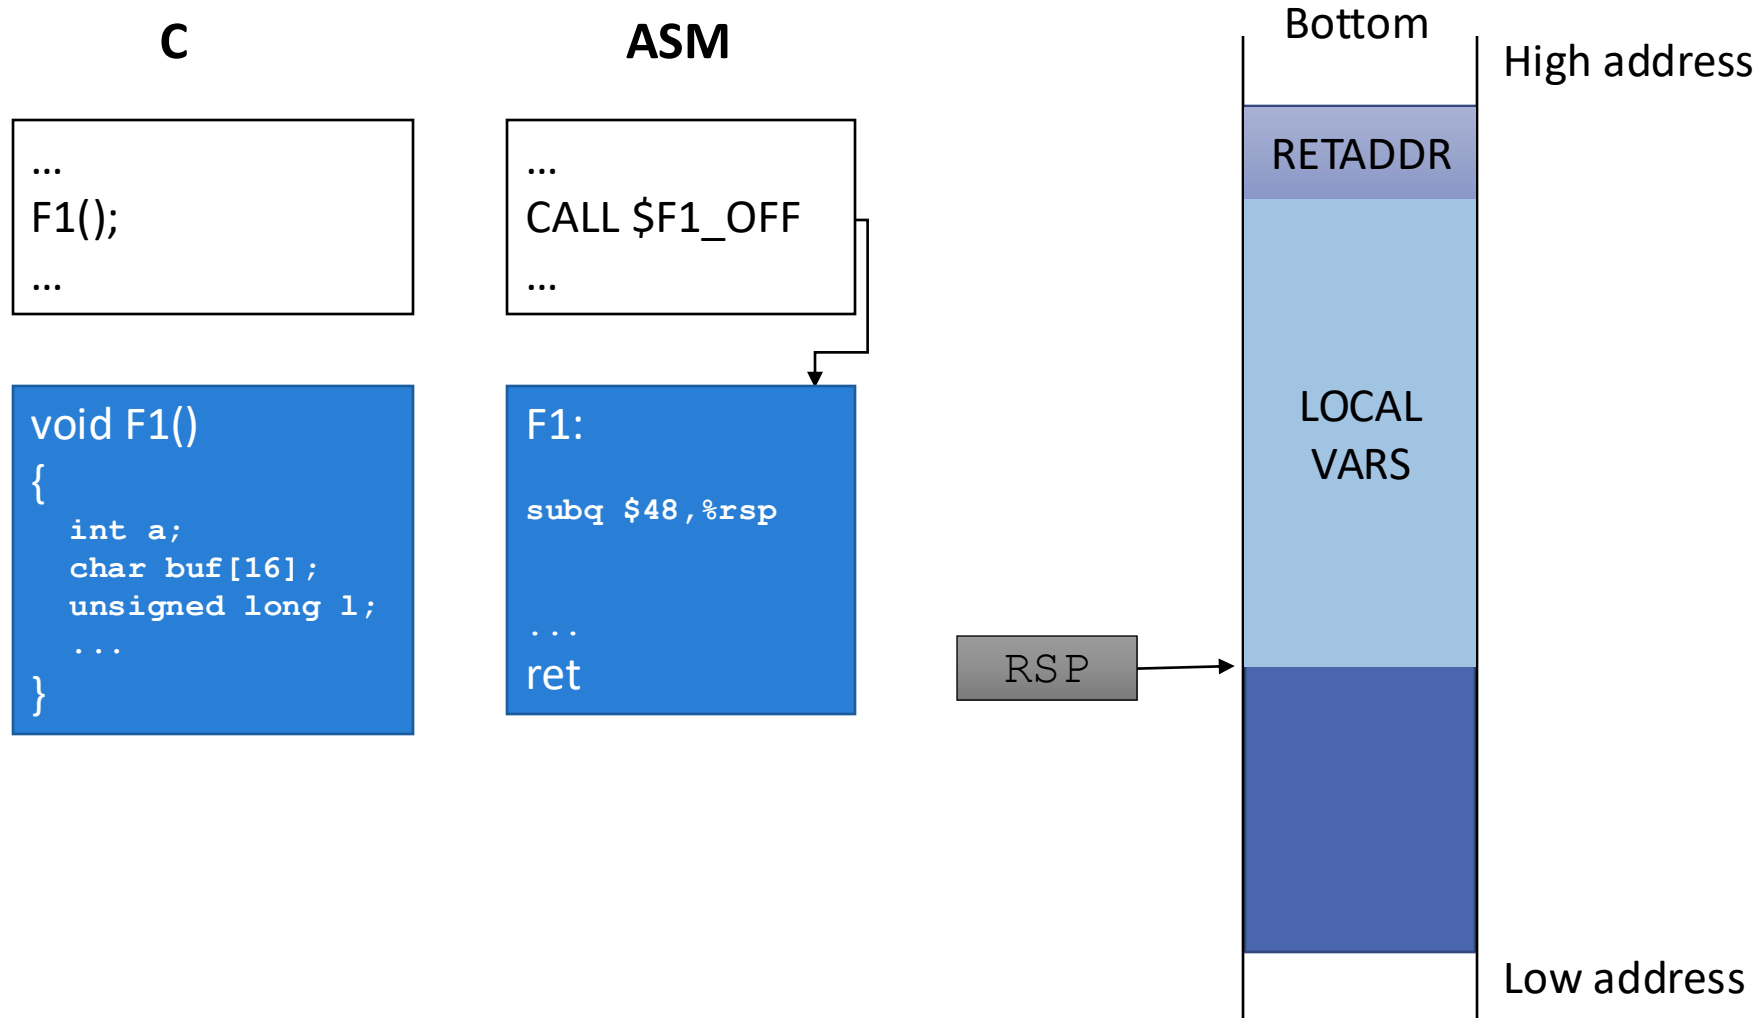
# Local Variables



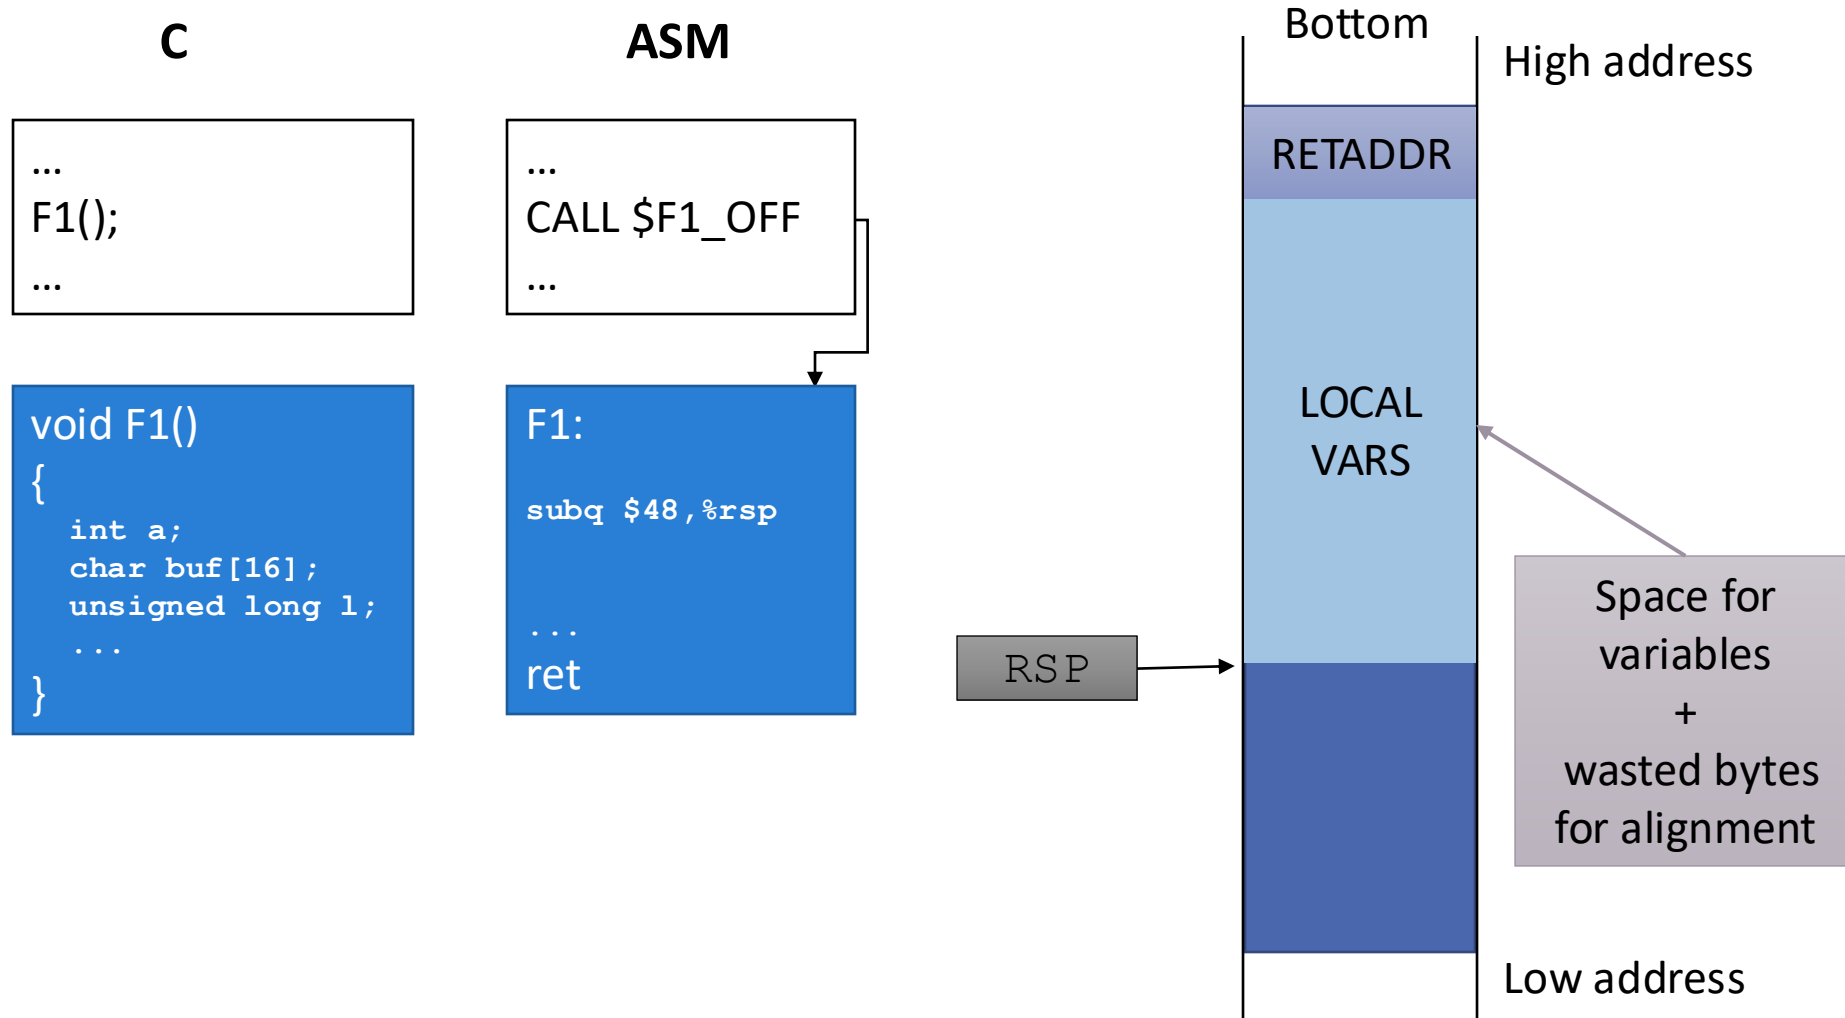
# Local Variables



# Local Variables



# Local Variables



# Alignment


- CPUs like aligned data
  - Better performance
- Compilers try to align data





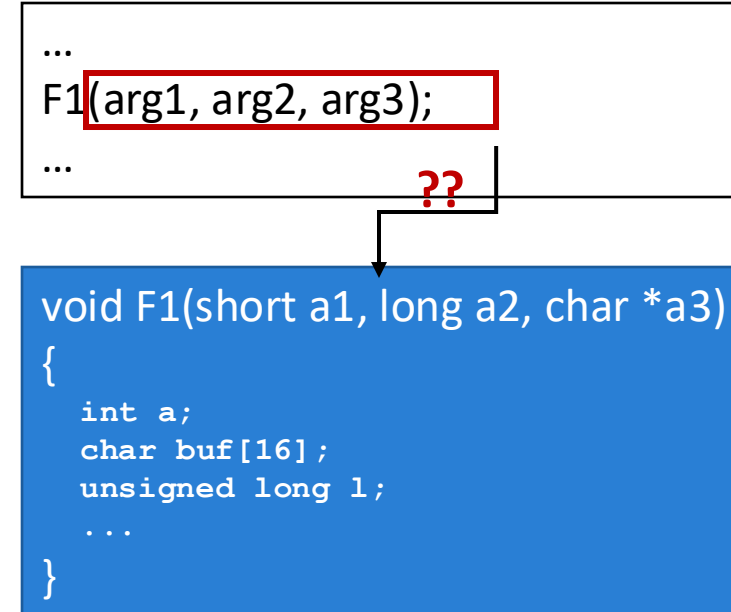
# The Stack Is Used...

- ...to store the return address of caller functions
  - Code pointers!
- ...to store local variables
  - Aka stack variables
- ...to pass function arguments
  - Mostly on 32-bit architectures



# Calling Conventions

- Defines the standard for passing arguments
- Caller and callee need to agree on the standard
- Enforced by compiler
- Important when using 3rd party libraries
  - Hence, also referred to as the Application Binary Interface (ABI)
- Different styles  $\leftrightarrow$  different advantages




## Popular conventions:

- cdecl (32-bit)
- System V AMD64 ABI

# System V AMD64 ABI


- Arguments are passed using registers
  - First 6 integer or pointer arguments are passed in registers RDI, RSI, RDX, RCX, R8, and R9

```
...  
F1(0xff, UINT_MAX, argv[0]);  
...  
  
...  
movq    (%rsi), %rdx  
movl    $4294967295, %esi  
movl    $255, %edi  
call    F1  
...
```



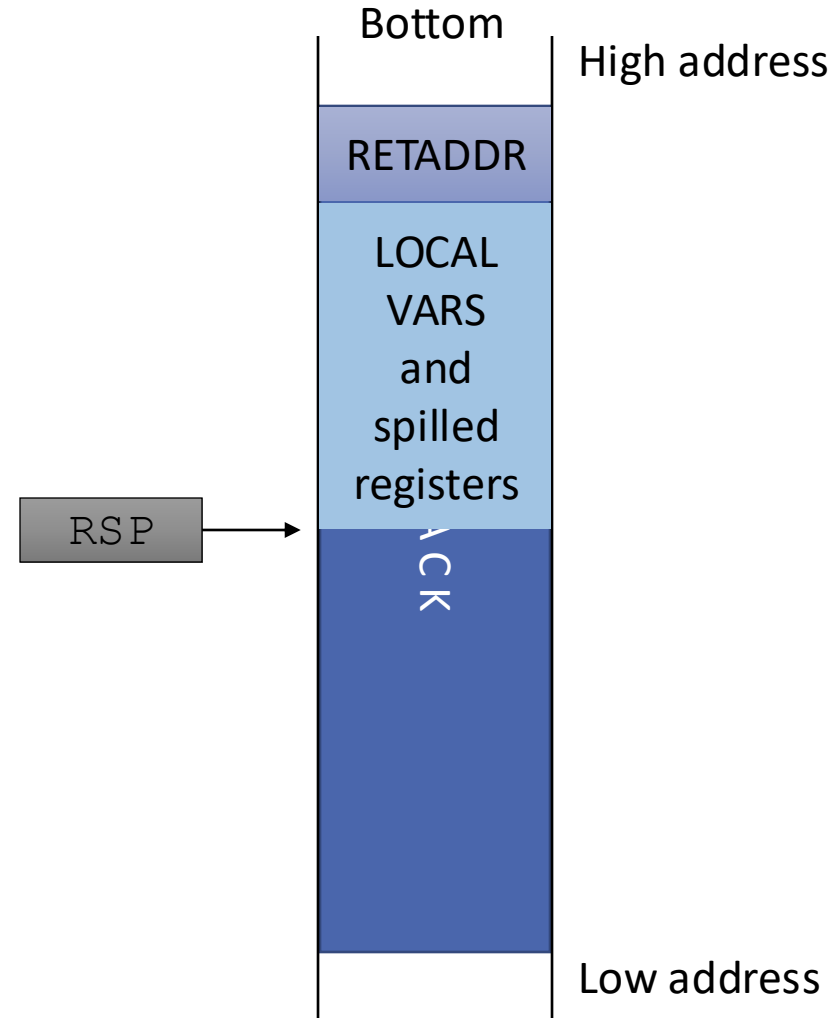
# System V AMD64 ABI

- Arguments are passed using registers
  - First 6 integer or pointer arguments are passed in registers RDI, RSI, RDX, RCX, R8, and R9
- RBP, RBX, and R12–R15 are callee saved



# System V AMD64 ABI

- Arguments are passed using registers
  - First 6 integer or pointer arguments are passed in registers RDI, RSI, RDX, RCX, R8, and R9
- RBP, RBX, and R12–R15 are callee saved
- RAX used for function return



# Popular Conventions Summary

## cdecl (mostly 32-bit)

- Arguments are passed on the stack
  - Pushed right to left
- eax, edx, ecx are caller saved
  - callee can overwrite without saving
- ebx, esi, edi are callee saved
  - callee must ensure they have same value on return
- eax used for function return value

## System V AMD64 ABI


- Arguments are passed using registers
  - First 6 integer or pointer arguments are passed in registers RDI, RSI, RDX, RCX, R8, and R9
- RBP, RBX, and R12–R15 are callee saved
- RAX used for function return

Conventions include additional information, consult reading material for thorough description

- Example: handling of floating point regs

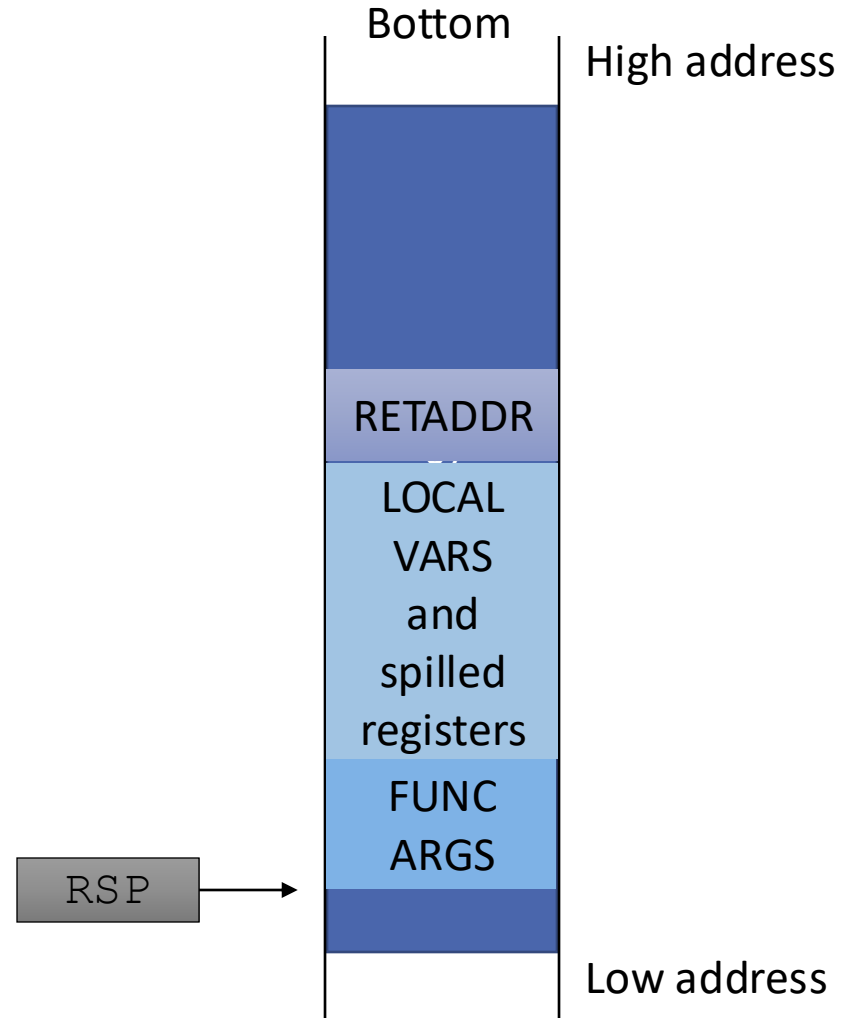
# The Stack Is Used...

- ...to store the return address of caller functions
  - Code pointers!
- ...to store local variables
  - Aka stack variables
- ...to pass function arguments
- ...to temporarily store register values
- ...to store the frame pointer



# Stack Frame

- A stack frame includes all function-local data
  - Local variables
  - Spilled registers
  - Function arguments pushed to the stack to make calls
- More of a logical entity
- Can grow as function executes

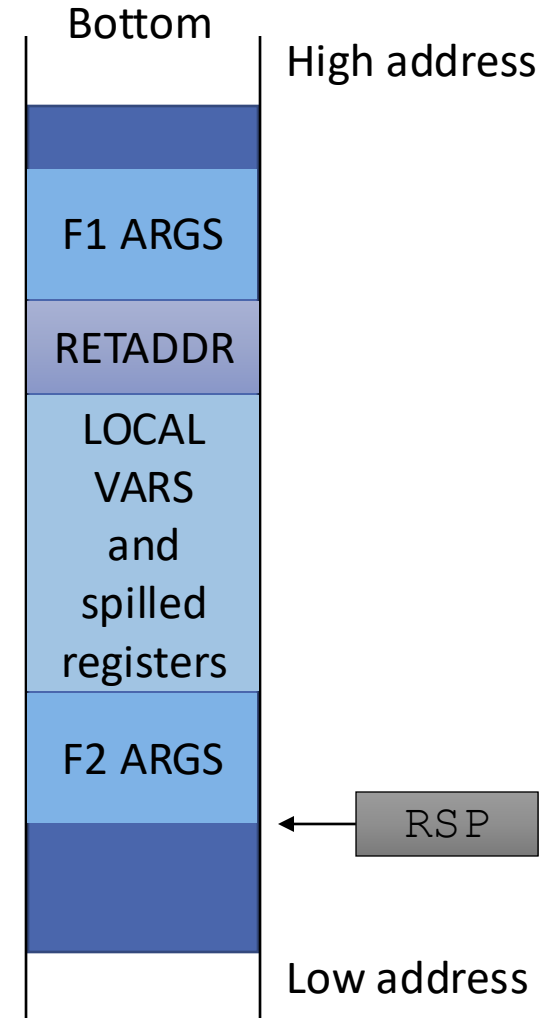




# Stack Frame Boundaries

- Start below return address
- Stop at stack pointer


```
void F1(short a1, long a2, char *a3)
{
    int a;
    char buf[16];
    unsigned long l;
    ...
    long l2 = F2(a);
    ...
}
```



# Stack Frame Boundaries


- Start below return address
- Stop at stack pointer

```
void F1(short a1, long a2, char *a3)
{
    int a;
    char buf[16];
    unsigned long l;
    ...
    long l2 = F2(a);
    ...
}
```

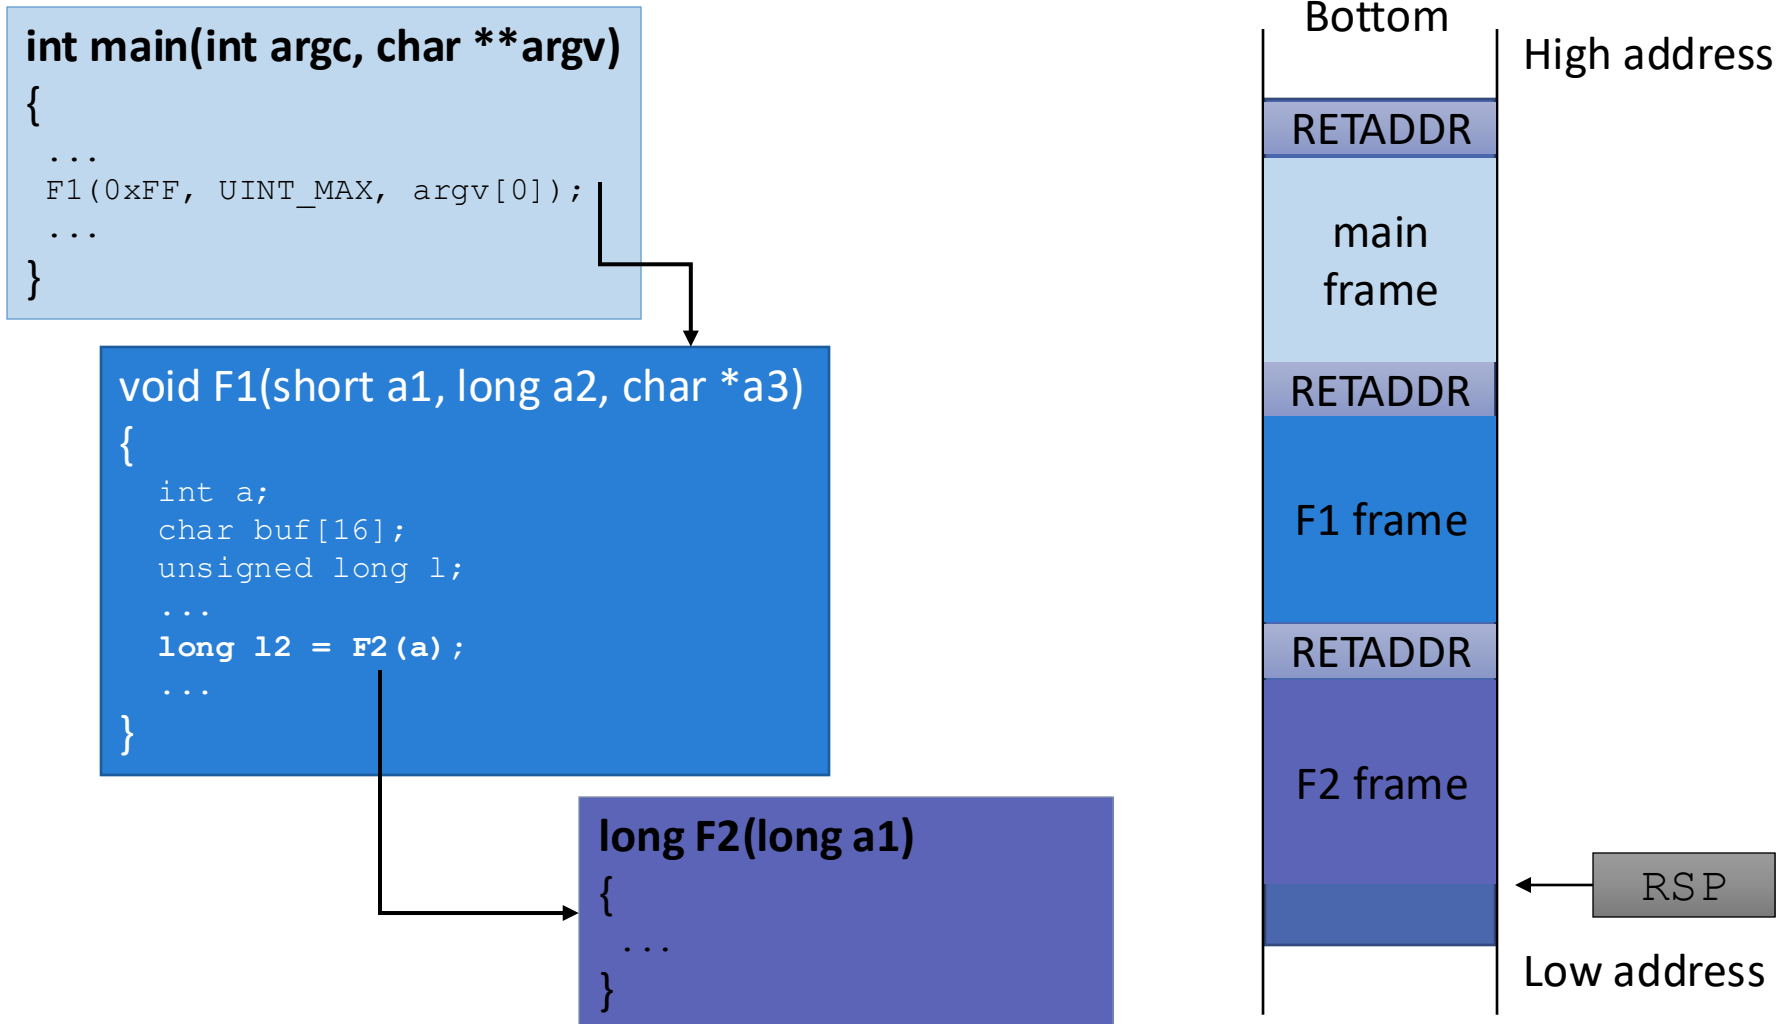


# RETADDR and Stack Frames

- The return address may also be considered part of the frame
- We will not consider it for simplicity




# Stack Frames Example




# Frame Pointer (FP)

- Marks the highest address in the frame
  - Bottom of the frame
- Aka Base Pointer

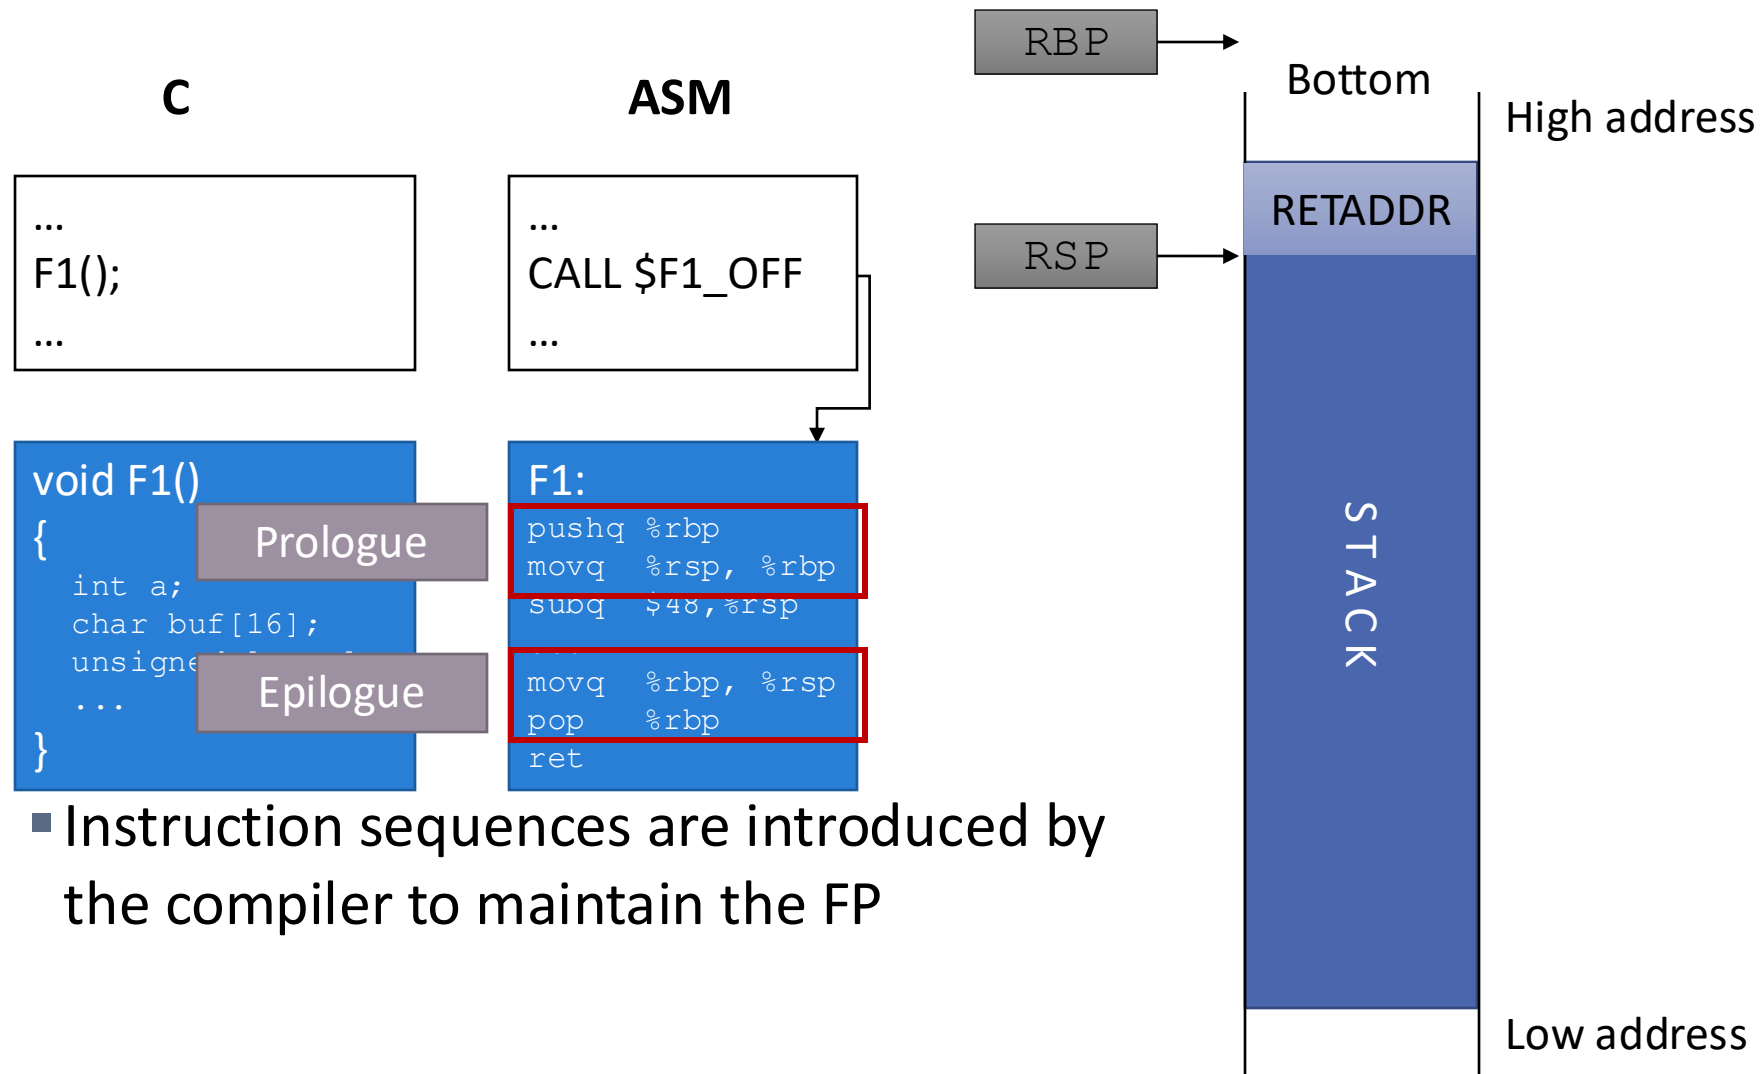


# Frame Pointer (FP)

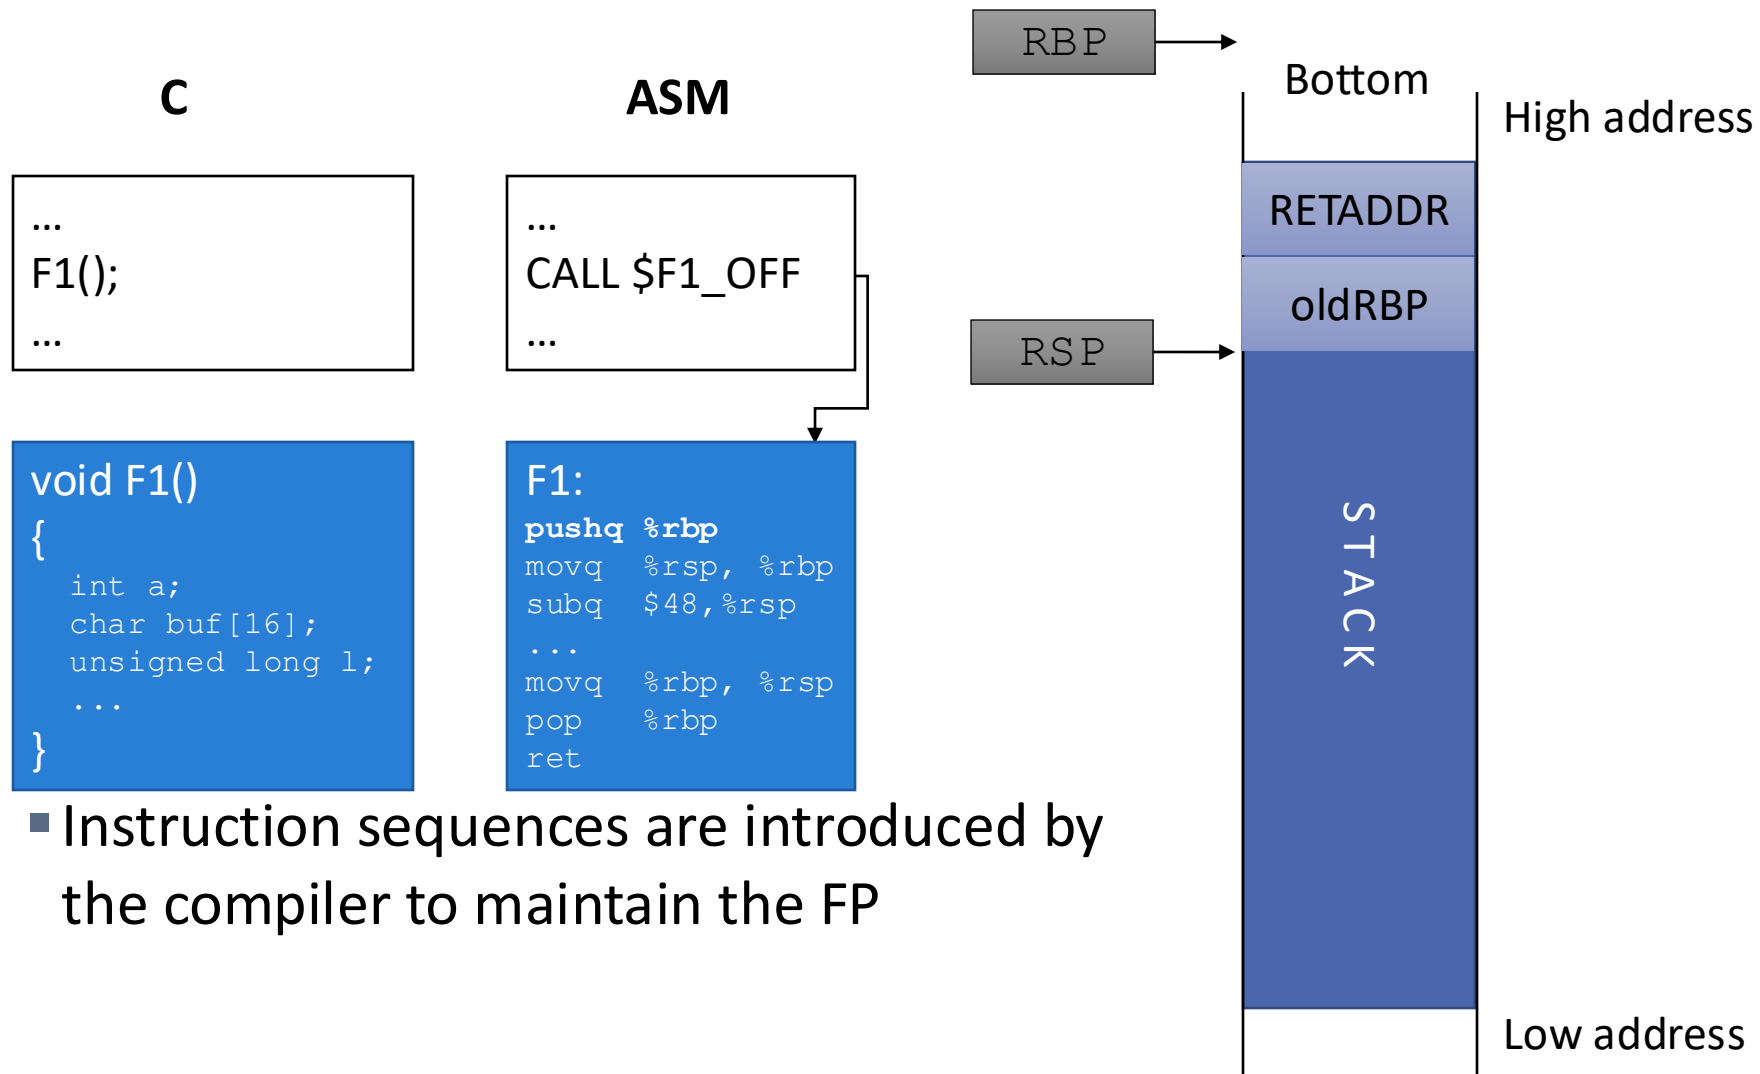
- Marks the highest address in the frame
  - Bottom of the frame
- Aka Base Pointer
- The RBP/EBP register commonly contains the FP
- RBP needs to be updated upon entry/exit of function



# Maintaining the FP



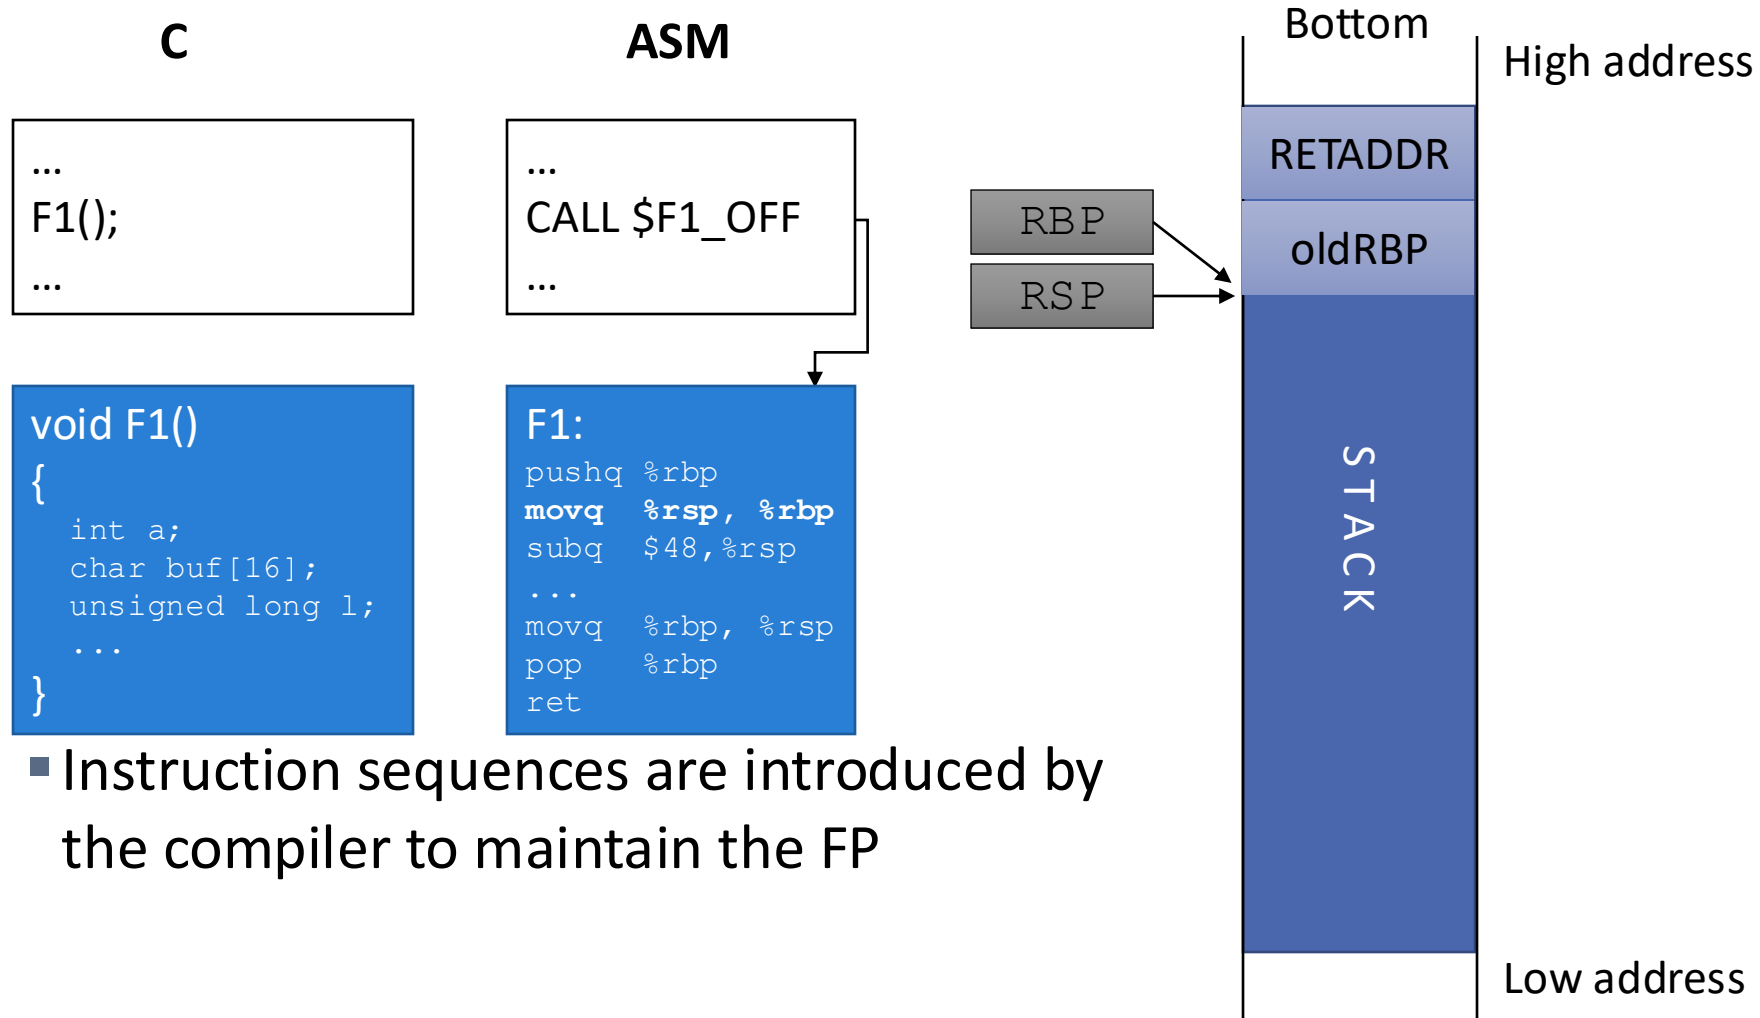
# Maintaining the FP



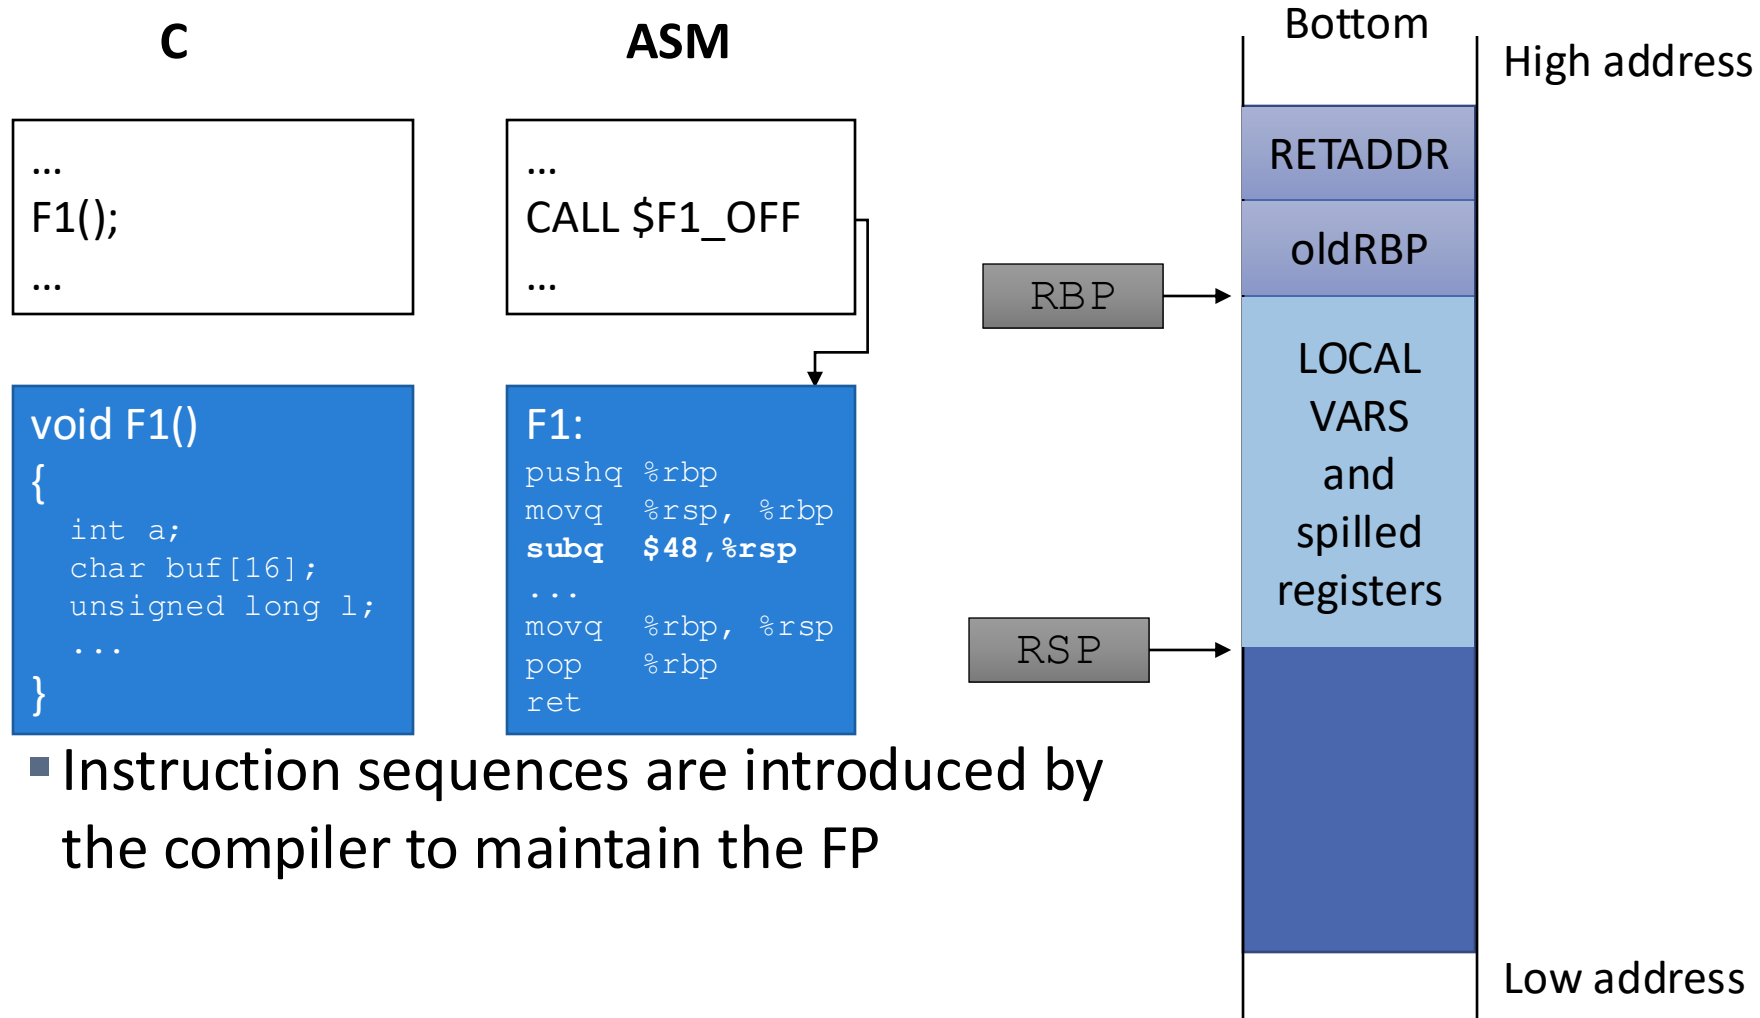
- Instruction sequences are introduced by the compiler to maintain the FP



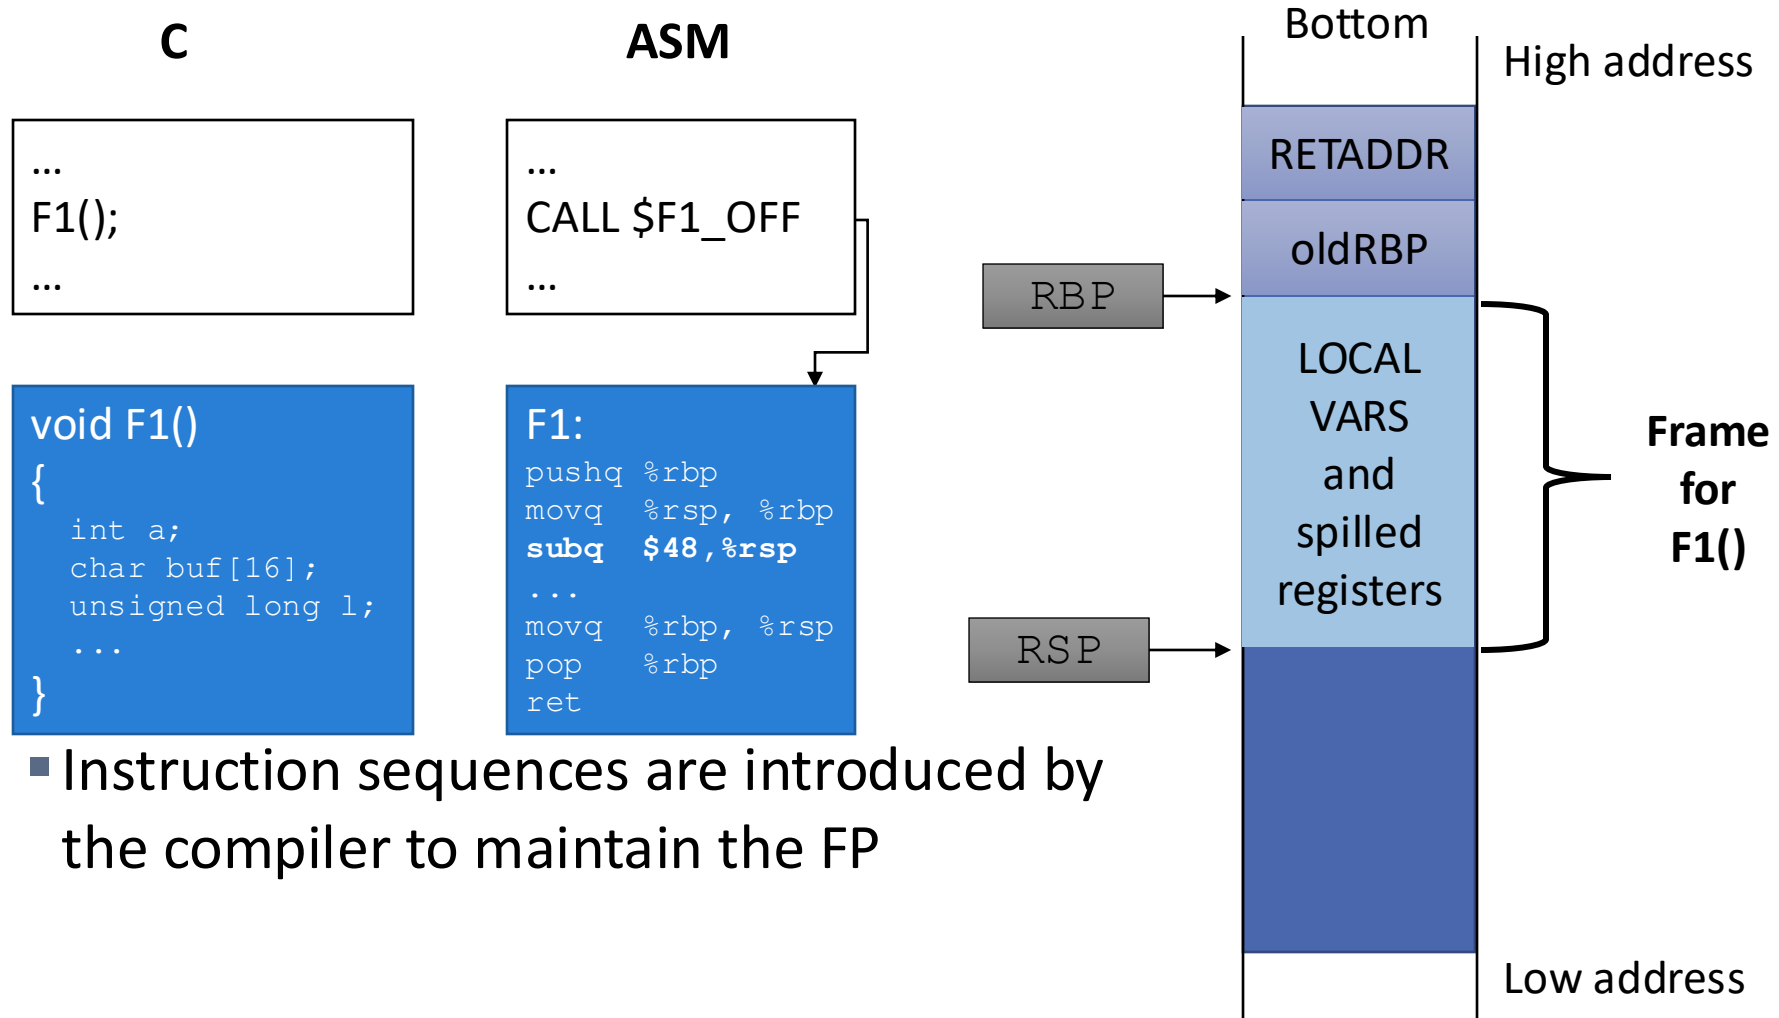
# Maintaining the FP



# Maintaining the FP

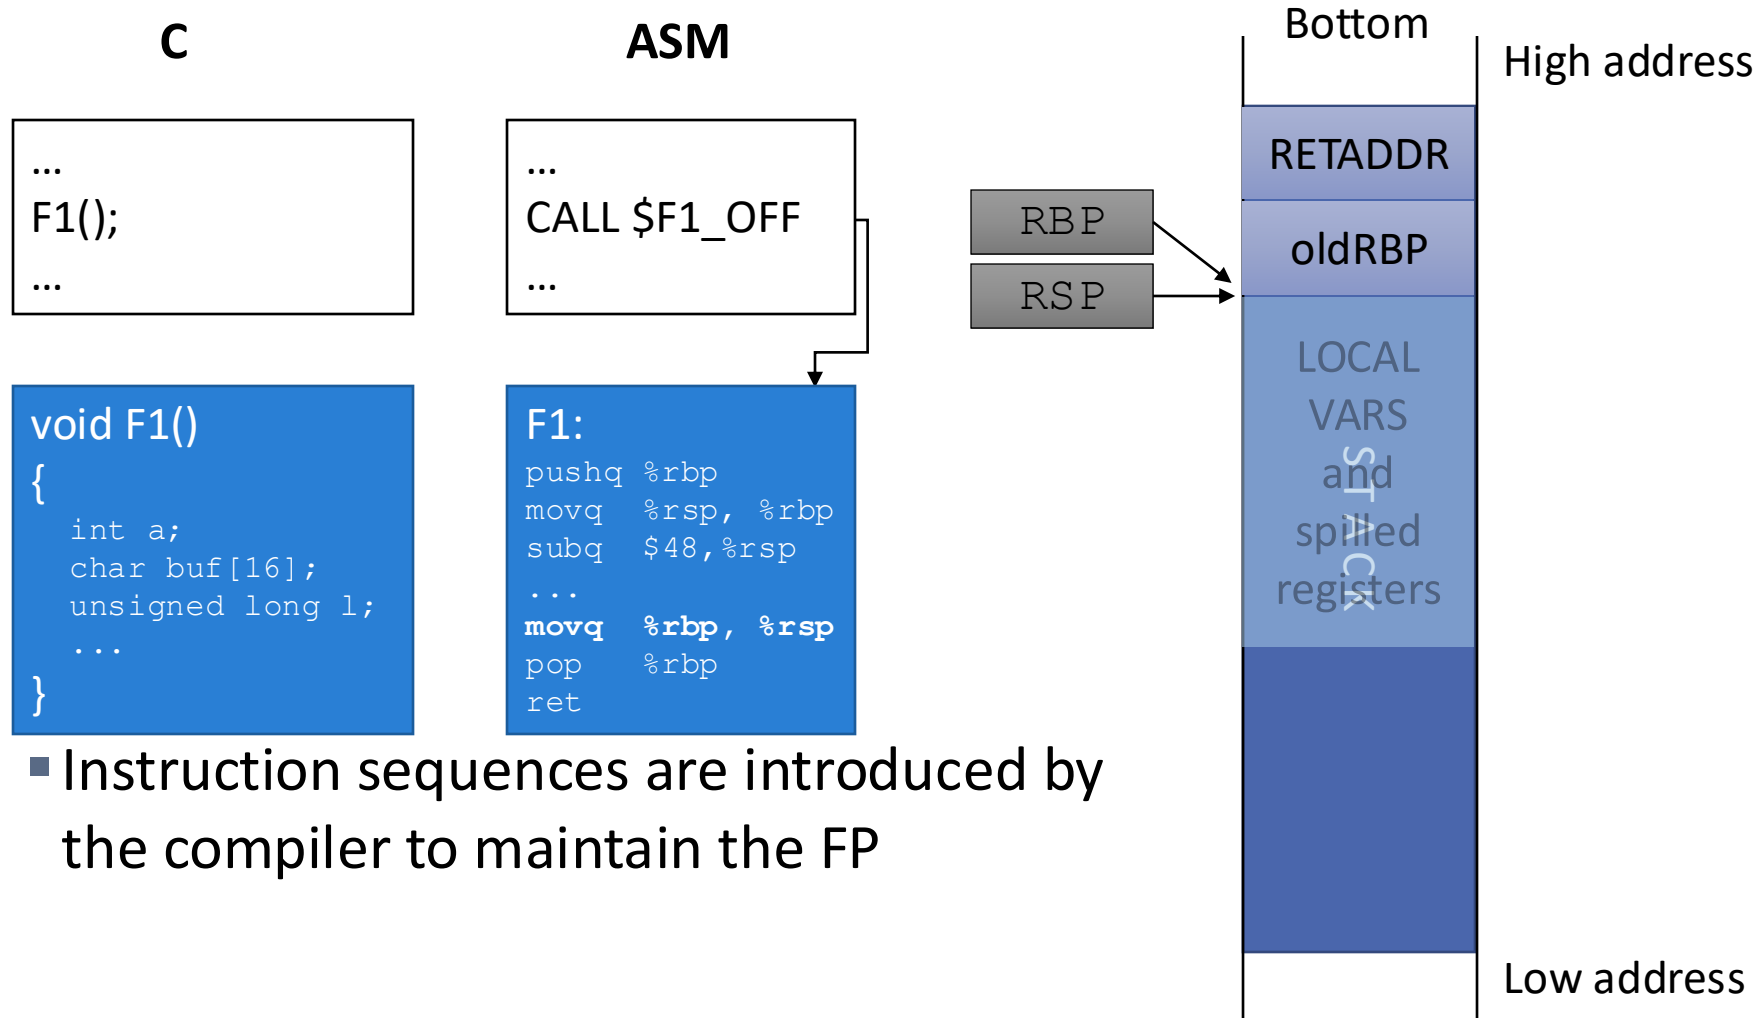


# Maintaining the FP

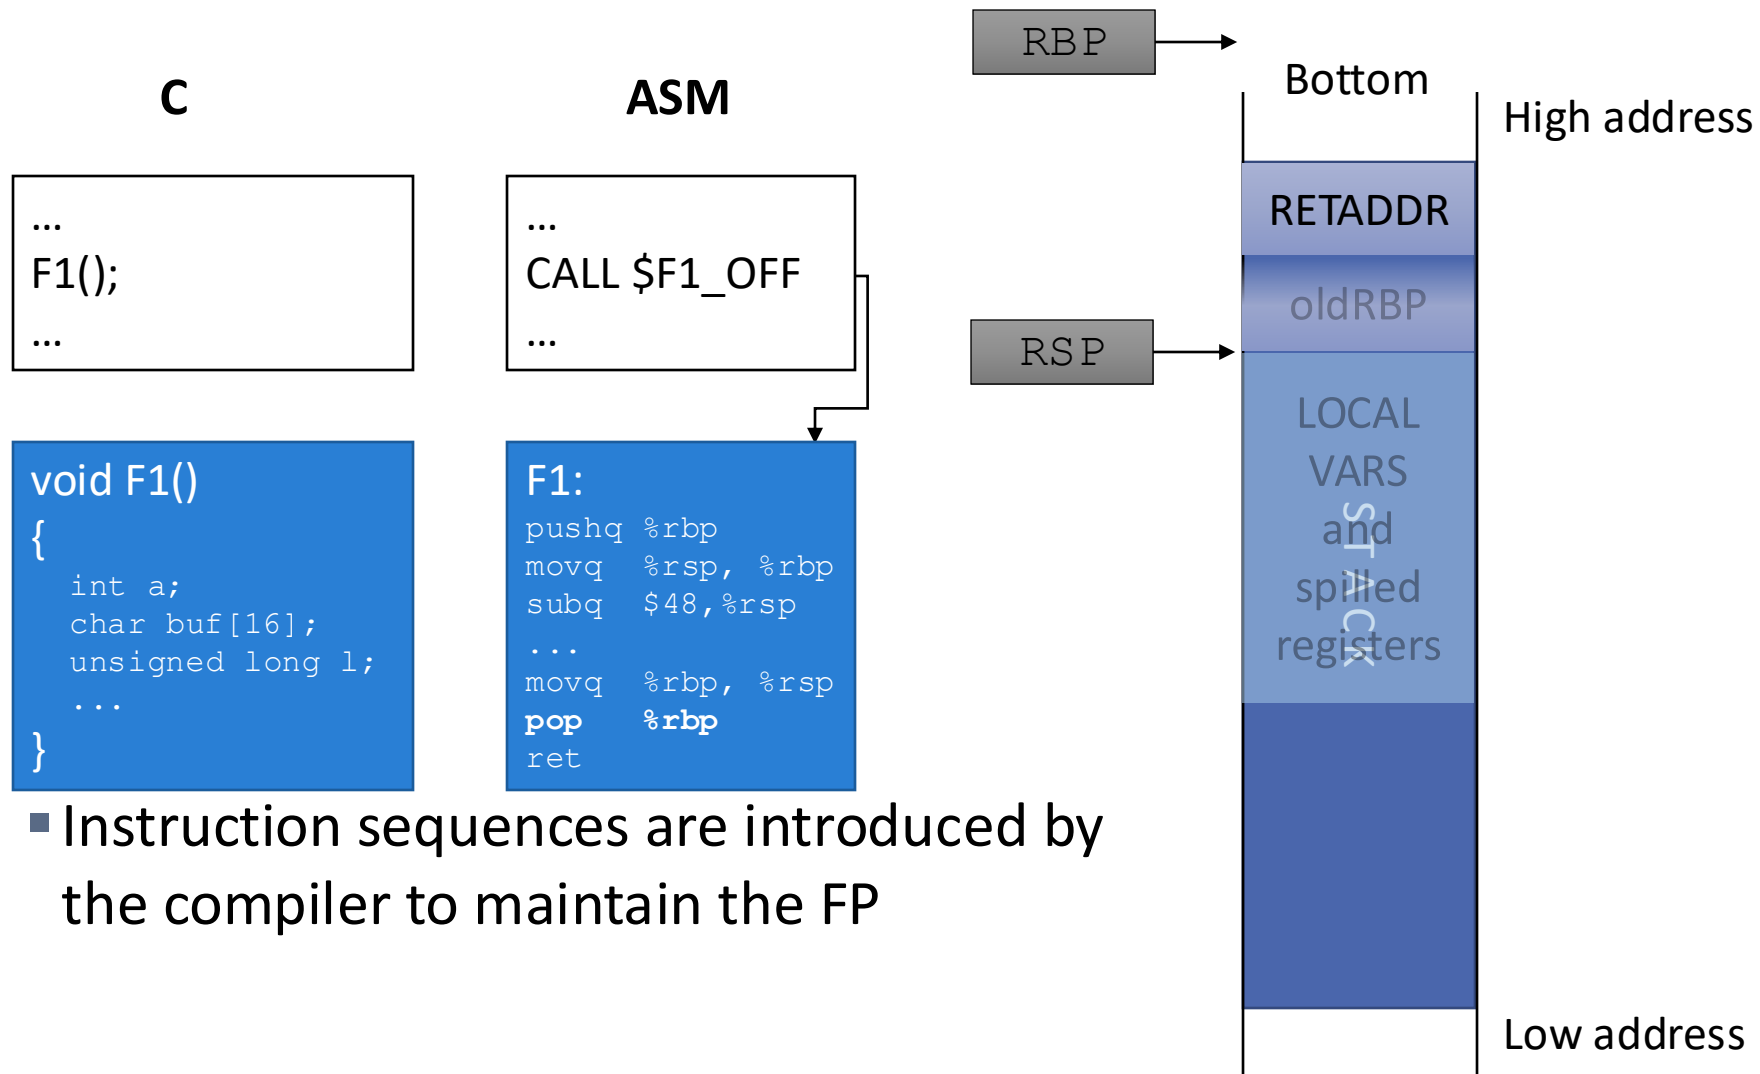


- Instruction sequences are introduced by the compiler to maintain the FP

# Maintaining the FP

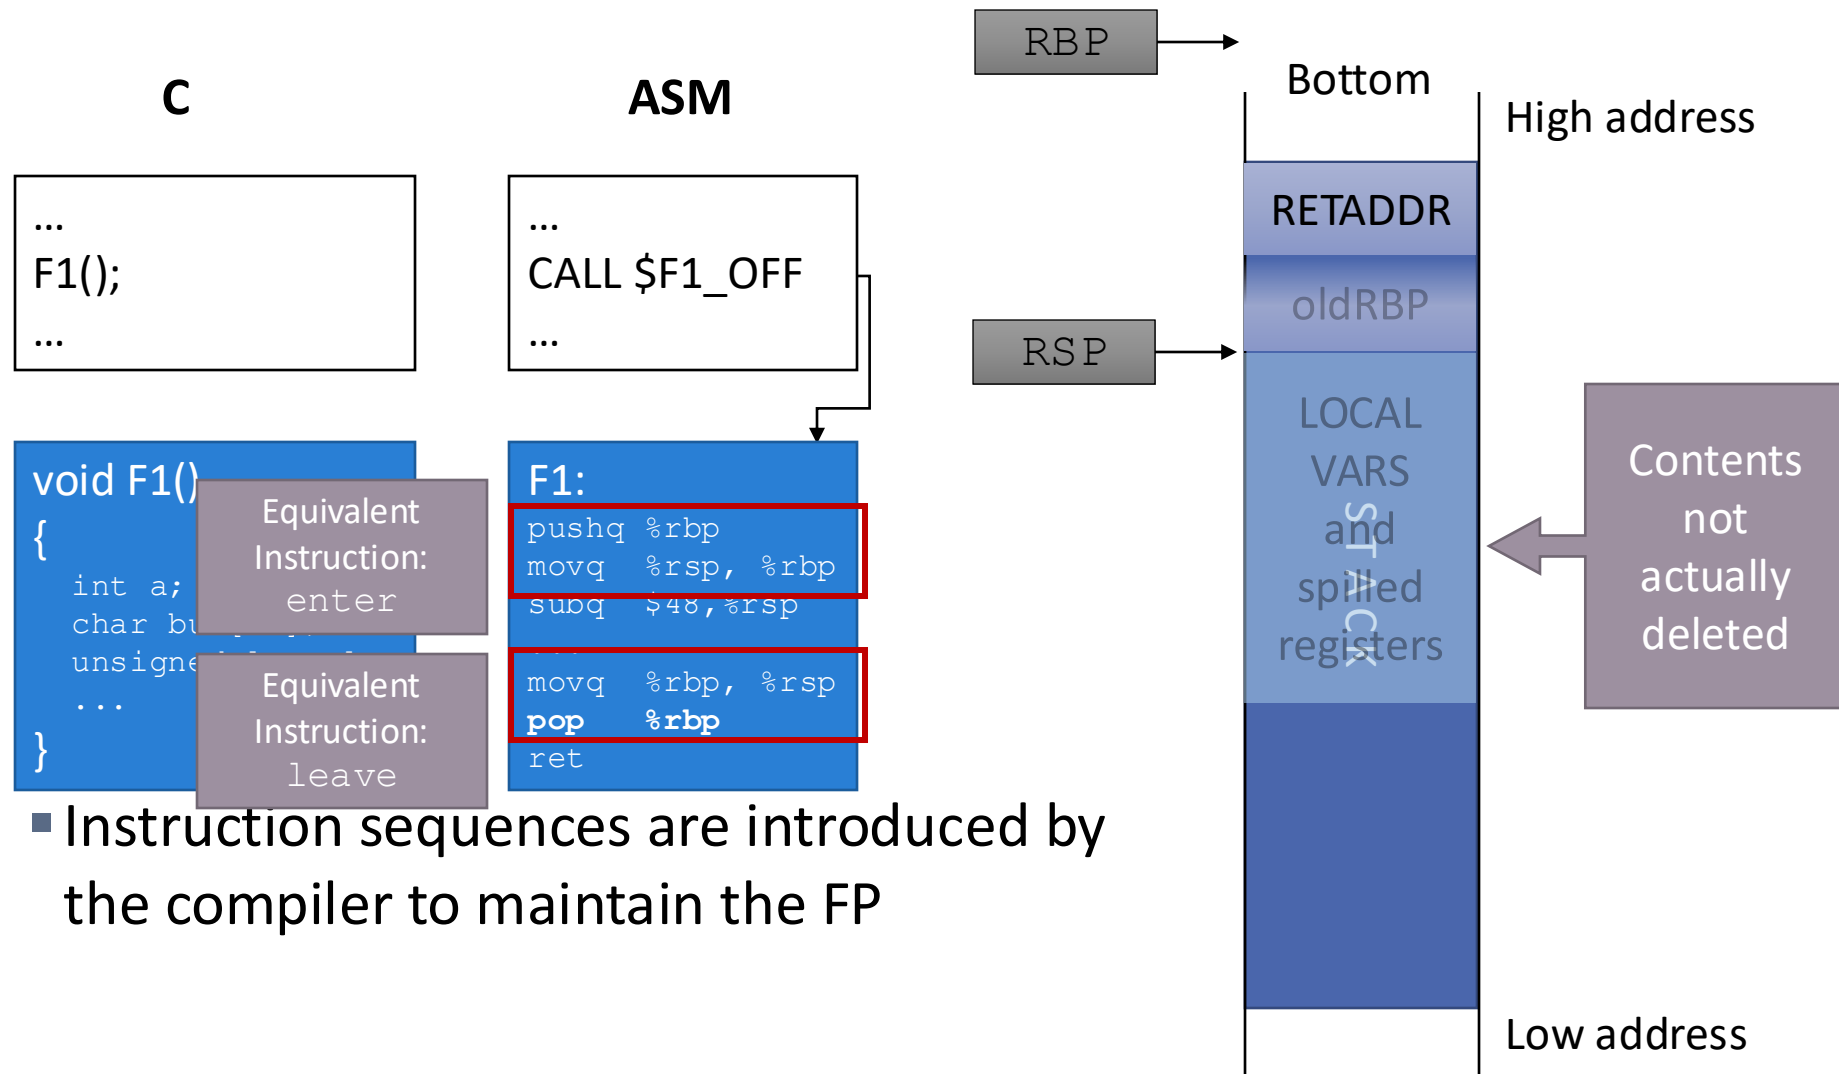


# Maintaining the FP



- Instruction sequences are introduced by the compiler to maintain the FP

# Maintaining the FP



# Putting It All Together

---

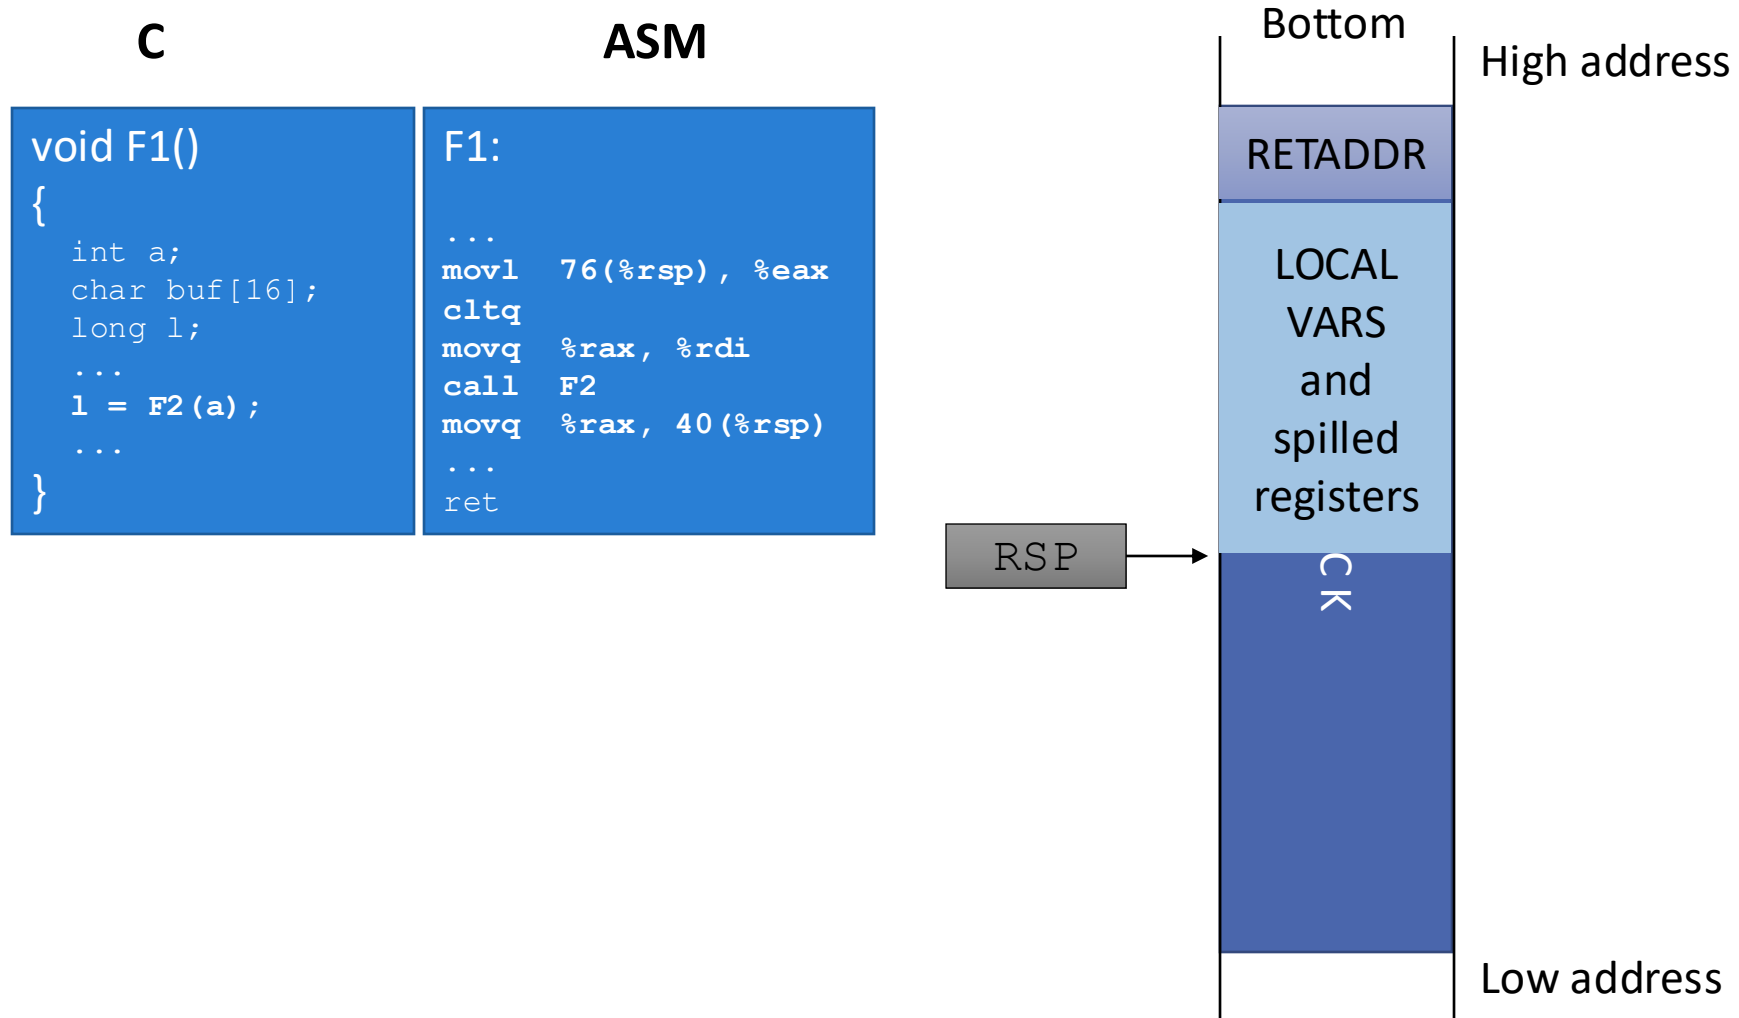
## Function Call

- Prepare function call arguments
- Make the call
- Function prologue
  - Save RBP/EBP
  - Setup new RBP/EBP
- Callee saves registers that need to be preserved
- Callee allocates stack space

## Function Return

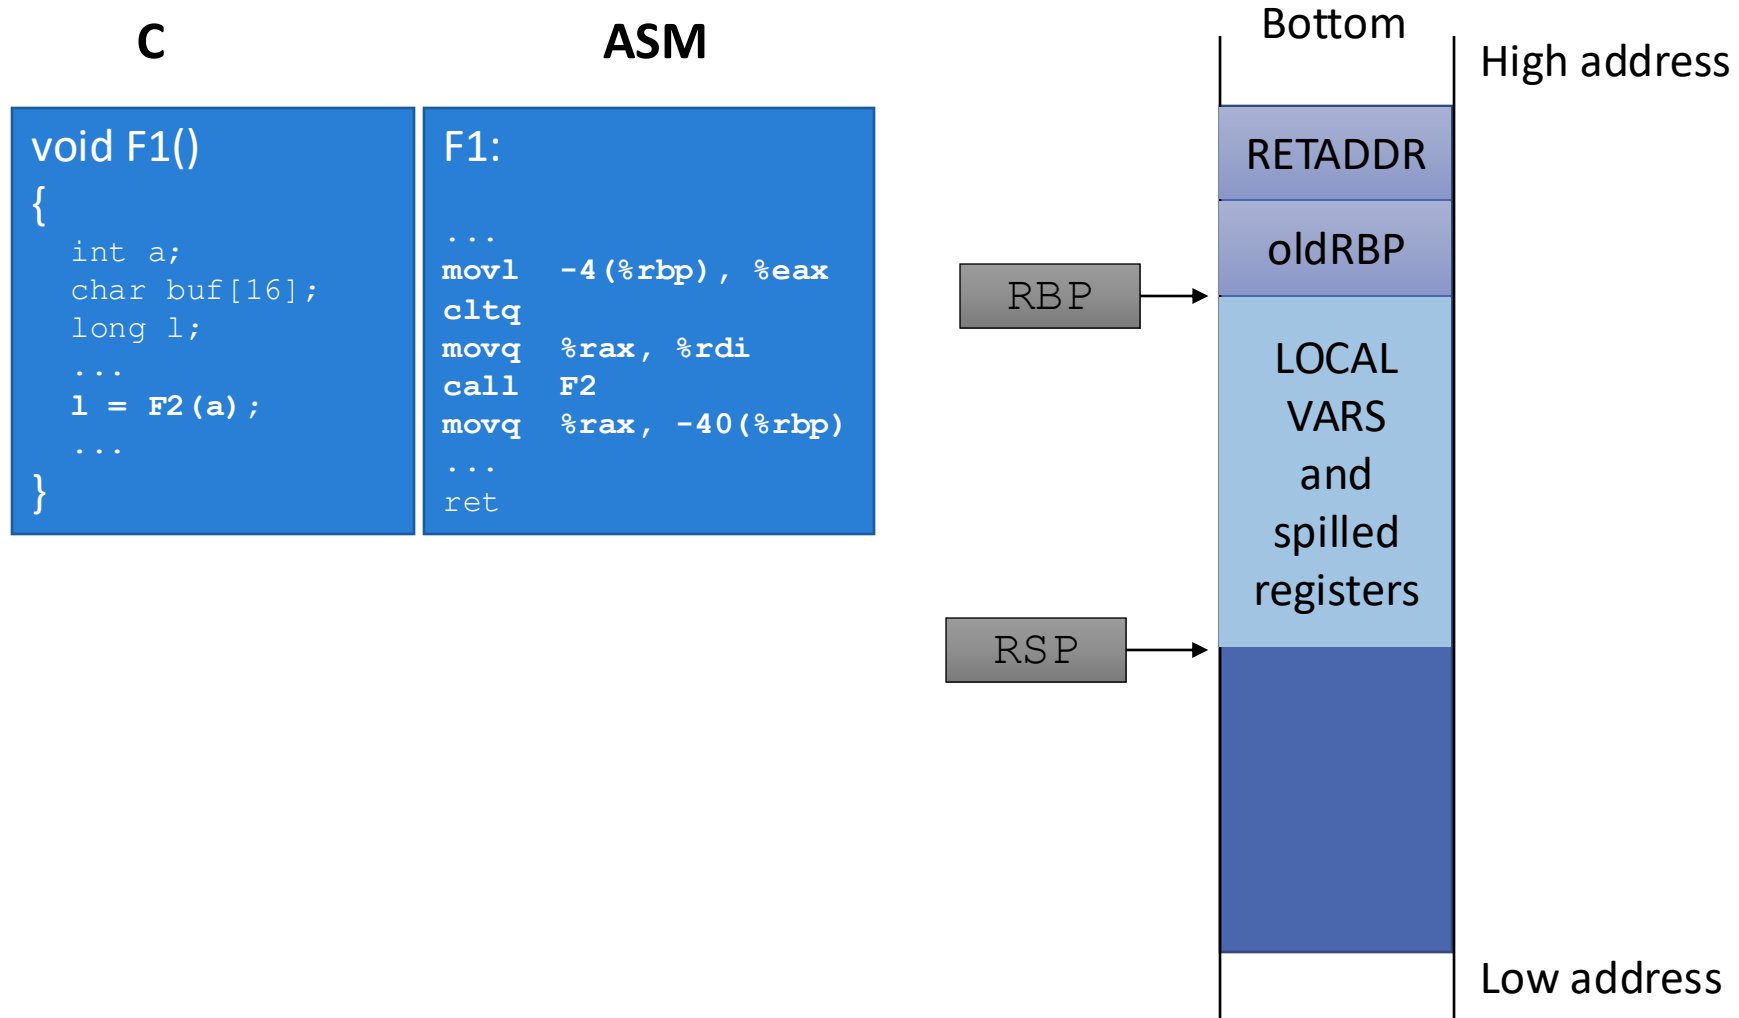
- Function epilogue
  - Release stack space
  - Restore BP
- Return

# Accessing Stack Variables (no FP)





# Accessing Stack Variables (with FP)



# Stack Smashing Attacks

---

# Stack Overflow Example

```
void copy(const char *str)
{
    char buf[16];

    strcpy(buf, str);

    puts(buf);
}
```

High address/stack bottom

RETADDR

buf

buf

STACK

Low address/stack top

# Stack Overflow Example

```
void copy(const char *str)
{
    char buf[16];

    strcpy(buf, str);

    puts(buf);
}
```

**`./copy AAAAAAAAA`**

High address/stack bottom

RETADDR

buf

buf

STACK

Low address/stack top

# Stack Overflow Example

```
void copy(const char *str)
{
    char buf[16];

    strcpy(buf, str);

    puts(buf);
}
```

**`./copy AAAAAAAAA`**

High address/stack bottom

RETADDR

\0???????

AAAAAAAAA

STACK

Low address/stack top

# Stack Overflow Example

```
void copy(const char *str)
{
    char buf[16];

    strcpy(buf, str);

    puts(buf);
}
```

```
./copy AAAAAAAAAAAAAAAAAAAAAA
```

High address/stack bottom

AAAAAAA\0

AAAAAAA

AAAAAAA

STACK

Low address/stack top

# Stack Overflow Example

```
void copy(const char *str)
{
    char buf[16];

    strcpy(buf, str);

    puts(buf);
}
```

```
subq    $40, %rsp
movq    %rdi, 8(%rsp)
movq    8(%rsp), %rdx
leaq    16(%rsp), %rax
movq    %rdx, %rsi
movq    %rax, %rdi
call    strcpy@PLT
leaq    16(%rsp), %rax
movq    %rax, %rdi
call    puts@PLT
nop
addq    $40, %rsp
ret
```

High address/stack bottom

AAAAAA\0

AAAAAA

AAAAAA

STACK

Low address/stack top

# Stack Overflow Example

- This stack overflow allows a to control the return address stored in the stack
- When `ret` executes, the control-flow of the program will be redirected to an arbitrary address → control-flow hijacking

```
subq    $40, %rsp
movq    %rdi, 8(%rsp)
movq    8(%rsp), %rdx
leaq    16(%rsp), %rax
movq    %rdx, %rsi
movq    %rax, %rdi
call    strcpy@PLT
leaq    16(%rsp), %rax
movq    %rax, %rdi
call    puts@PLT
nop
addq    $40, %rsp
ret
```

High address/stack bottom

AAAAAAA\0

AAAAAAA

AAAAAAA

STACK

Low address/stack top



# Control-Flow Hijacking Attacks

---

- Untrusted inputs that lead to corruption of a code pointer, which will be later dereferenced, lead to **control-flow hijacking attacks**


# Original Stack Smashing Attack

- Appeared at Phrack magazine  
<http://phrack.org/issues/49/14.html#article>
- Exploits the fact that stack used to be executable
  - Stores binary code in the controlled buffer
    - **Any** executable, controlled buffer will do!
  - Redirect program to inject code
- Performs arbitrary **code injection!**

```
# write(1, message, 13)
mov    $1, %rax          # system call 1 is write
mov    $1, %rdi          # file handle 1 is stdout
mov    $message, %rsi
mov    $13, %rdx         # number of bytes
syscall                    # invoke operating system to do the write


# exit(0)
mov    $60, %rax         # we want return code 0
xor    %rdi, %rdi        # invoke operating system to exit
syscall

message:
.ascii "Hello, world\n"
```



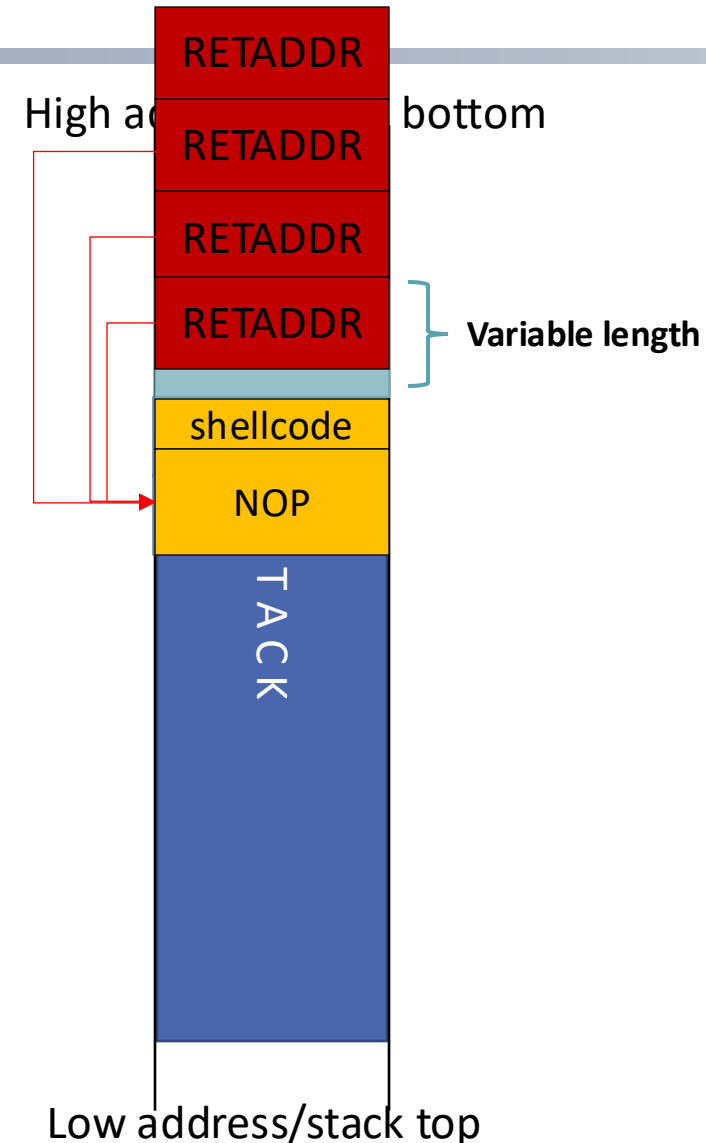
# Making Exploits More Robust

- Observation: Different compiler may use different alignment, spill different register, etc.
- Problems:
  - Exact distance of return address may be different between binaries
  - Exact address of buffer may be different



# Making Exploits More Robust

- Observation: Different compiler may use different alignment, spill different register, etc.
- Problems:
  - Exact distance of return address may be different between binaries
  - Exact address of buffer may be different
- Solutions:
  - Use multiple copies of the target address
  - Prepend a **NOP sled** to shellcode
    - NOPs → No operations are special one byte instructions to do nothing
  - Aim for target address pointing into NOP sled
    - Execution will slide into shellcode



# Appendix: cdecl Calling Convention (Optional)


---

# cdecl

- Arguments are passed on the stack
  - Pushed right to left

```
...  
F1(0xff, UINT_MAX, argv[0]);  
...
```

```
...  
pushl    (%eax)  
pushl    $-1  
pushl    $255  
call     F1  
...
```




# cdecl

- Arguments are passed on the stack
  - Pushed right to left


```
...  
F1(0xff, UINT_MAX, argv[0]);  
...
```

```
...  
pushl    (%eax)  
pushl    $-1  
pushl    $255  
call     F1  
...
```



# cdecl

- Arguments are passed on the stack
  - Pushed right to left
- `eax`, `edx`, `ecx` are caller saved
  - callee can overwrite without saving
- `ebx`, `esi`, `edi` are callee saved
  - callee must ensure they have same value on return



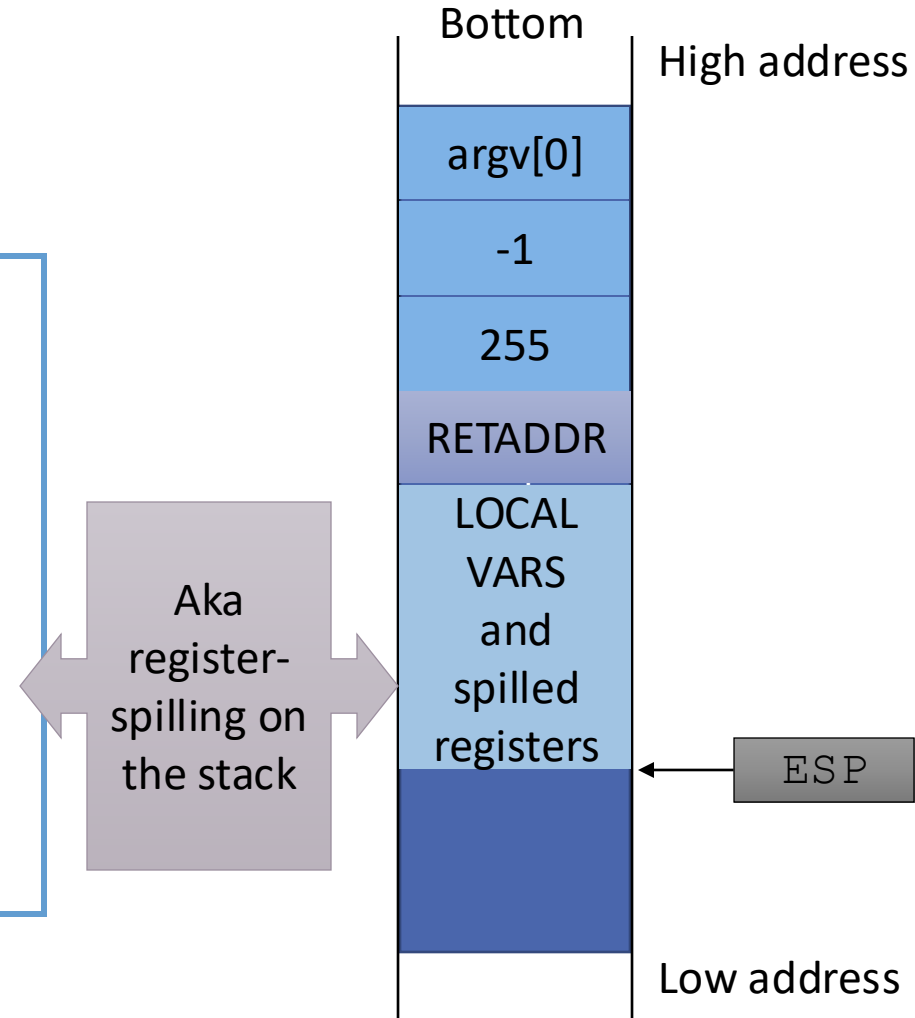


# cdecl

- Arguments are passed on the stack
  - Pushed right to left

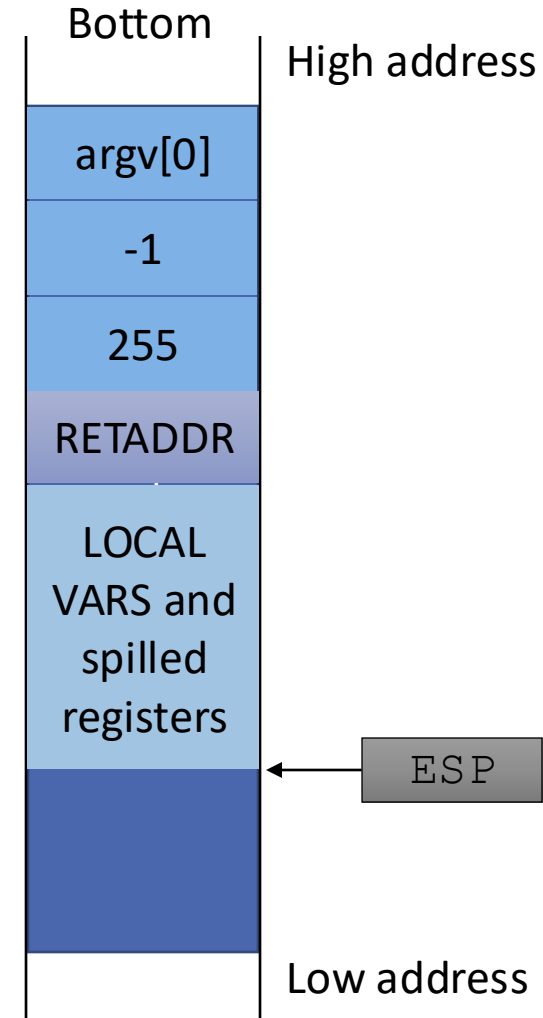
- `eax`, `edx`, `ecx` are caller saved
  - callee can overwrite without saving

- `ebx`, `esi`, `edi` are callee saved
  - callee must ensure they have same value on return



# cdecl

- Arguments are passed on the stack
  - Pushed right to left
- `eax`, `edx`, `ecx` are caller saved
  - callee can overwrite without saving
- `ebx`, `esi`, `edi` are callee saved
  - callee must ensure they have same value on return
- `eax` used for function return value



## Part 3. Data Execution Prevention

---

# Writable and Executable Memory

---

- Code injection is possible because there is a memory area that is both writable and executable

# Writable and Executable Memory

- Code injection is possible because there is a memory area that is both writable and executable
- Can be eliminated if we introduce the following policy

## **W<sup>X</sup> Policy**

The Write XOR Execute (W<sup>X</sup>) policy mandates that in a program there are no memory pages that are both writable and executable

# Writable and Executable Memory

- Code injection is possible because there is a memory area that is both writable and executable
- Can be eliminated if we introduce the following policy

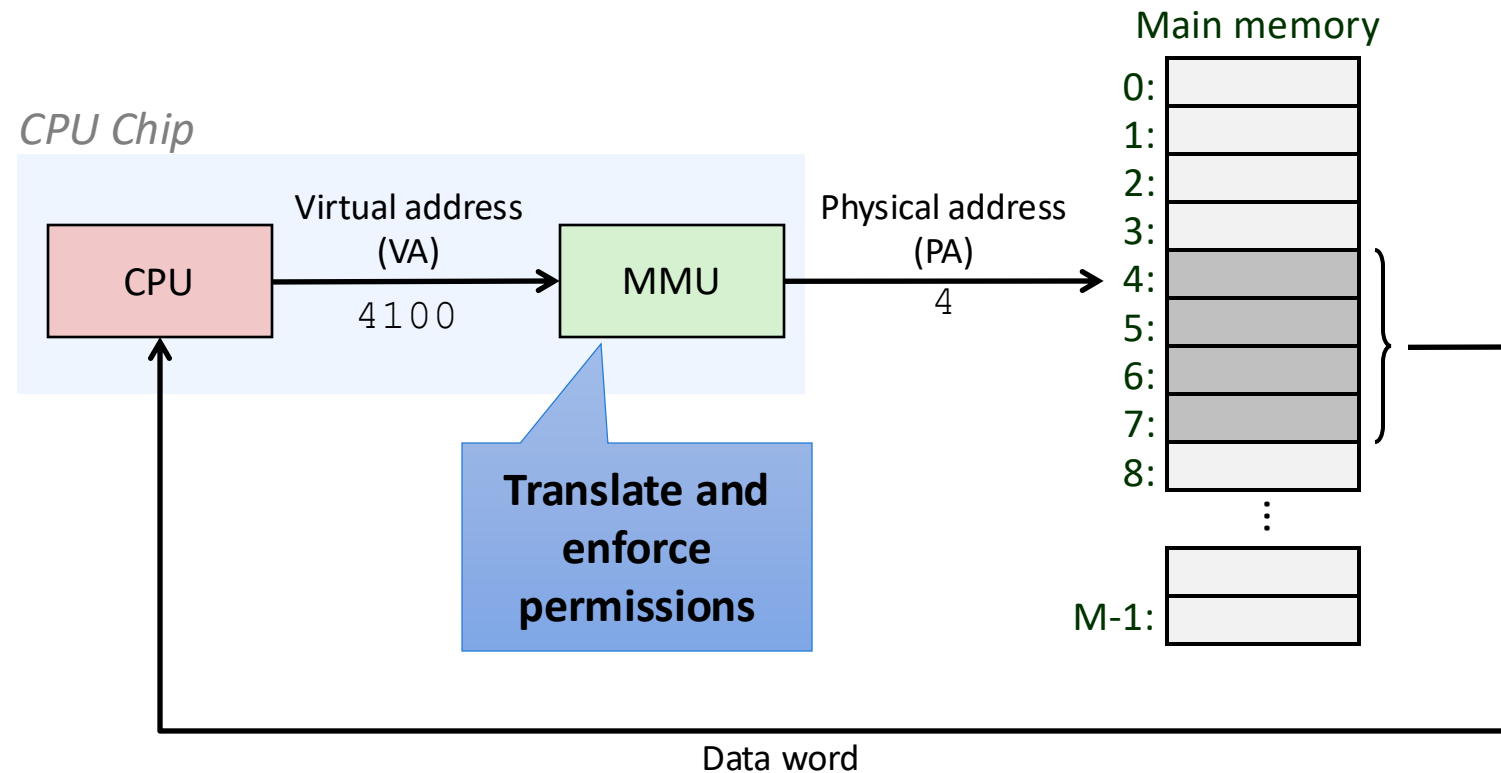
## W<sup>X</sup> Policy

How?

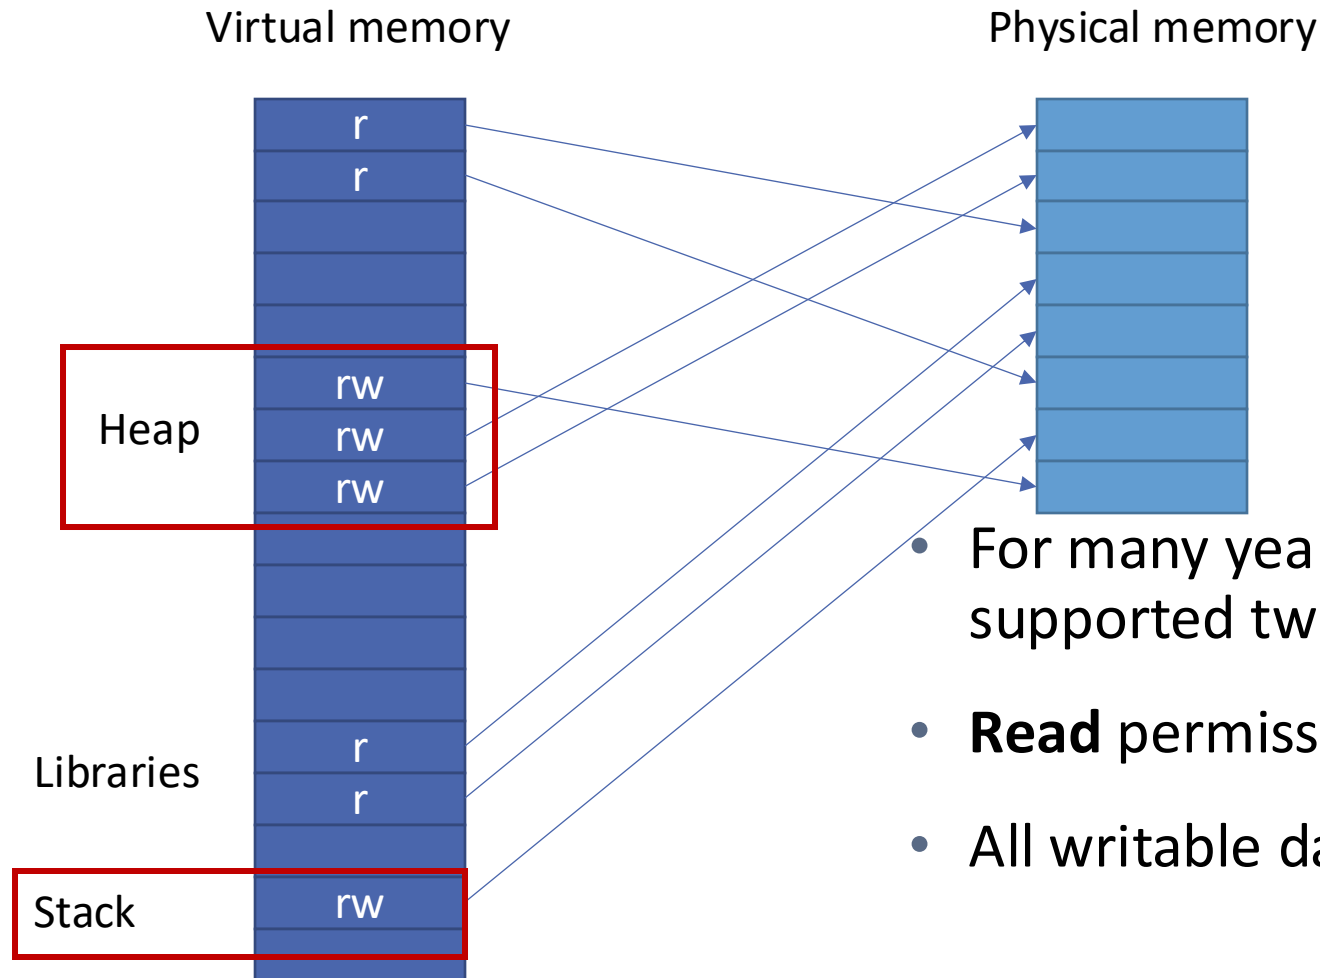
The Write XOR Execute (W<sup>X</sup>) policy mandates that in a program there are no memory pages that are both writable and executable

# The Memory Management Unit (MMU) - Paging

- Used in all modern servers, laptops, and smart phones
- One of the great ideas in computer science



# Page Permissions



- For many years (<2000), processors supported two permissions (bits)
- **Read** permission implied **Execute**
- All writable data segments violated  $W^X$



# Hardware Support: NX-bit

---

- Processor manufacturers introduced a new bit in page permissions to prevents code injections
- Coined **No-eXecute** or **Execute Never**
- The NX-bit (No-eXecute) was introduced first by AMD to resolve such issues in 2001
  - Asserting NX, makes a readable page non-executable
  - Frequently referred to as Data Execution Prevention (DEP) on Windows
- **Marketed as antivirus technology**

[Blog](#)[Bulletin](#)[VB](#)

# Enhanced virus protection

**Costin Raiu** *Kaspersky Lab*

[download slides](#) (PDF)

AMD Athlon 64 CPU Feature:

1. HyperTransport technology
2. Cool'n'Quiet technology
3. Enhanced Virus Protection for Microsoft Windows XP SP2

The AMD64 architecture is an affordable way of getting the power of 64-bit processing into a desktop computer. Interesting enough, AMD has not only designed an improved CPU core and longer registers, but they have also included a feature designed to significantly increase the security of modern operating systems.

The idea of hardware protection isn't new – every contemporary CPU includes at least a basic hardware mechanism for enforcing a security scheme, for instance, those from the Intel x86 family, based on

# Adoption

---

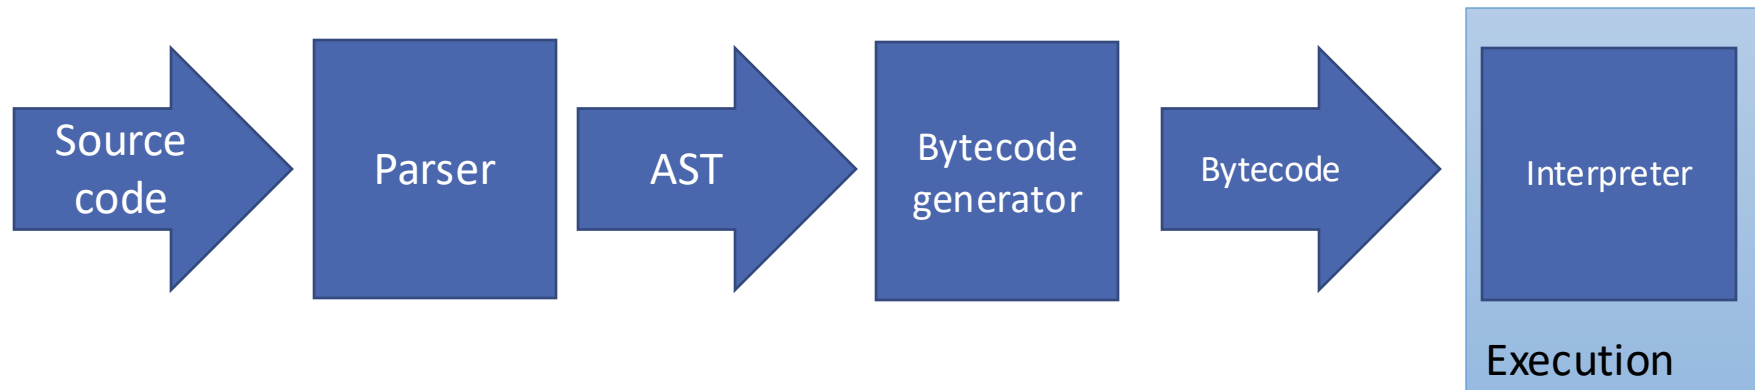
- A non-executable stack was not immediately adopted
- The OS occasionally needed to place code in the stack
  - For example, trampoline code for handling UNIX signals
- Widely adopted today

# Unless You Are a Browser...

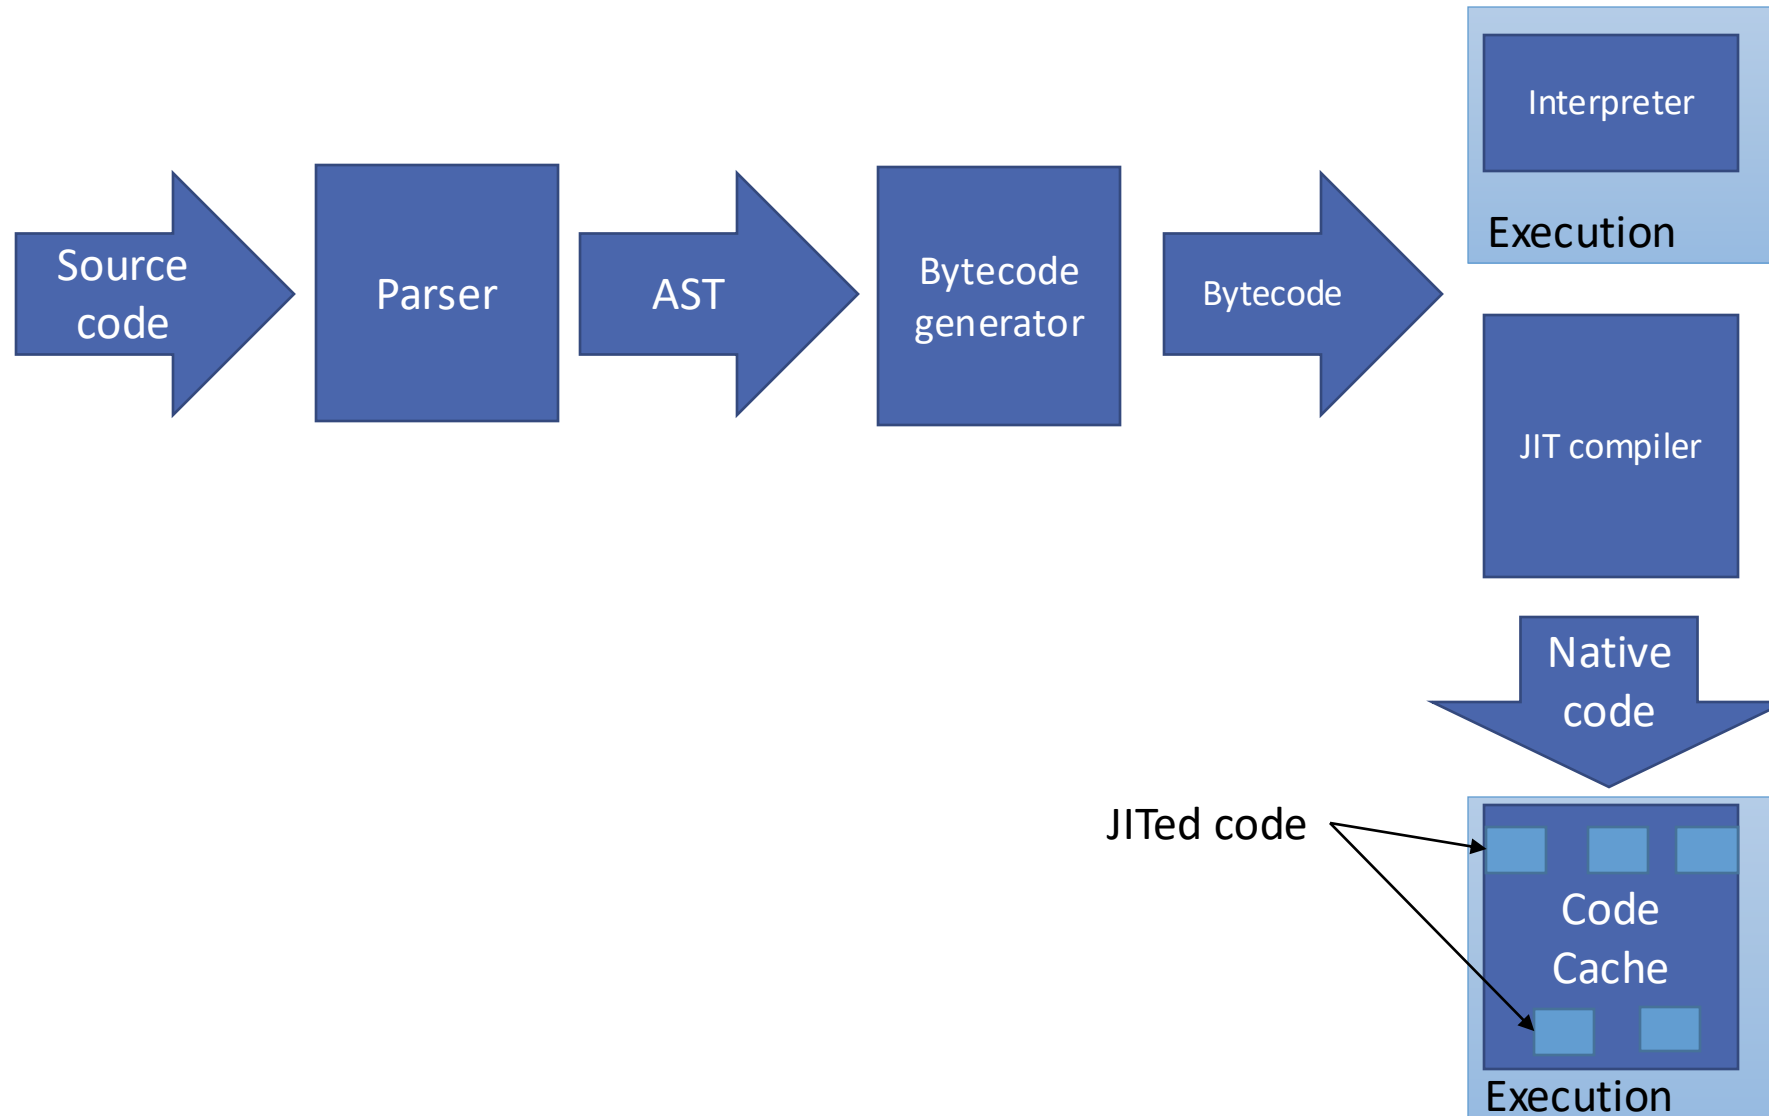
---

- Very popular software
  - Probably installed on every client device
- Large and complex software
- Execute JavaScript

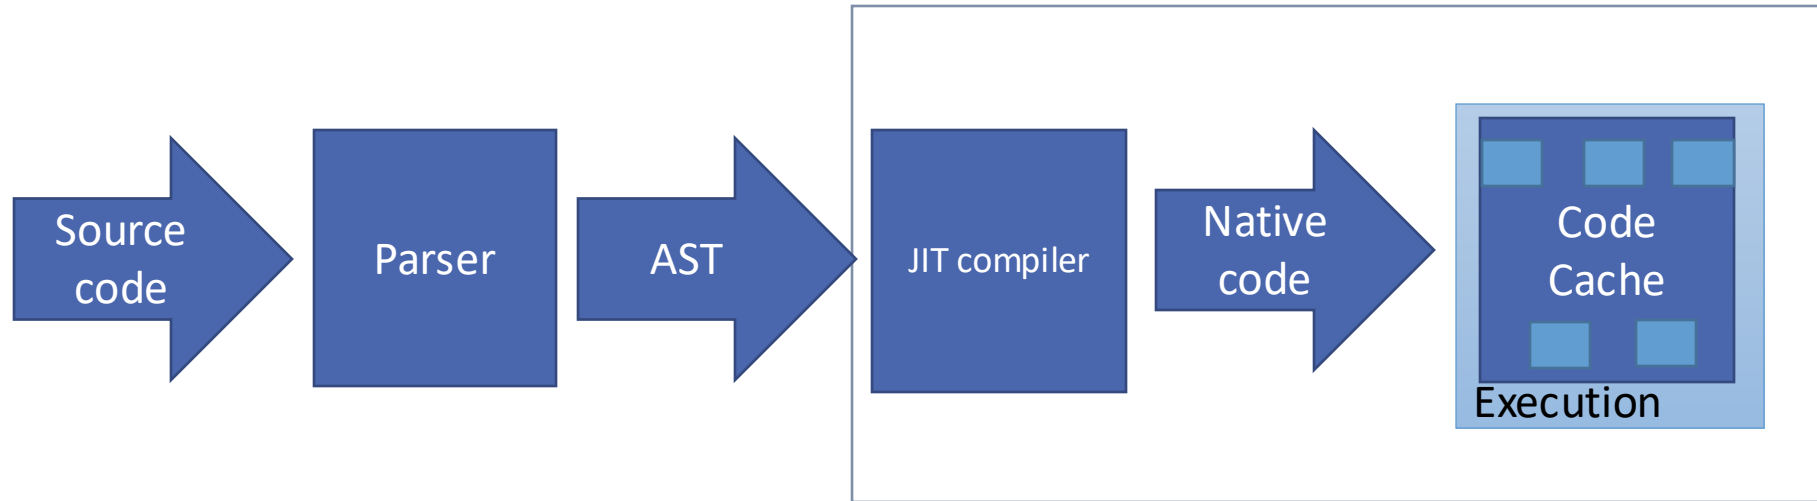
# How Does JavaScript Run



# How Does JavaScript Run




# How Does JavaScript Run



- Google V8 designed specifically to execute at speed.
- Bytecode generation skipped
- Directly emit native code
- Overall JavaScript execution improved by 150%

# Code Cache

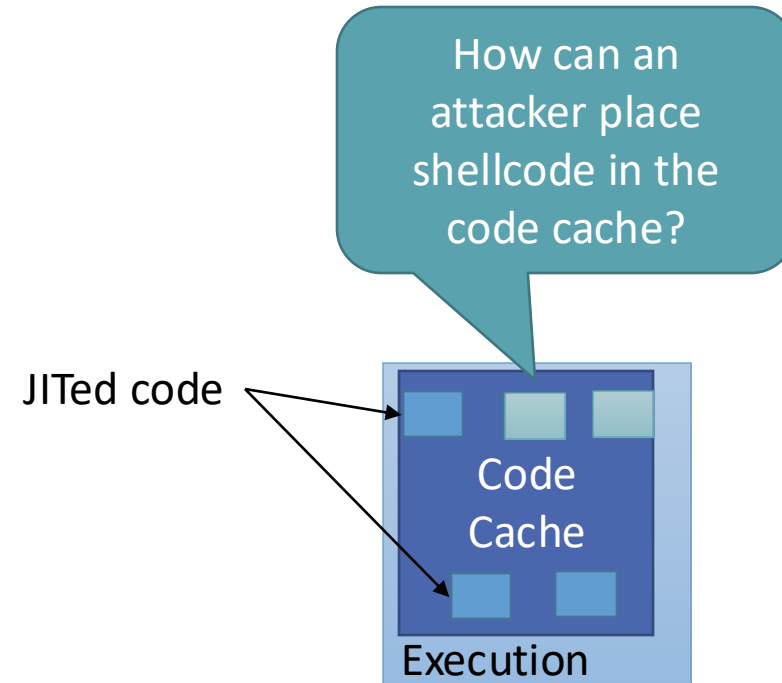
- JITed code and code cache have interesting properties from the perspective of the attacker
  - Code is continuously generated
  - Code needs to be executable
- **Violates the W<sup>X</sup> policy**





# Code Cache

- JITed code and code cache have interesting properties from the perspective of the attacker
  - Code is continuously generated
  - Code needs to be executable
- **Violates the W<sup>X</sup> policy**



# From JS to Code Cache

- JS code is JITed and placed in the code cache
- Some JS engines do not separate data and code

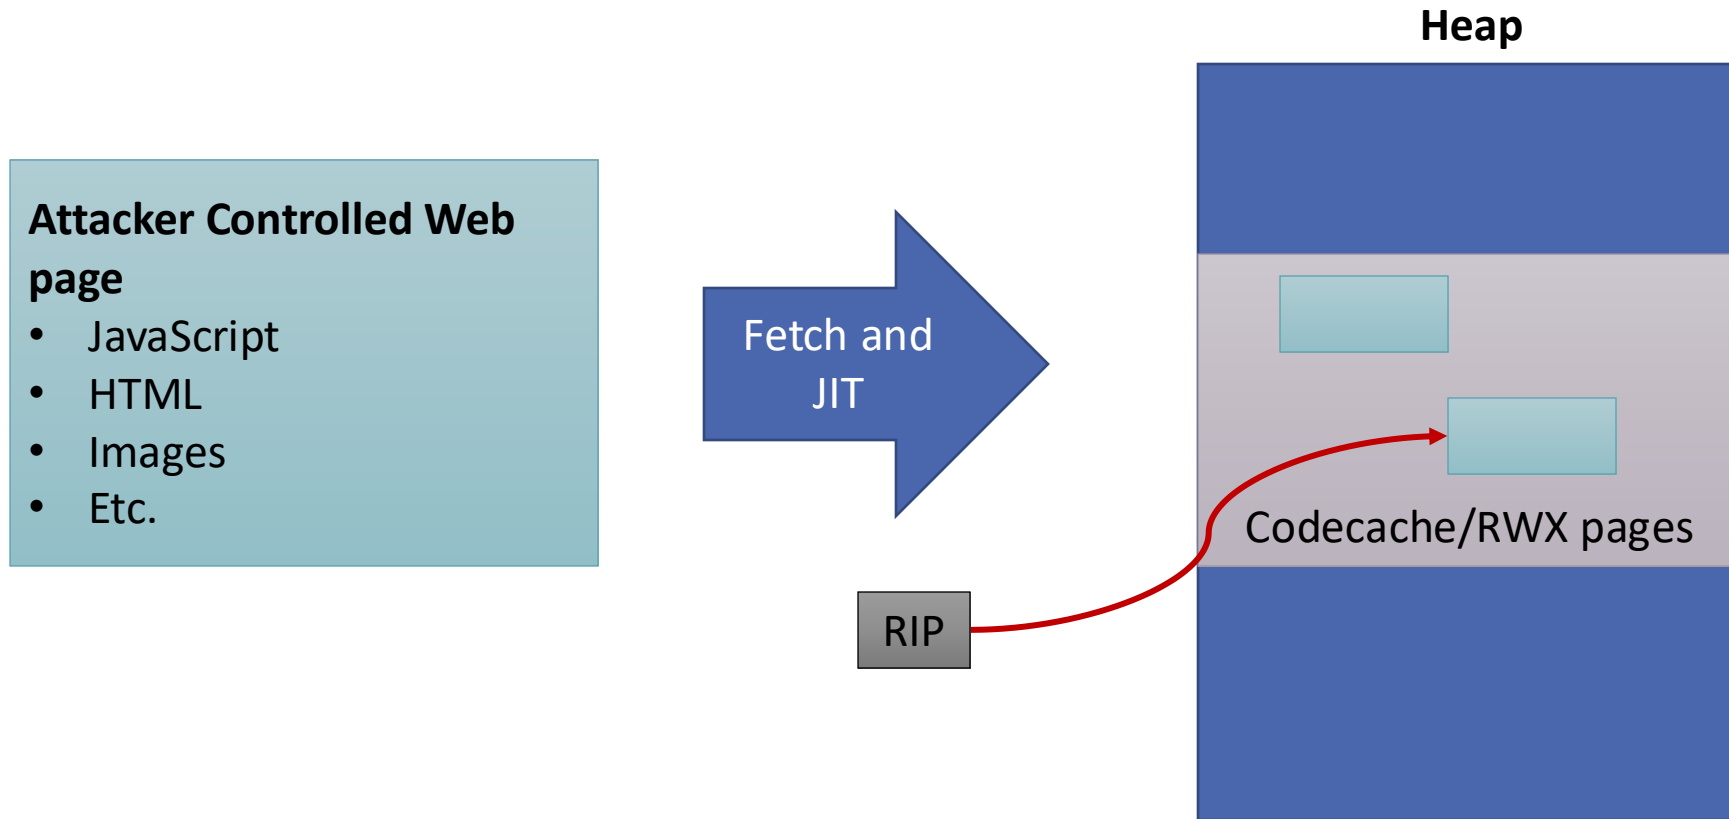
```
<html>
<body>
<script language='javascript'>

var myvar = unescape('%u\4F43%u\4552'); // CORE
myvar += unescape('%u\414C%u\214E'); // LAN!
alert("allocation done");

</script>
</body>
</html>
```

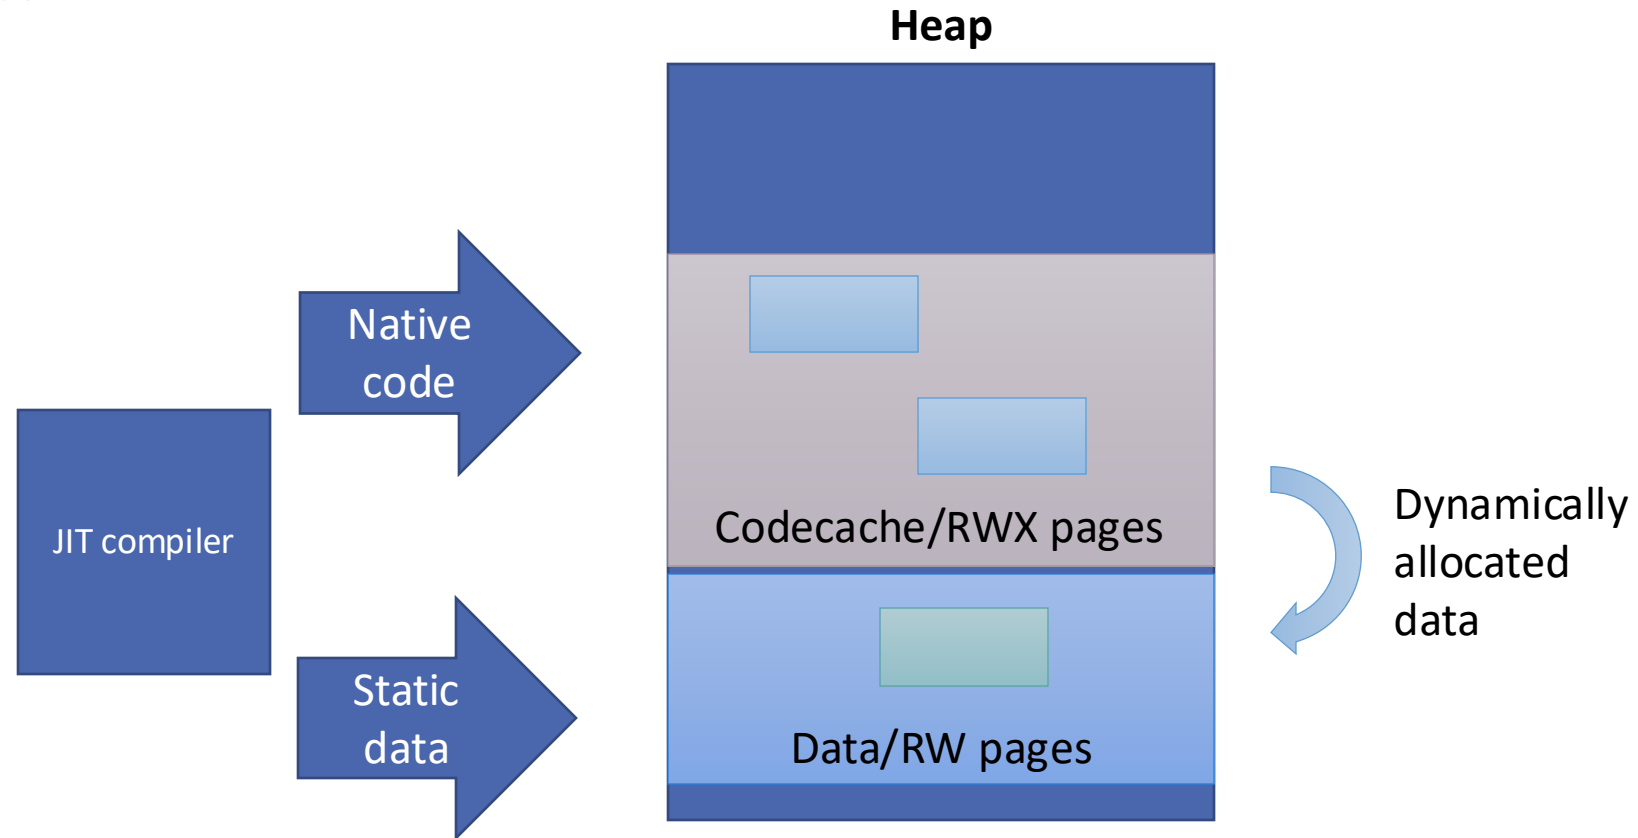
# Code-Injection Attacks Against Browsers

- Return to code injected in the codecache

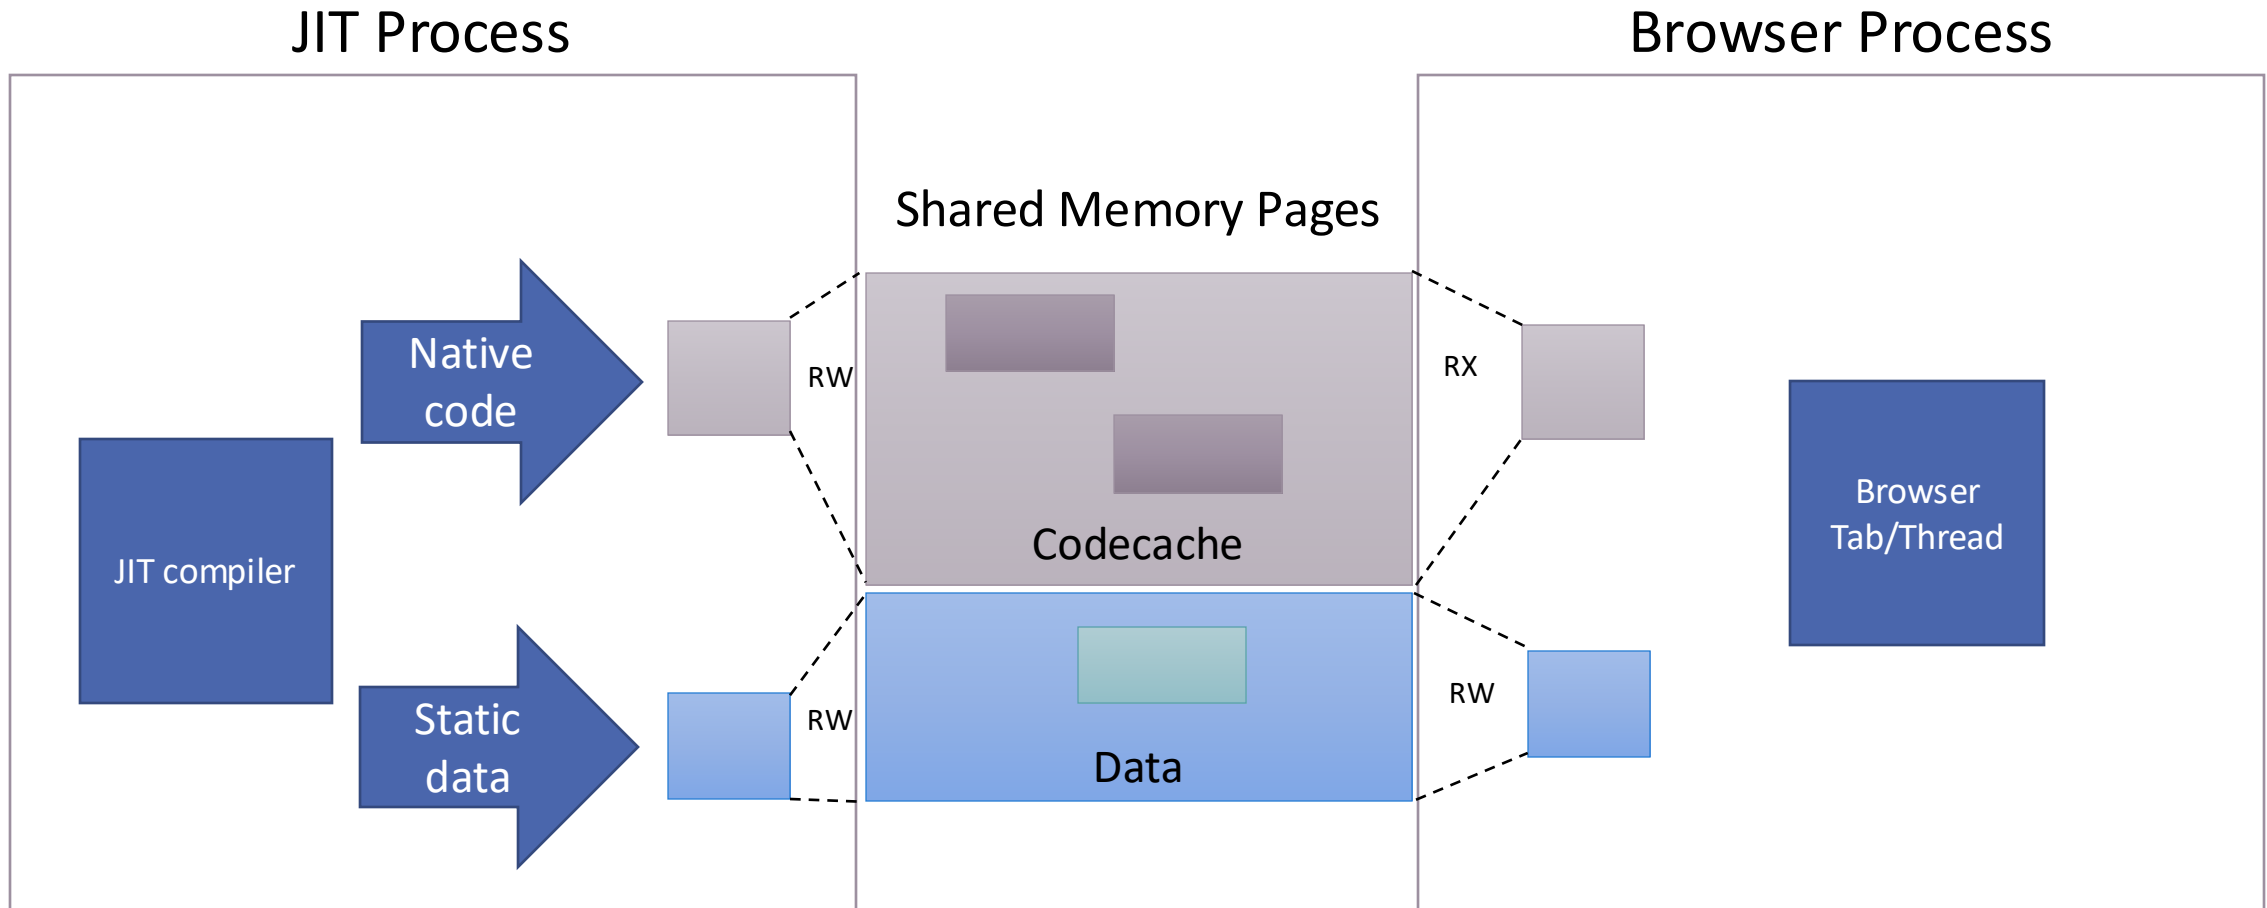


# Avoiding Code Injection in Browsers

- Separate code and data into separate memory areas
- Still violates W<sup>X</sup>



# W^X Semantics in Browser Processes



# Additional Reading (Optional)

---

- The Devil is in the Constants: Bypassing Defenses in Browser JIT Engines
  - [https://www.portokalidis.net/files/devilinconstants\\_ndss15.pdf](https://www.portokalidis.net/files/devilinconstants_ndss15.pdf)
- libmpk: Software Abstraction for Intel Memory Protection Keys (Intel MPK)
  - <https://www.usenix.org/system/files/atc19-park-soyeon.pdf>

# Appendix: Emulating NX-bit (Optional)

---


# Early Approaches: PAGEEXEC

---

- A Linux kernel patch emulating non-executable memory
- Introduced in 2000 by the PaX team
- PAGEEXEC refused code execution on selected writable pages
  - Heap and stack



# Emulating Non-Executable Memory



- Mark writable pages so that access causes a page fault
  - Not present → a page-fault will be raised on every access
  - With supervisor bit (S) → Access only allowed from the kernel
- Custom page-fault handler intercepts and checks accesses:
  - Fault caused by other instruction → data access → OK
  - Faulting address is being executed → code execution → Violation

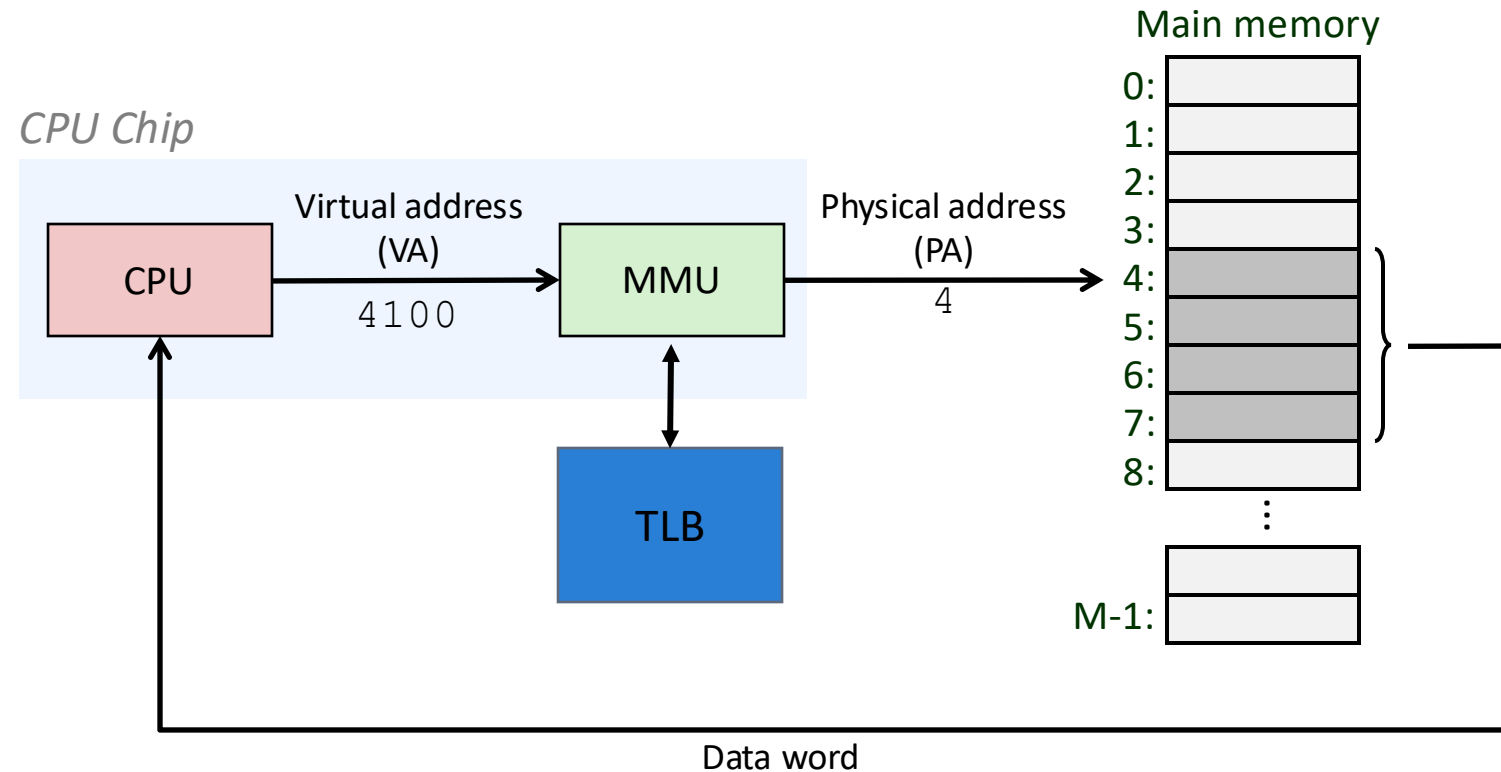
# Emulating Non-Executable Memory

**Should be very  
expensive....**

- Mark writable pages so that access causes a page fault
  - Not present → a page-fault will be raised on every access
  - With supervisor bit (S) → Access only allowed from the kernel
- Custom page-fault handler intercepts and checks accesses:
  - Fault caused by other instruction → data access → **OK**
  - Faulting address is being executed → code execution → **Violation**

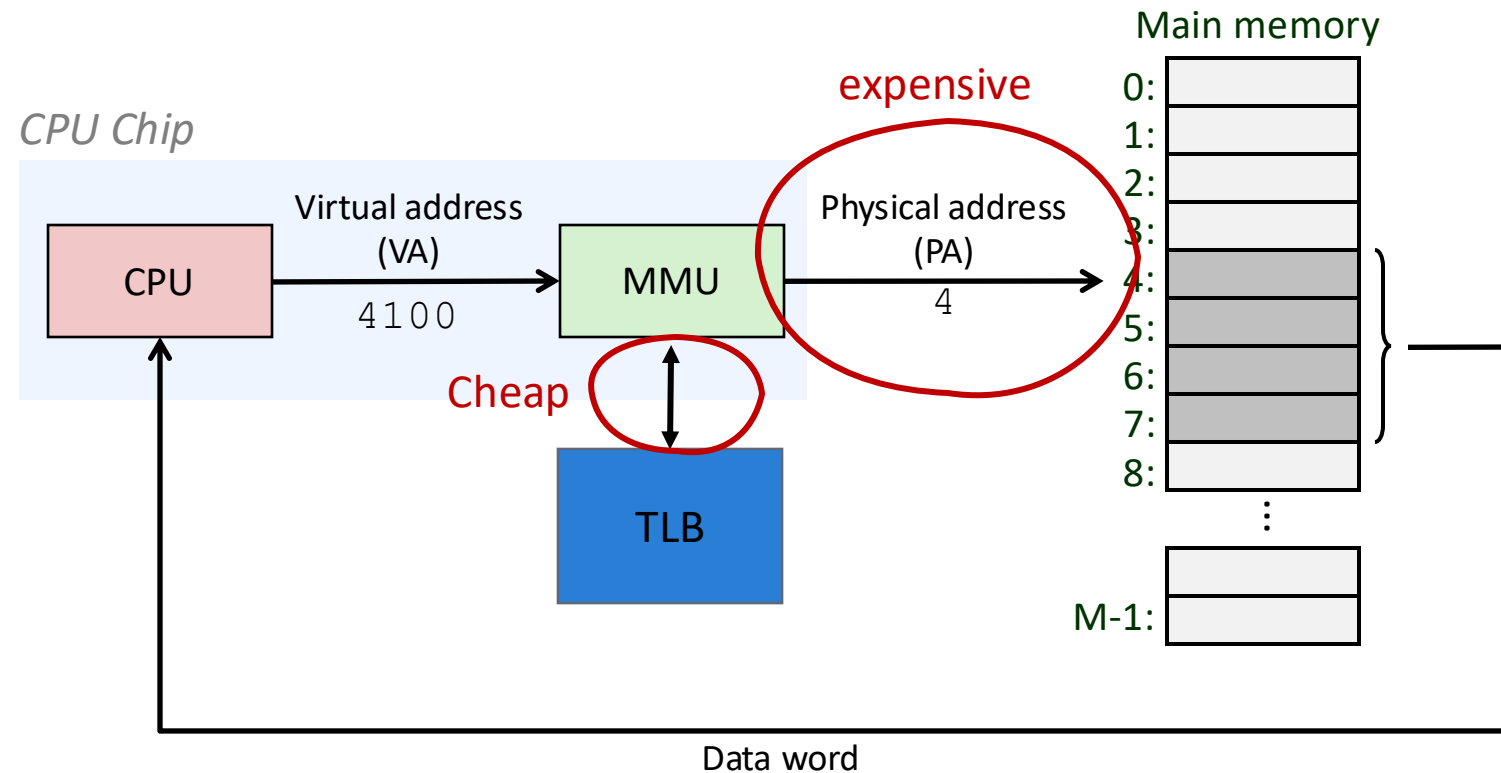
# Translation Lookaside Buffer (TLB)

- A cache for storing the translations for the most frequently accessed pages



# Translation Lookaside Buffer (TLB)

- A cache for storing the translations for the most frequently accessed pages



# Split TLBs

- Instruction TLB (ITLB) used when fetching bytes to be decoded and executed as an instruction
  - PC → memory addr
  - addr → ITLB
- Data TLB (DTLB) used when reading/write bytes required by the executing instruction
  - (addr) → memory addr
    - Example: mov (addr), reg
  - addr → DTLB

# Split TLBs and PAGEEXEC

---

- Fault caused because PC points within data area → Violation
- Fault caused by other access
  - Remove supervisor bit from page
  - Complete load which will be added to the DTLB
  - Add supervisor bit to page
  - Subsequent accesses to address will be served by the DTLB
    - Until it is flushed or the entry for the address evicted